



UNIVERSITÀ  
DEGLI STUDI  
DI BERGAMO

# Tecnologie Cloud & Mobile

TECNOLOGIE CLOUD  
E MOBILE

CLOUD SYSTEMS AND INFRASTRUCTURE

[mauro.pelucchi@gmail.com](mailto:mauro.pelucchi@gmail.com)

Mauro Pelucchi

21059 - Tecnologie cloud e mobile

# Cosa vedremo?

- Cloud computing
- AWS
- Micro-services
- Server-less architecture

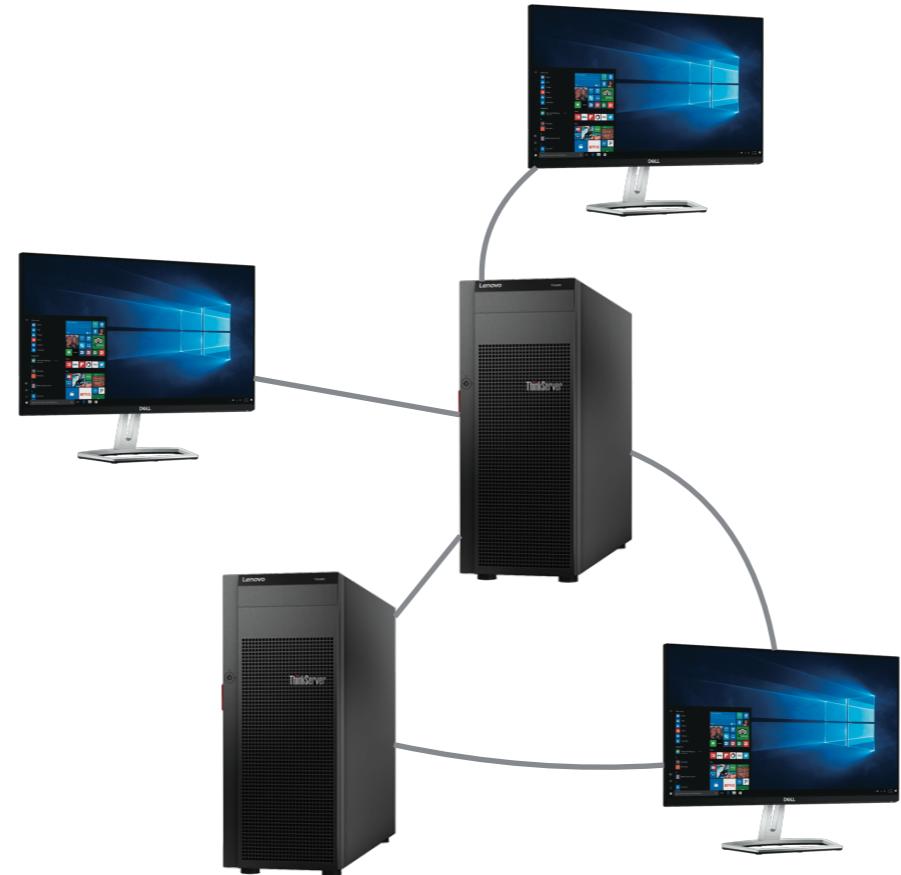
# Cloud computing

# Computing paradigms

“The network is the computer”  
(John Gage, Sun Microsystems,  
1984)



Mainframe computing  
1 computer / molti utenti

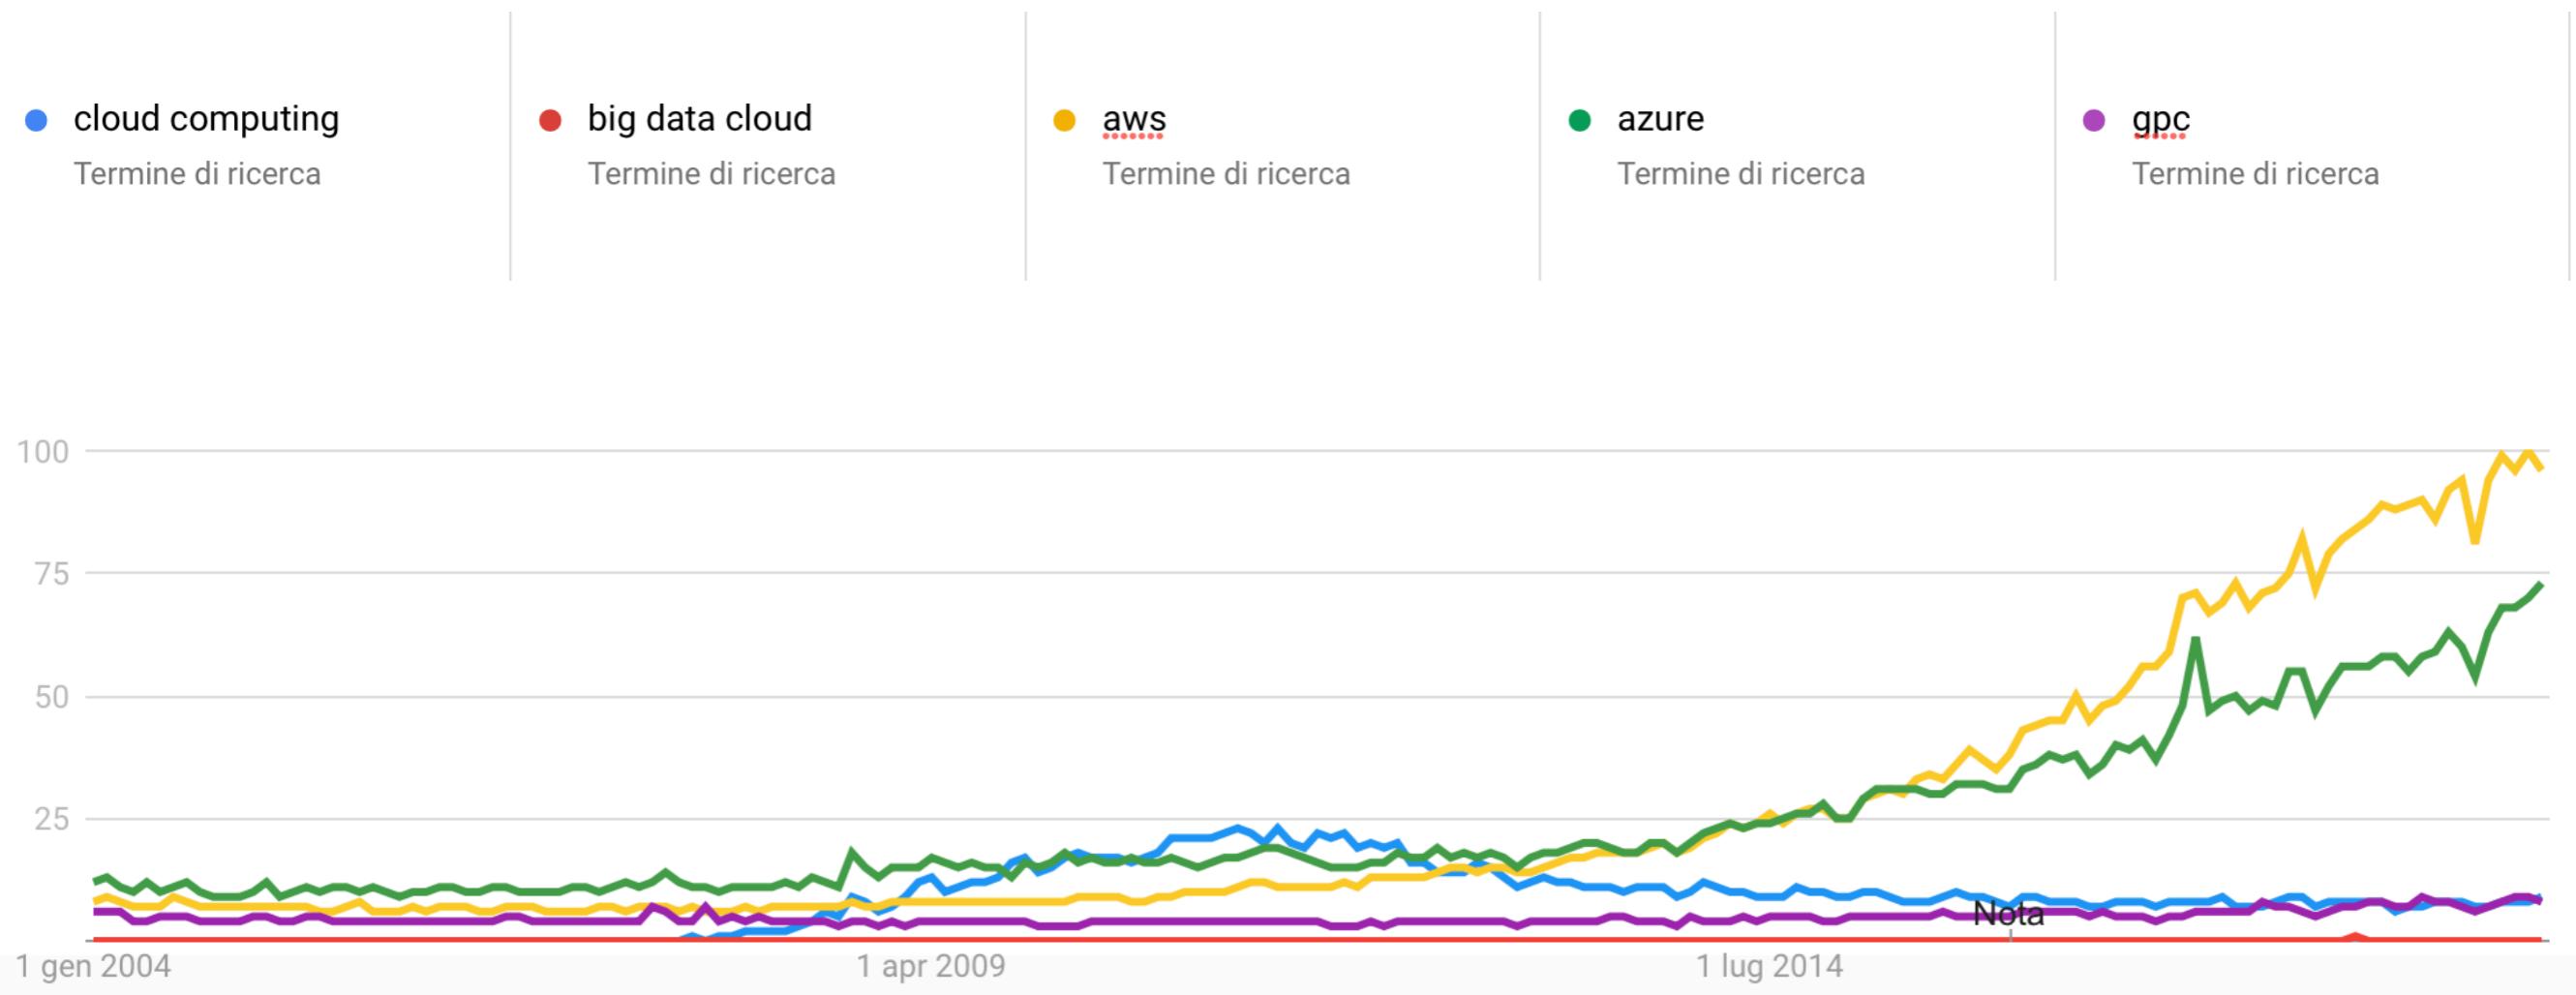


Client-server computing  
molti computer / molti utenti

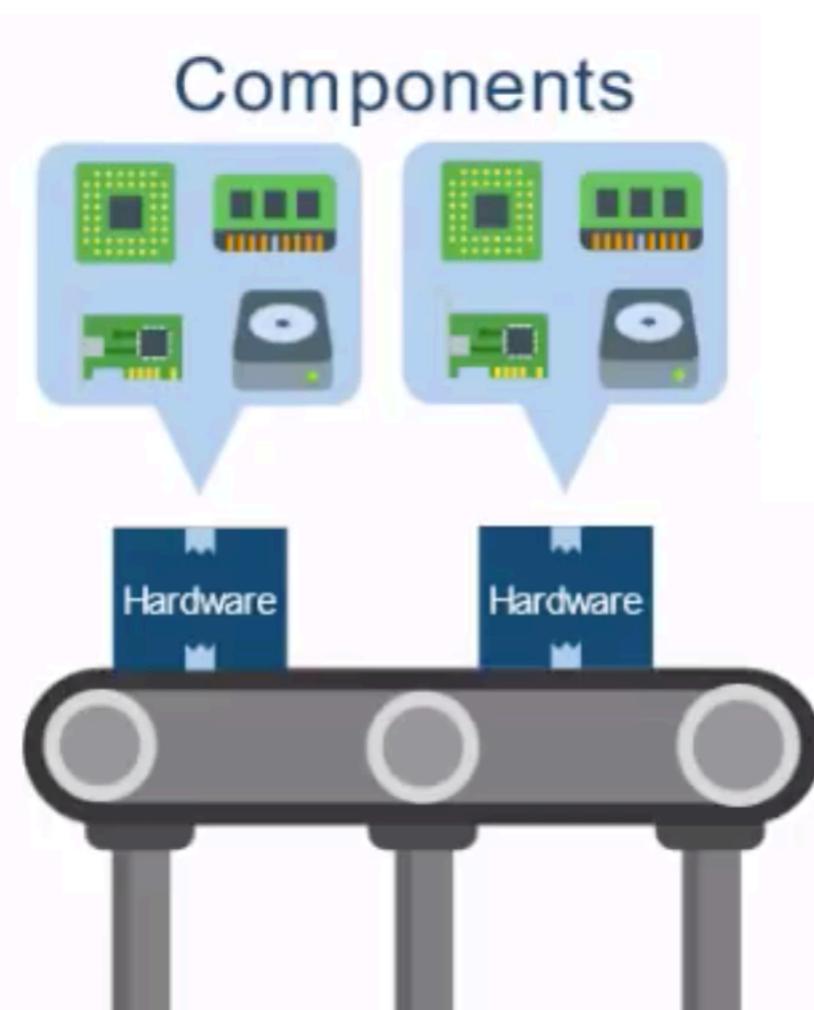
“The cloud is the computer”



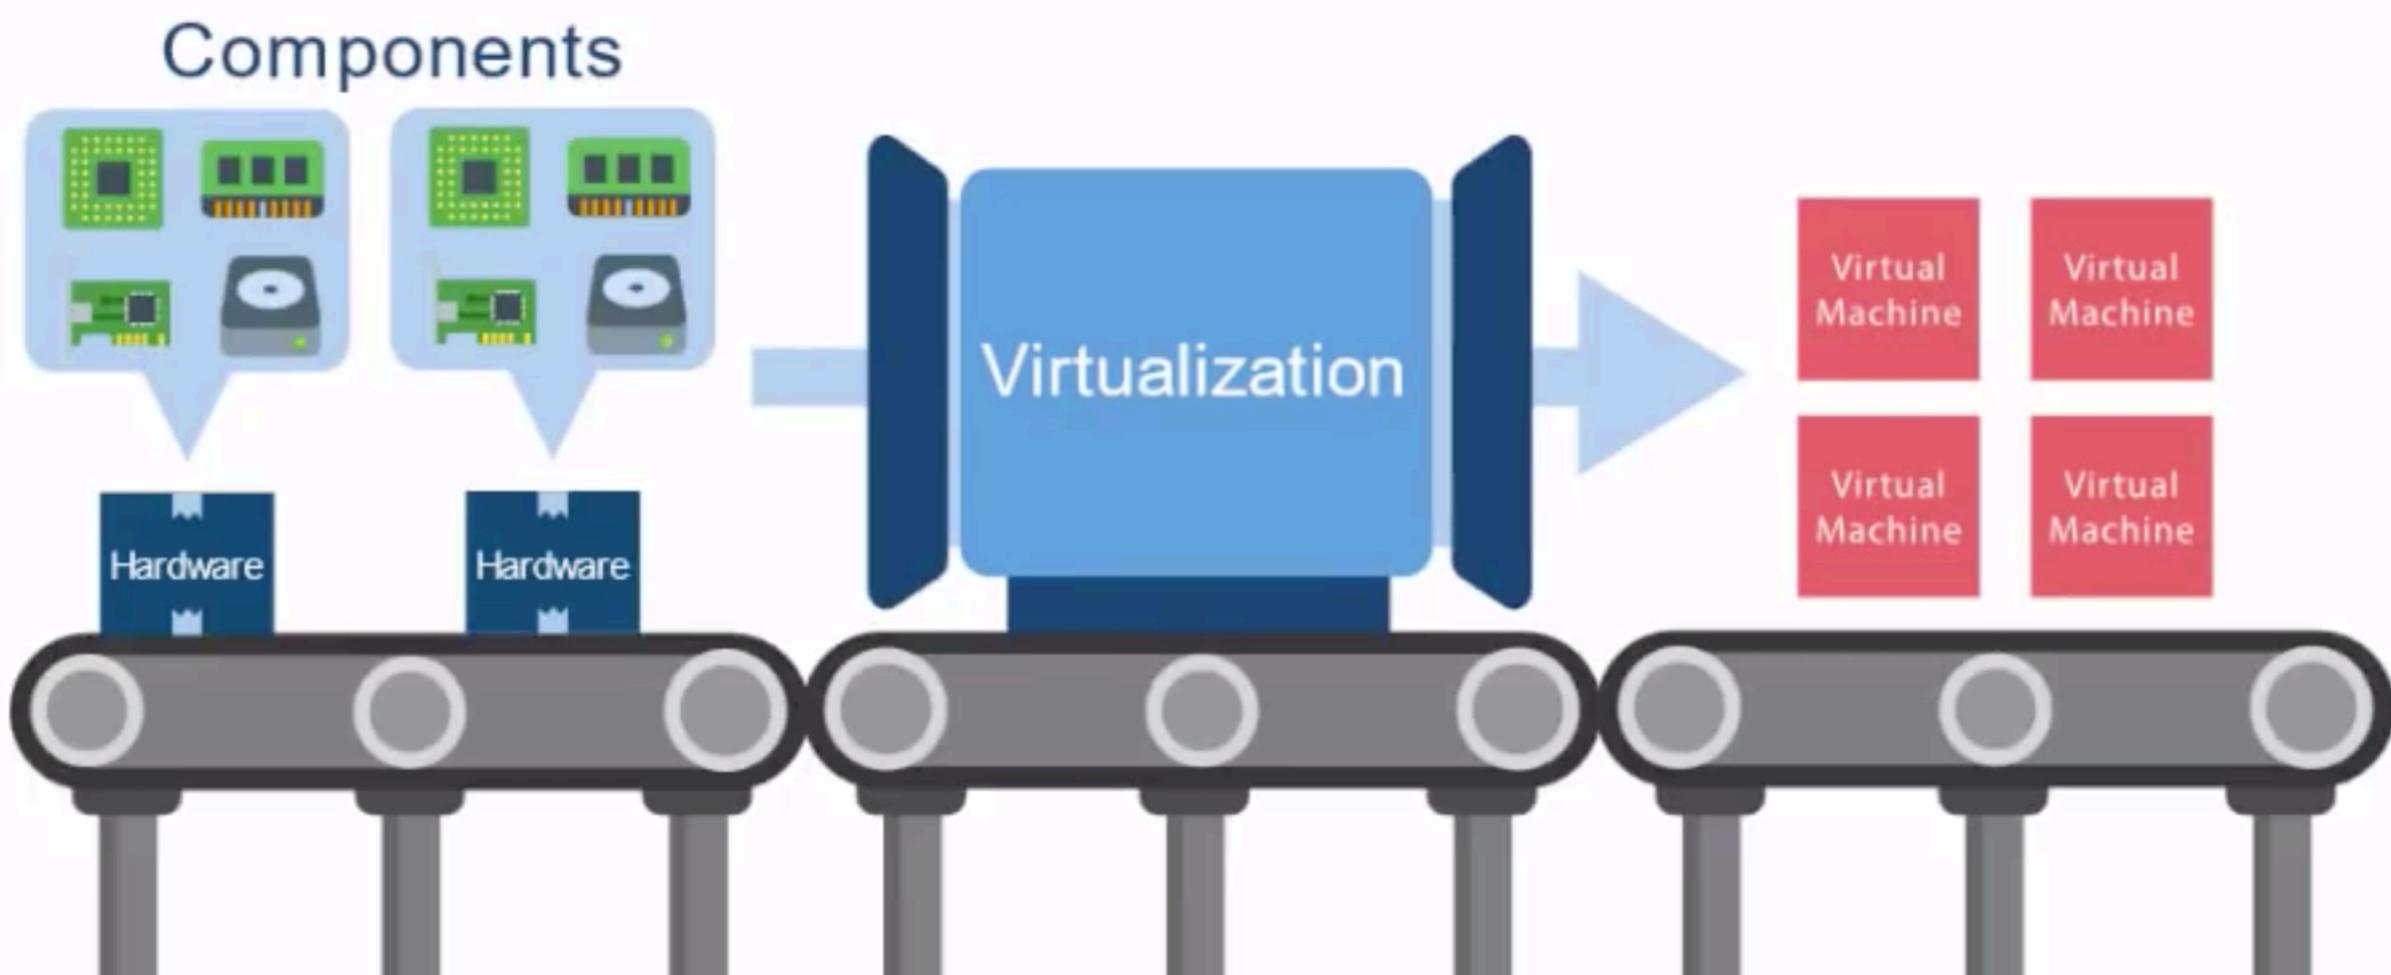
Cloud computing



Primo passo...  
la virtualizzazione....



Primo passo...  
la virtualizzazione....



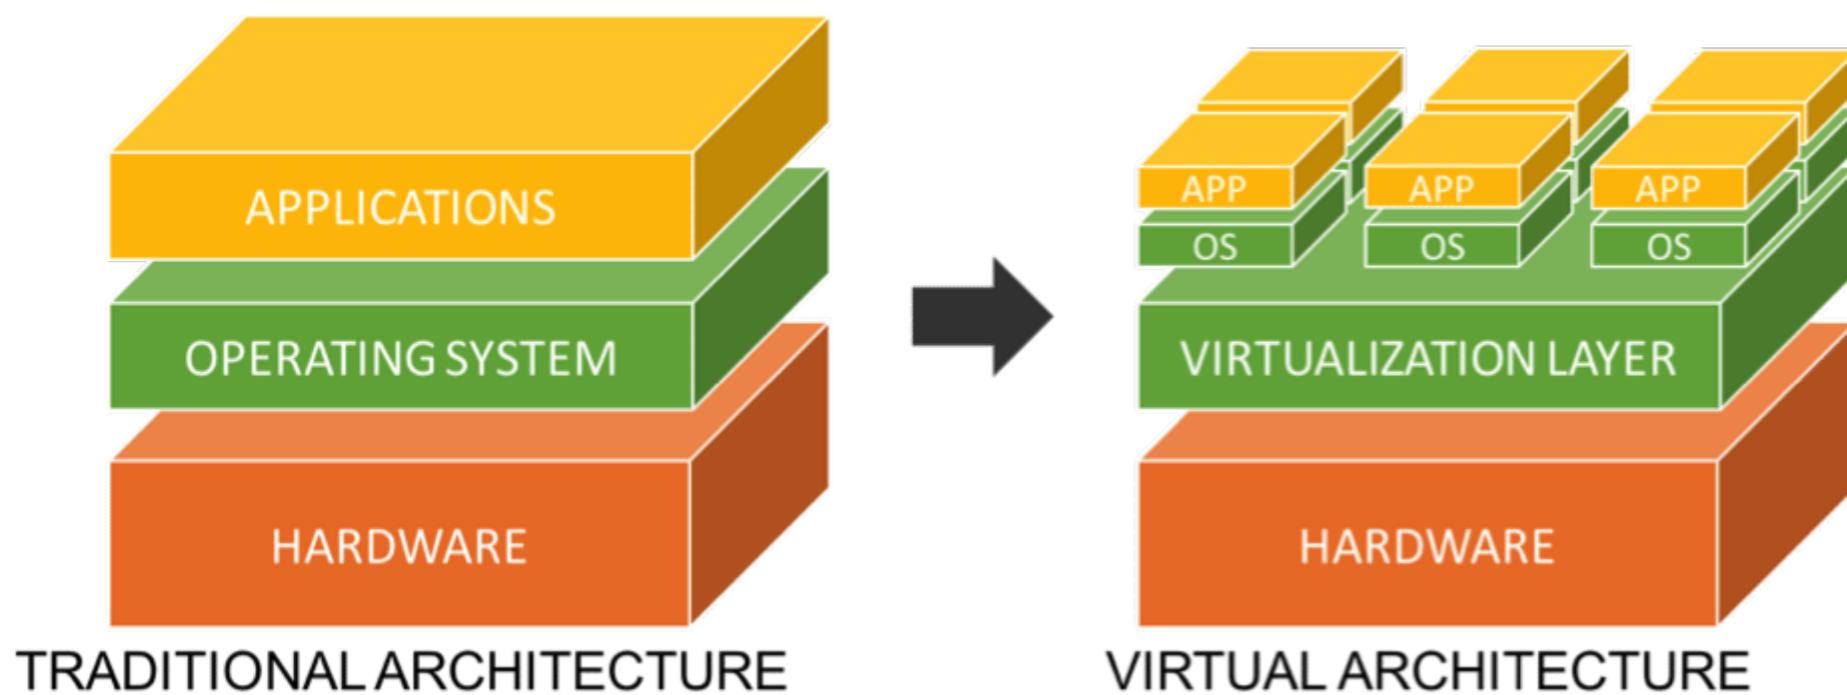
# Virtualizzazione

Attraverso la **virtualizzazione** vado a scomporre un grosso server fisico in tanti server virtuali.

Ciascun server virtuale è dedicato ad un componente del mio sistema.

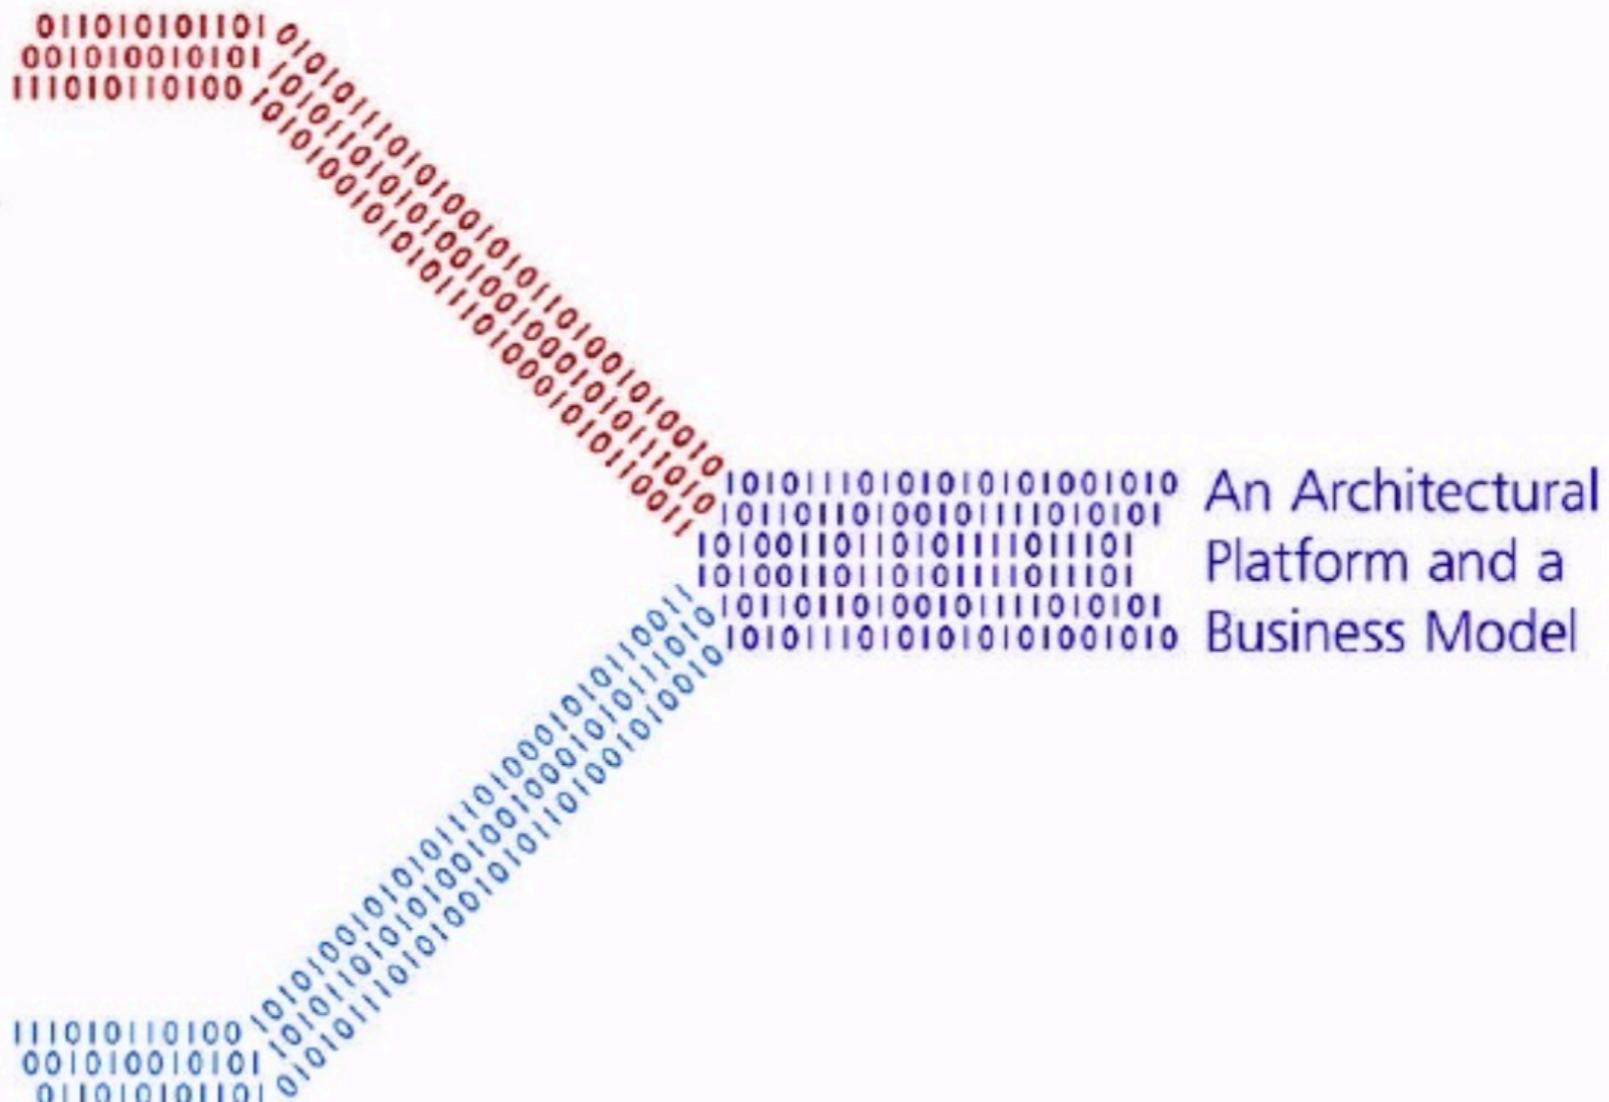
In altre parole, vado ad **astrarre un servizio IT**.

Poche grosse macchine fisiche ospitano tante piccole macchine virtuali indipendenti (diverso sistema operativo, diverse applicazioni, ....).



## A Business Model

- Pay as you go
- Virtualized resources
- On-demand
- Multi-tenancy
- Scalability, etc.



## A Delivery Platform

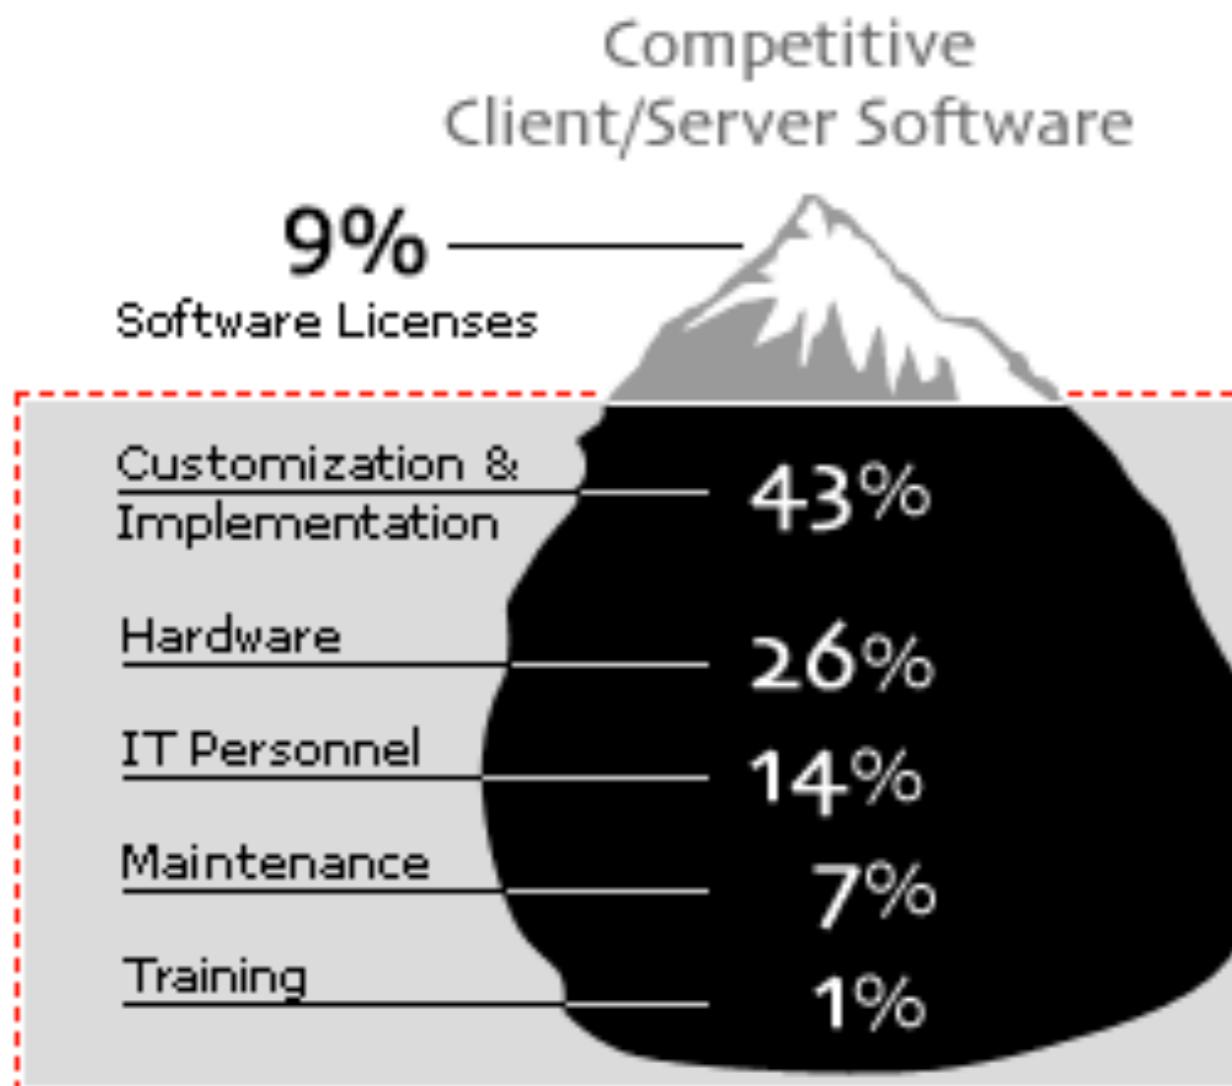
Enable users to access IP services and on-demand applications

A style of computing in which dynamically scalable  
and often virtualized resources are provided as a  
service over the Internet.

[Wikipedia]

Perché?

## Avoid the hidden costs of traditional CRM software



# Cloud computing

Il cloud computing fornisce tramite **WEB** risorse di elaborazione che spaziano dai server e lo storage fino ad intere applicazioni di livello enterprise (email, security, backup ecc.)

Il Cloud è in grado di effettuare provisioning  
di risorse **IT ‘as a service’**.

# 5 Caratteristiche fondamentali

- Servizi su richiesta
- Accesso di rete
- Pooling di risorse
- Elasticità rapida
- Misura dei servizi

# Benefici e sfide cloud computing

Pay per Use

Risorse infinite

Sicurezza  
(Compliance)

Service  
Assurance

Scalabilità

Servizi gestiti

Trasparenza

Integrazione  
con IT

Complessità  
gestita

Innovazione  
continua

Portabilità

Licenze

Riduce il time  
to market

Updated

Confronto dei  
costi

# Cloud computing

## Vantaggi

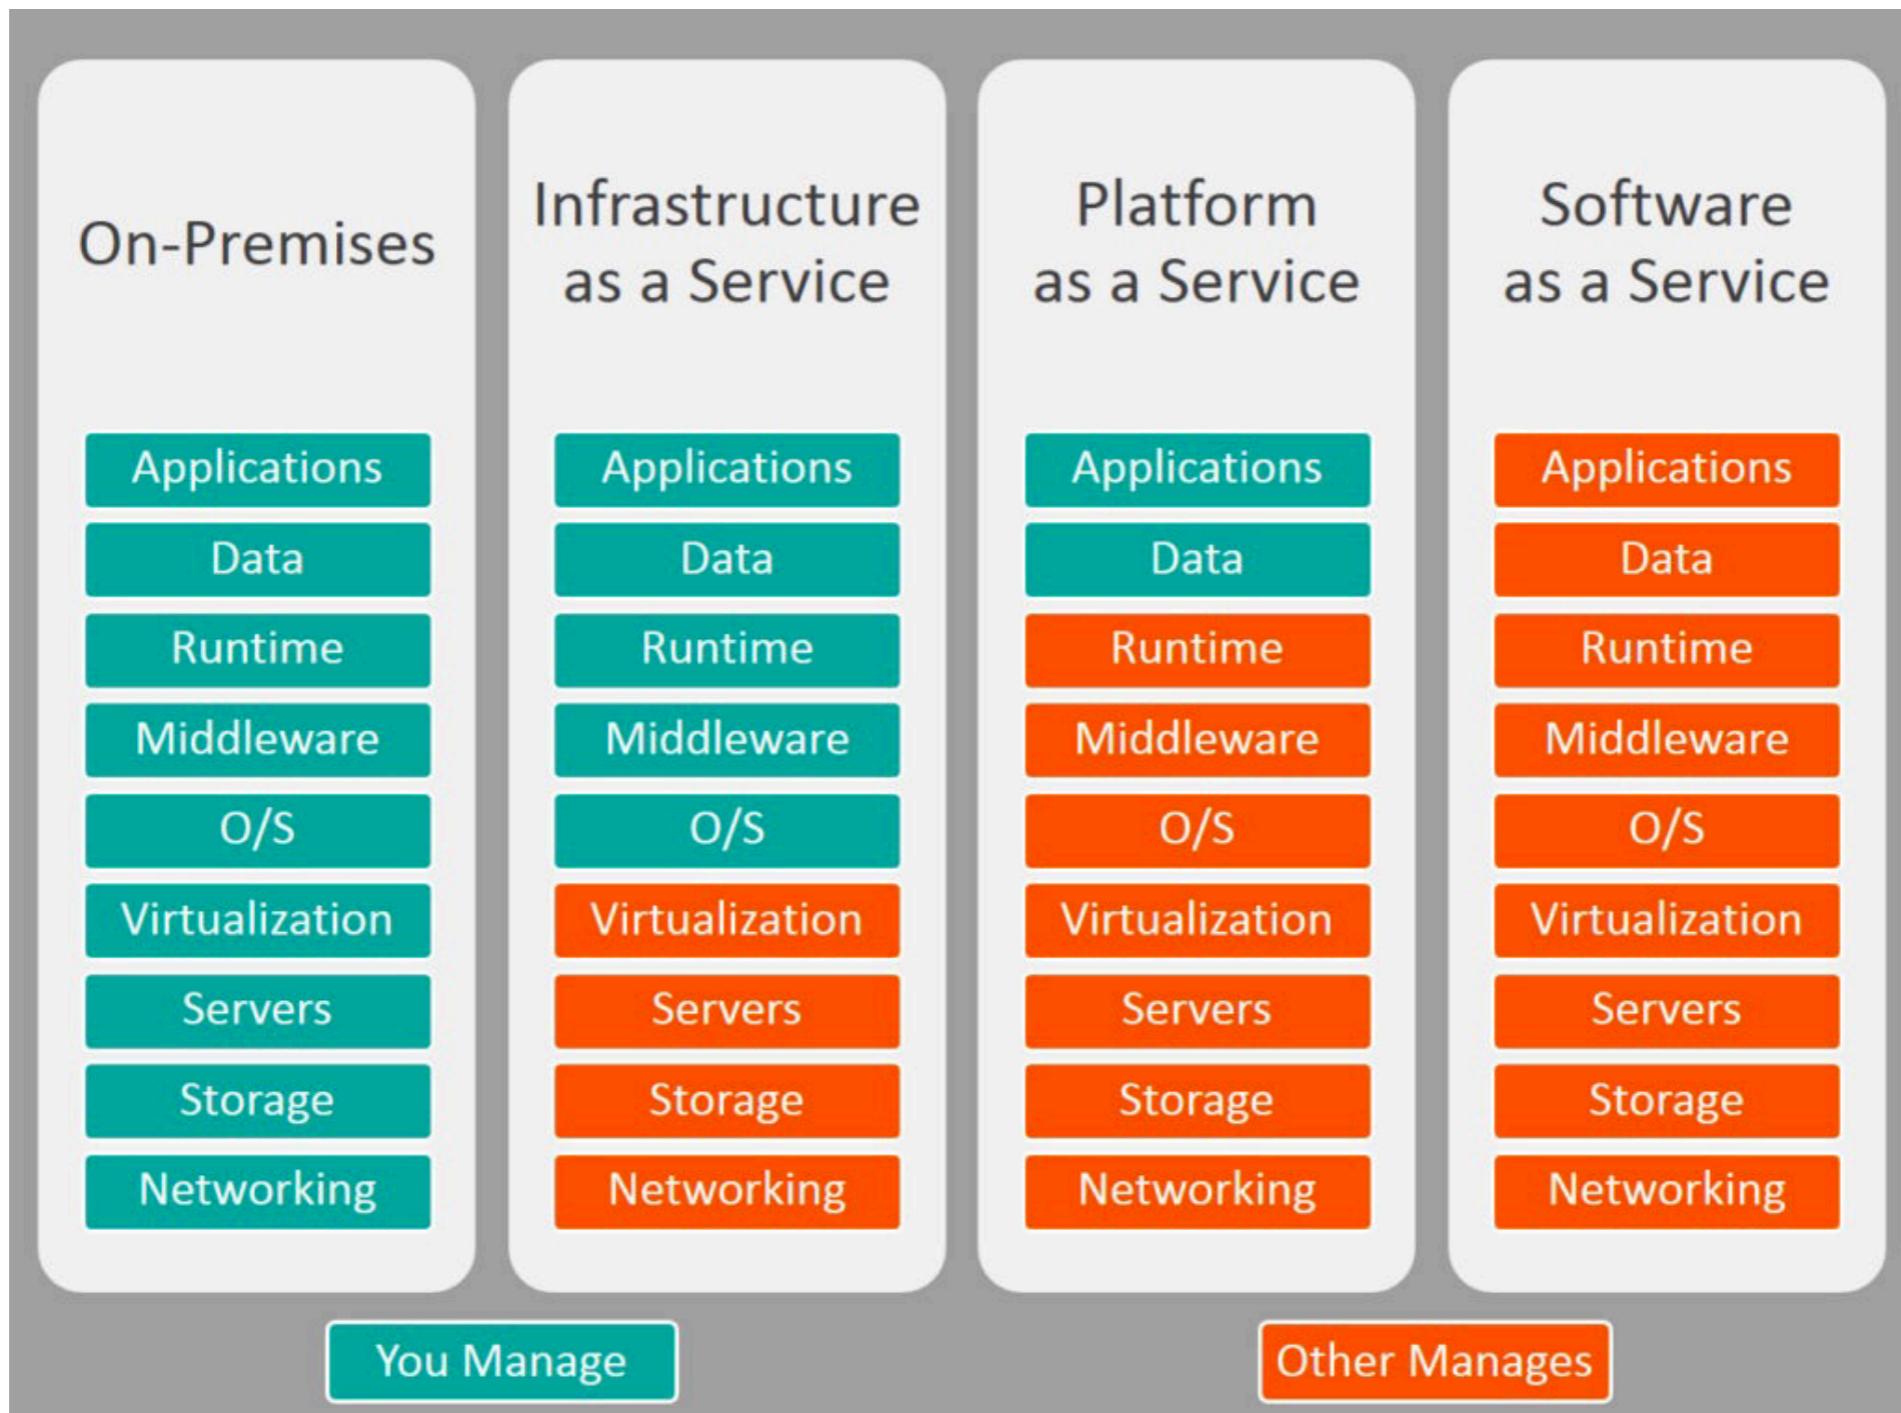
Increase  
Productivity  
of my business  
processes

Lower Cost  
Both CAPEX  
and OPEX

No Risk  
Keep all  
my assets safe

No Pain  
KISS principle  
(Keep it simple  
stupid)

# Modello di servizio



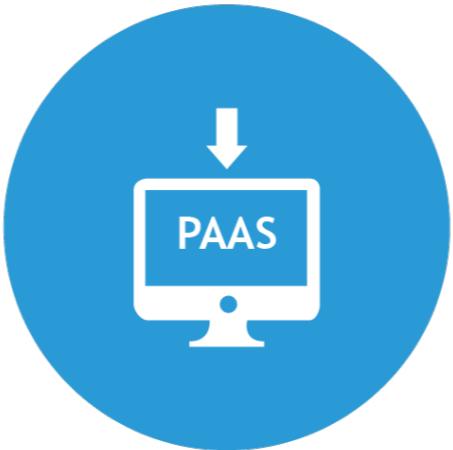
# Modelli di servizio



SaaS

Software as a Service

Il cliente utilizza applicazioni su una infrastruttura accessibile da vari dispositivi client attraverso un'interfaccia (API, interfaccia web, client dedicato)



PaaS

Platform as a Service

Permette di sviluppare e distribuire applicazioni create utilizzando linguaggi di programmazione supportati dal fornitore

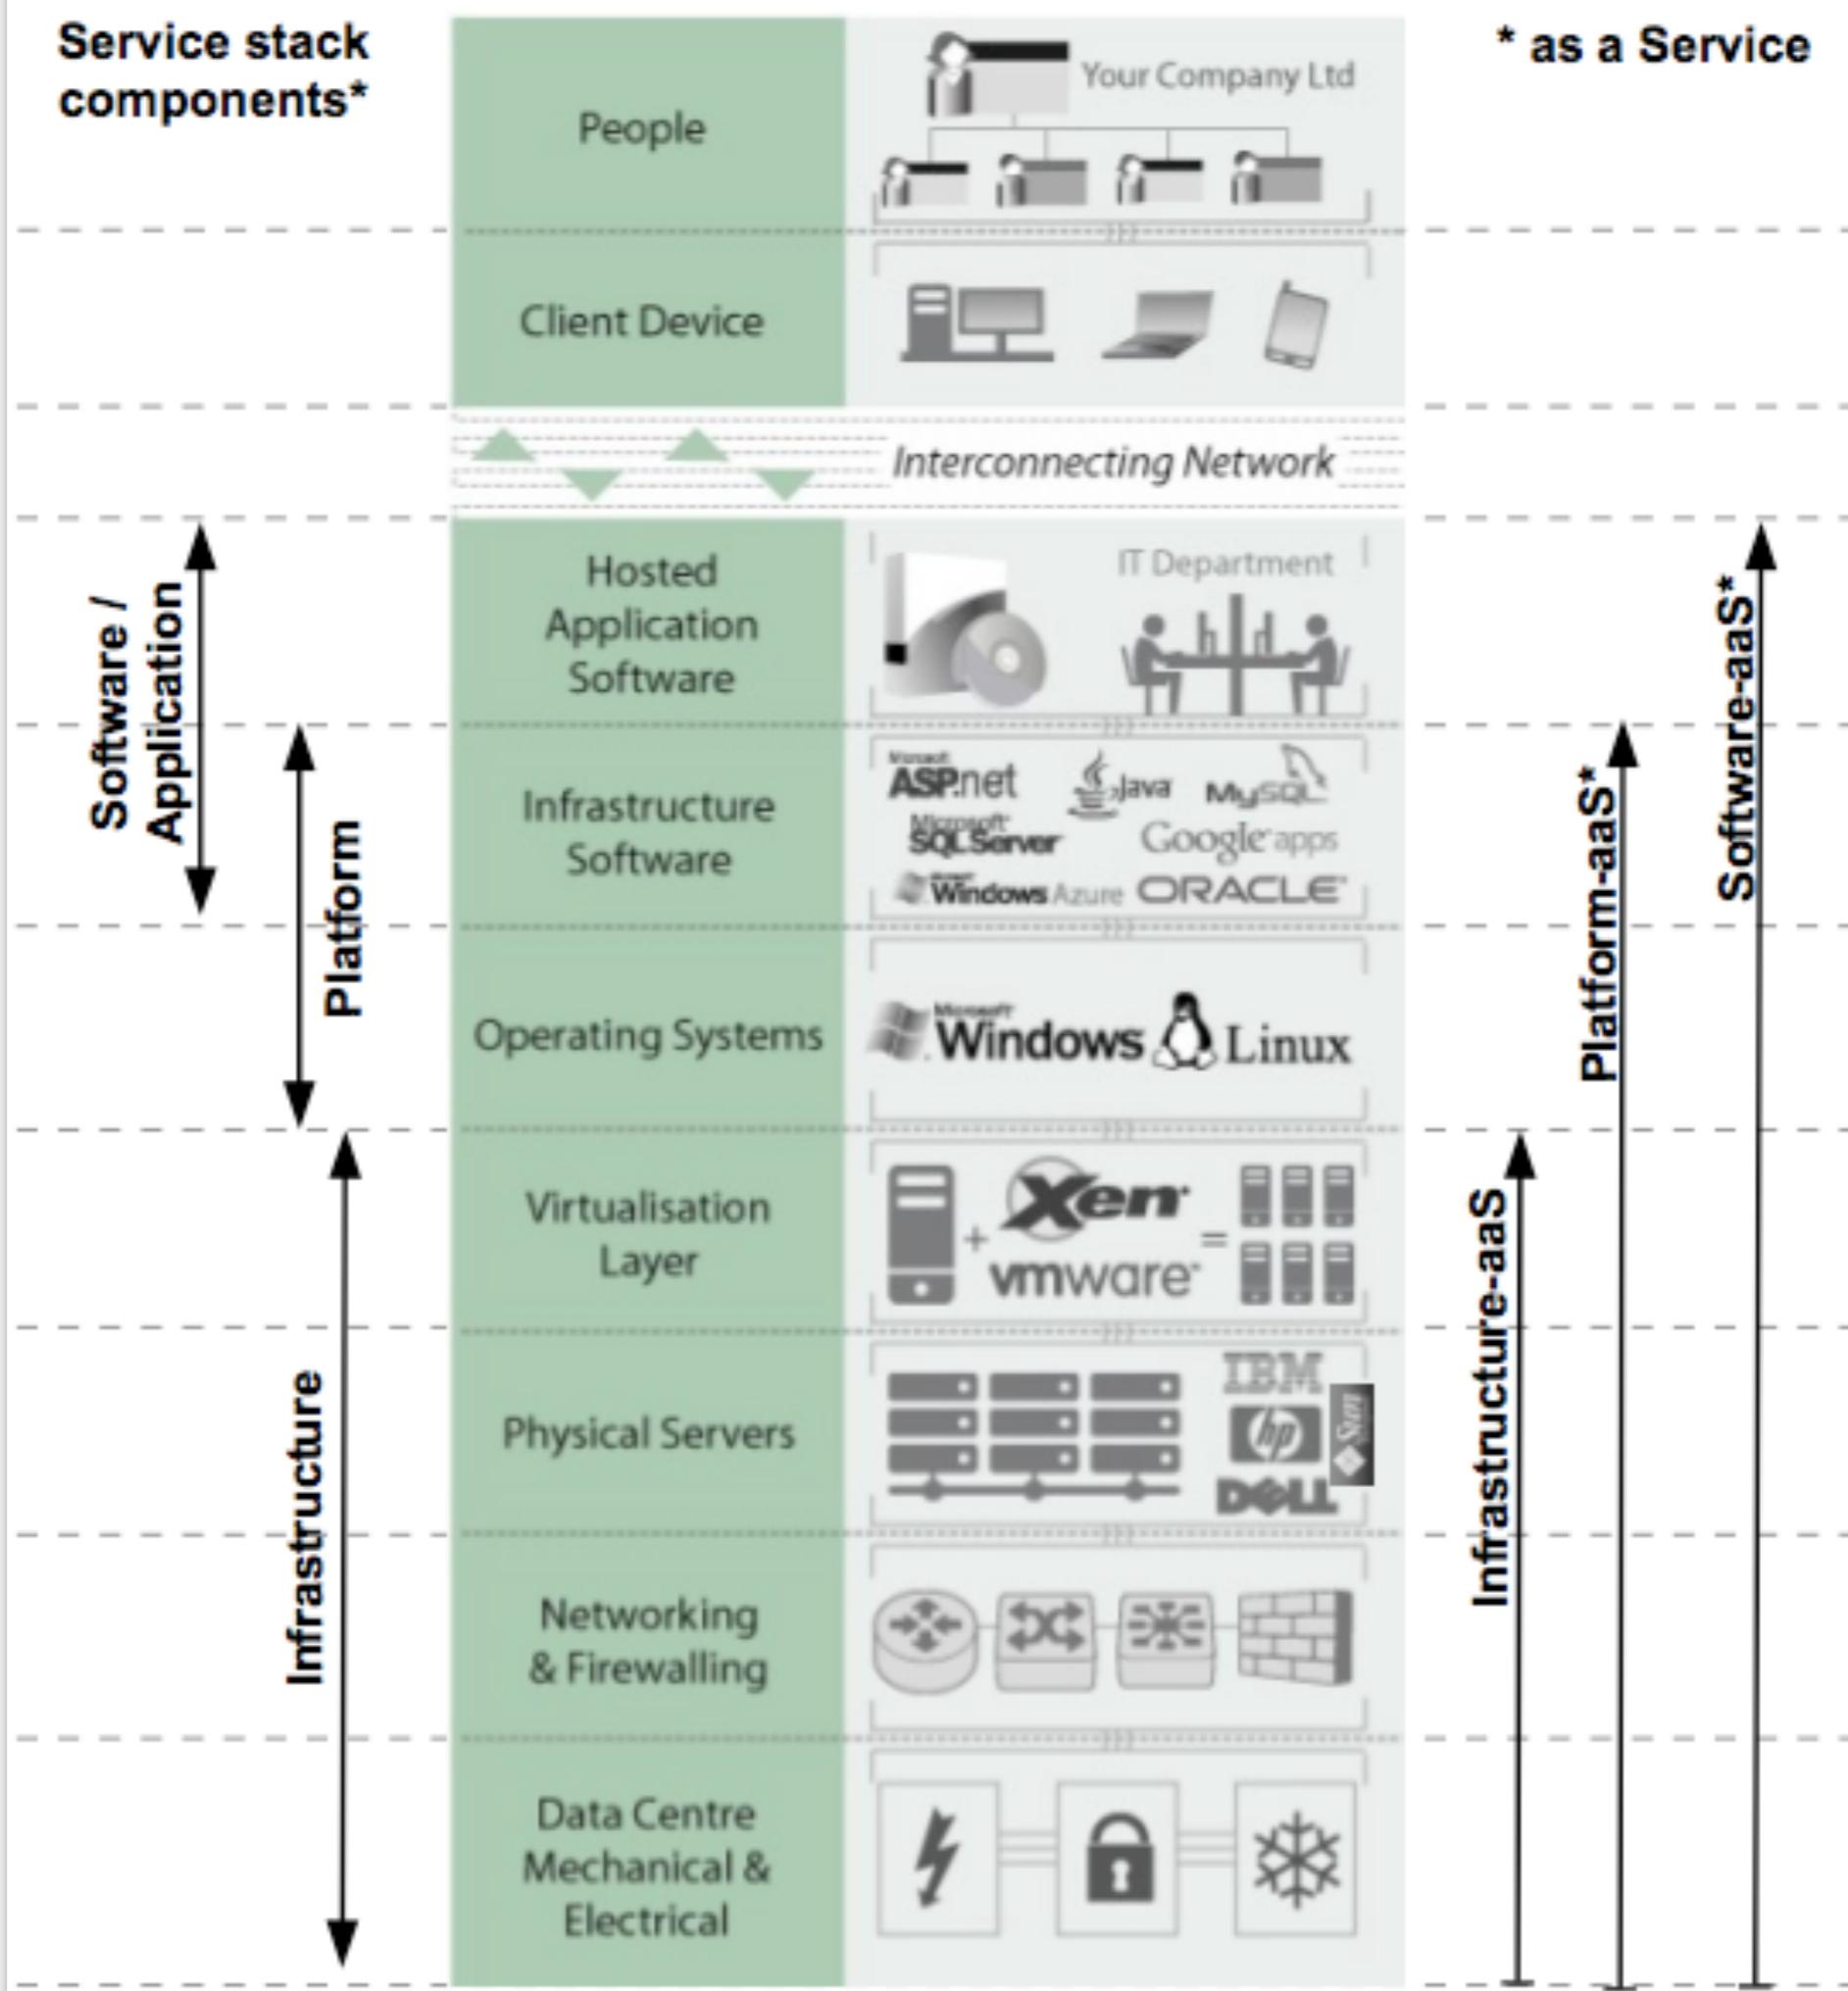


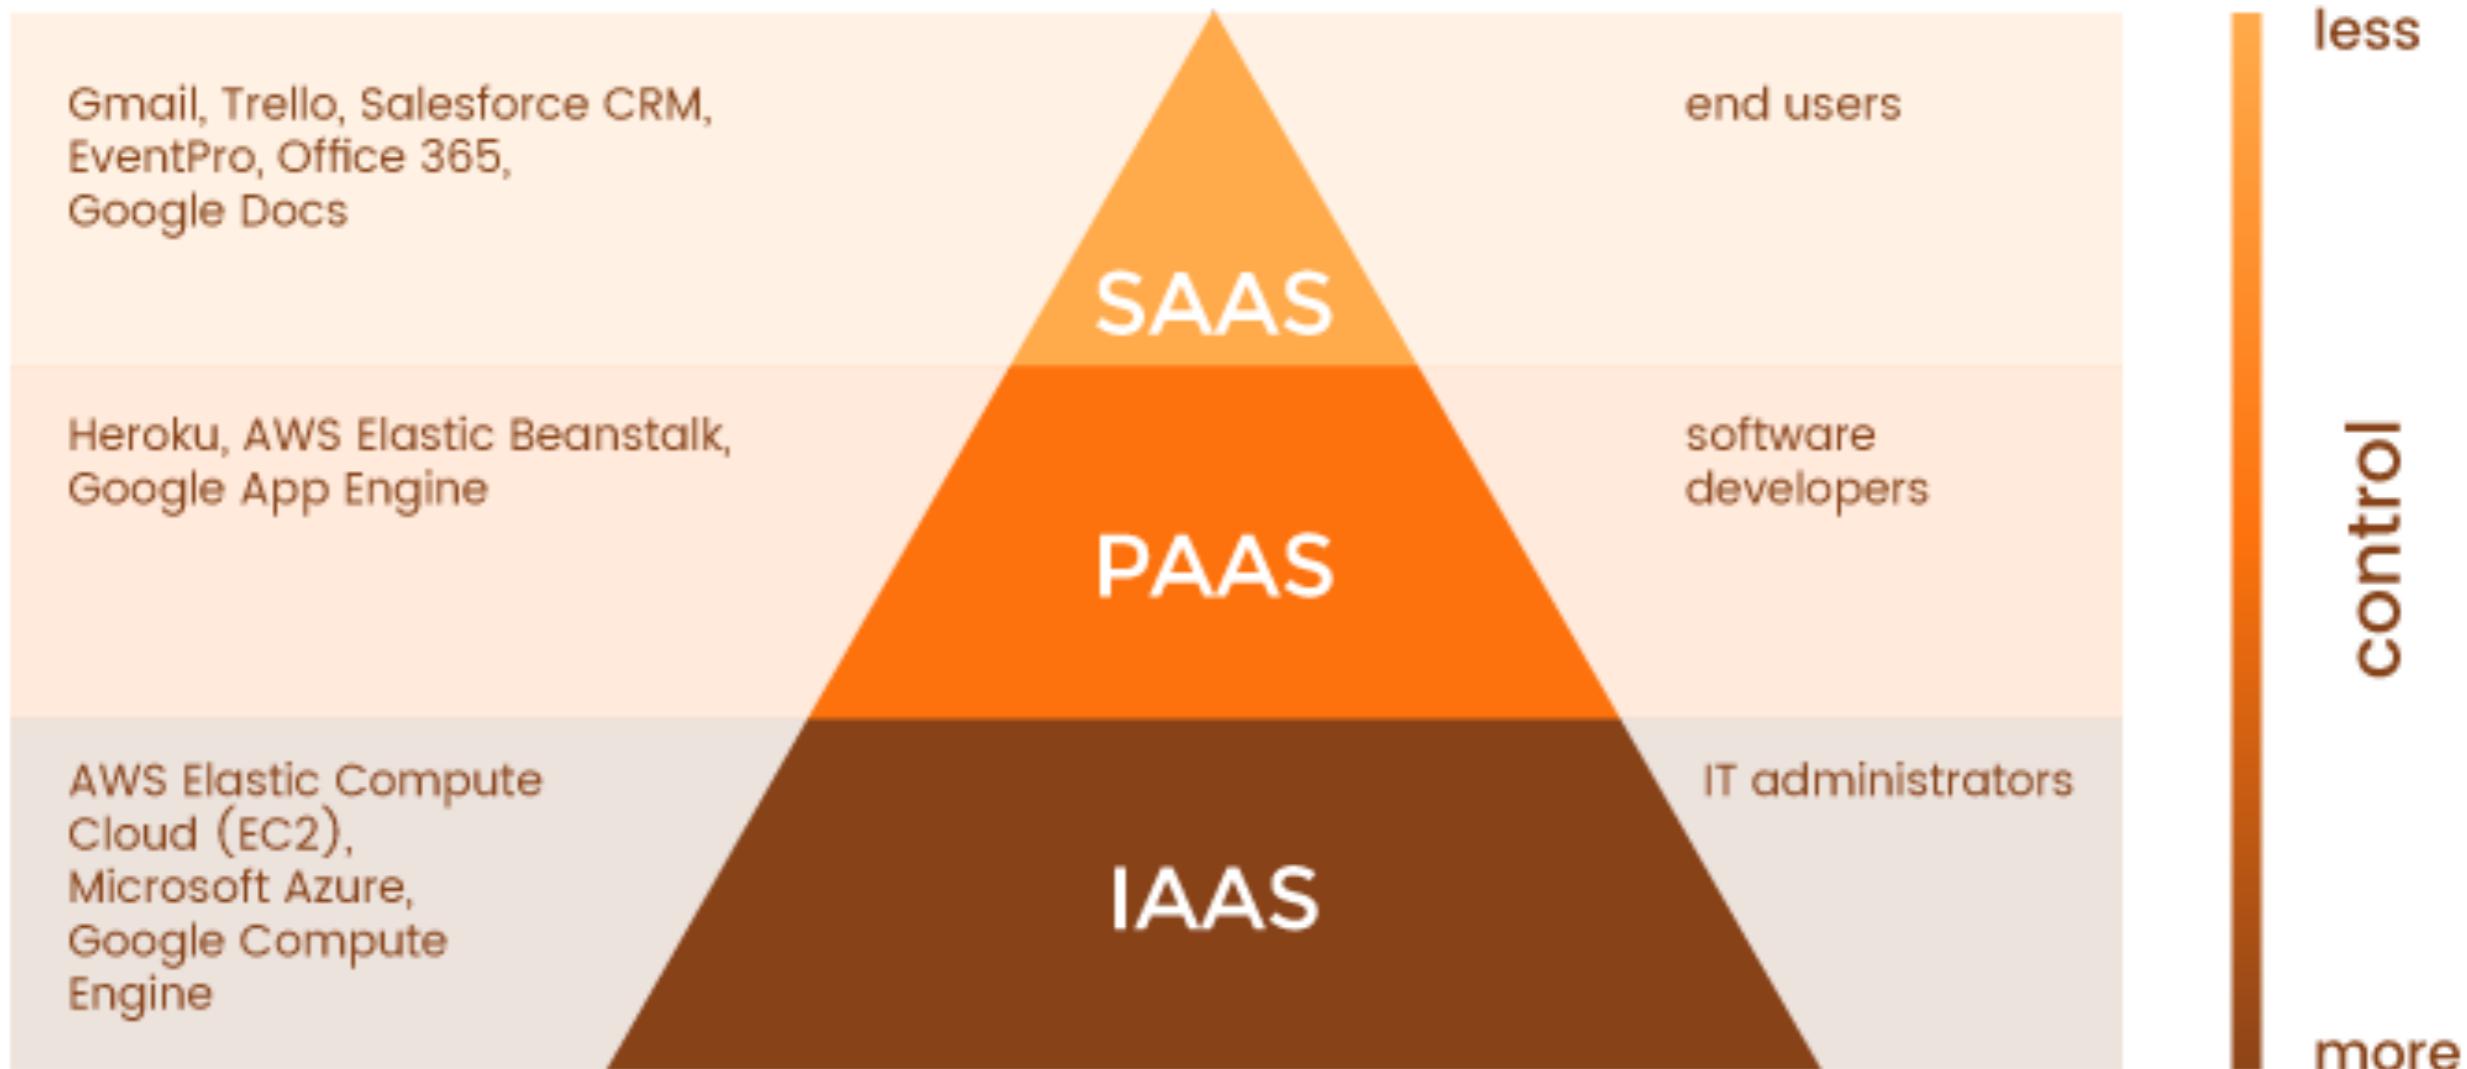
IaaS

Infrastructure as a Service

Noleggio capacità di CPU, storage, network e altre risorse come i sistemi operativi e le applicazioni

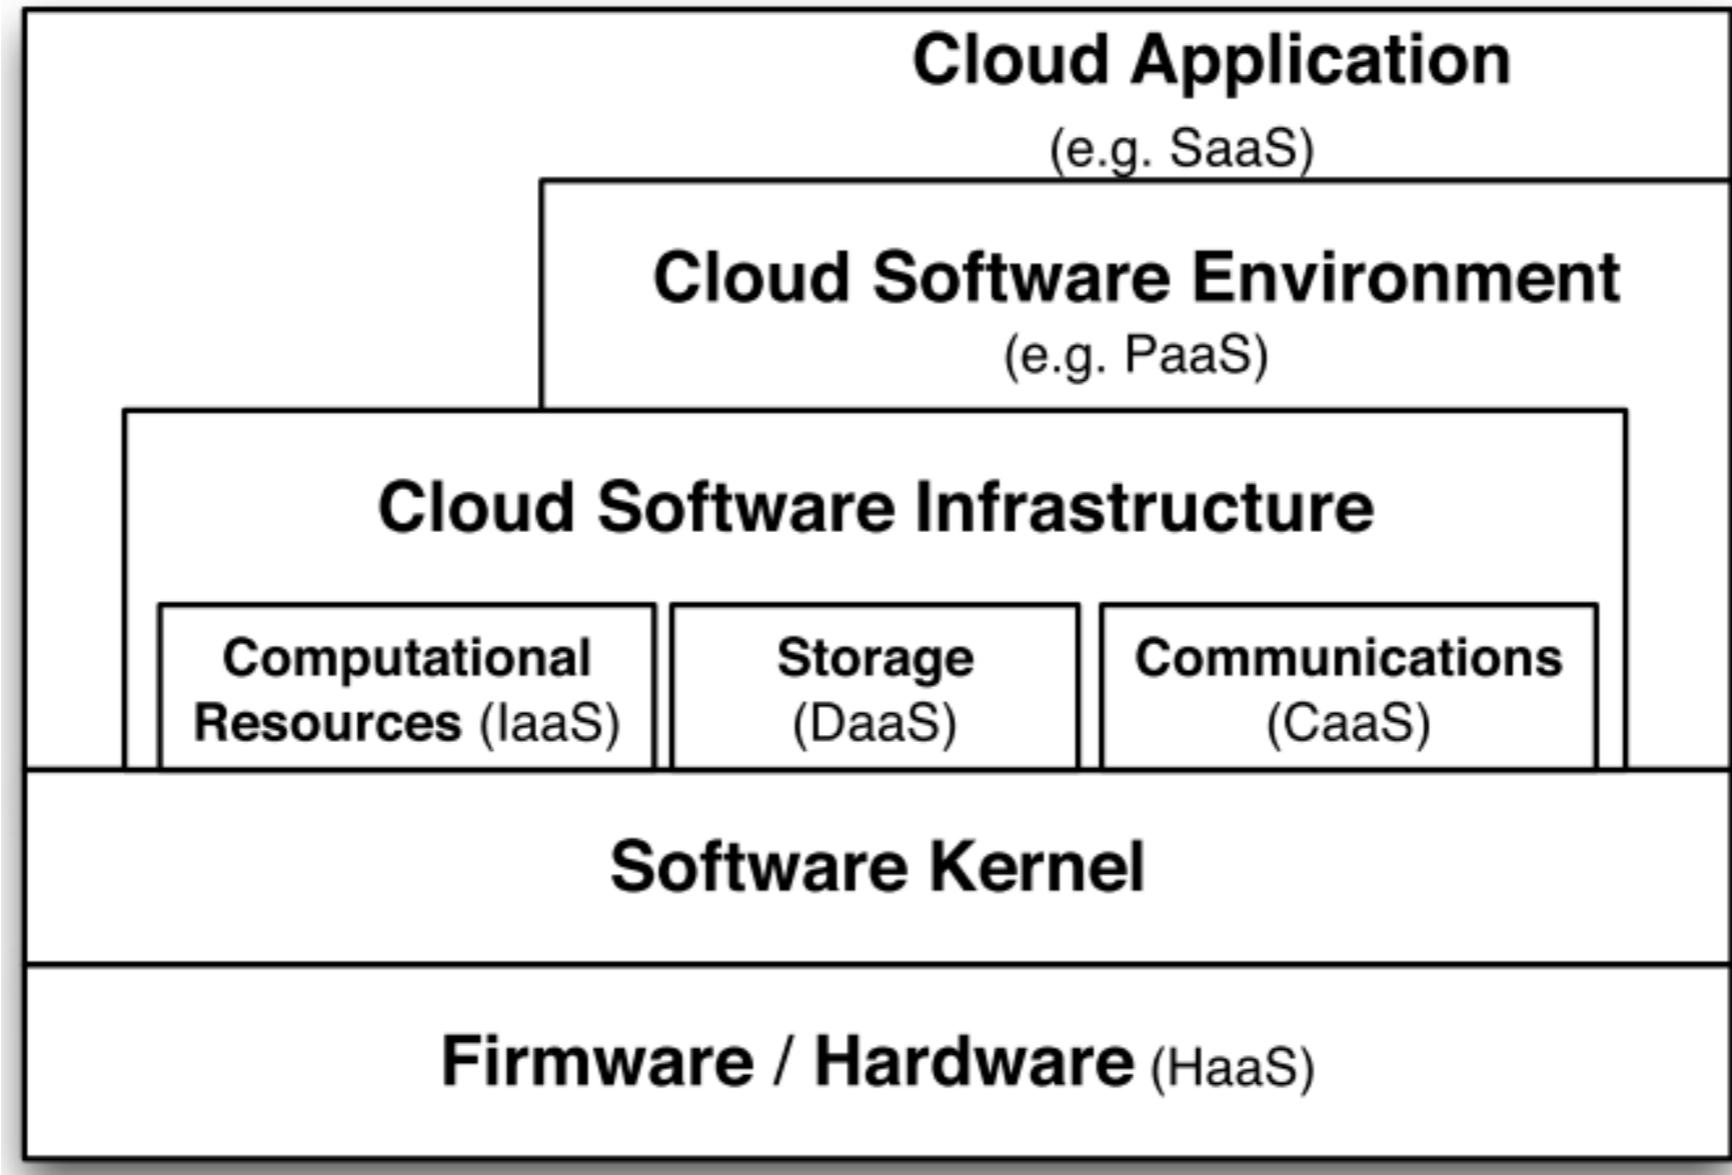
## Service stack components\*





AaaS	Architecture as a Service
BaaS	Business as a Service
CaaS	Computing as a Service
CRMaaS	CRM as a Service
DaaS	Data as a Service
DBaaS	Database as a Service
EaaS	Ethernet as a Service
FaaS	Frameworks as a Service
GaaS	Globalization or Governance as a Service
HaaS	Hardware as a Service
IaaS	Infrastructure or Integration as a Service
IDaaS	Identity as a Service
ITaaS	IT as a Service
LaaS	Lending as a Service
MaaS	Mashups as a Service
OaaS	Organization or Operations as a Service
SaaS	Software as a Service
StaaS	Storage as a Service
PaaS	Platform as a Service
TaaS	Technology or Testing as a Service
VaaS	Voice as a Service

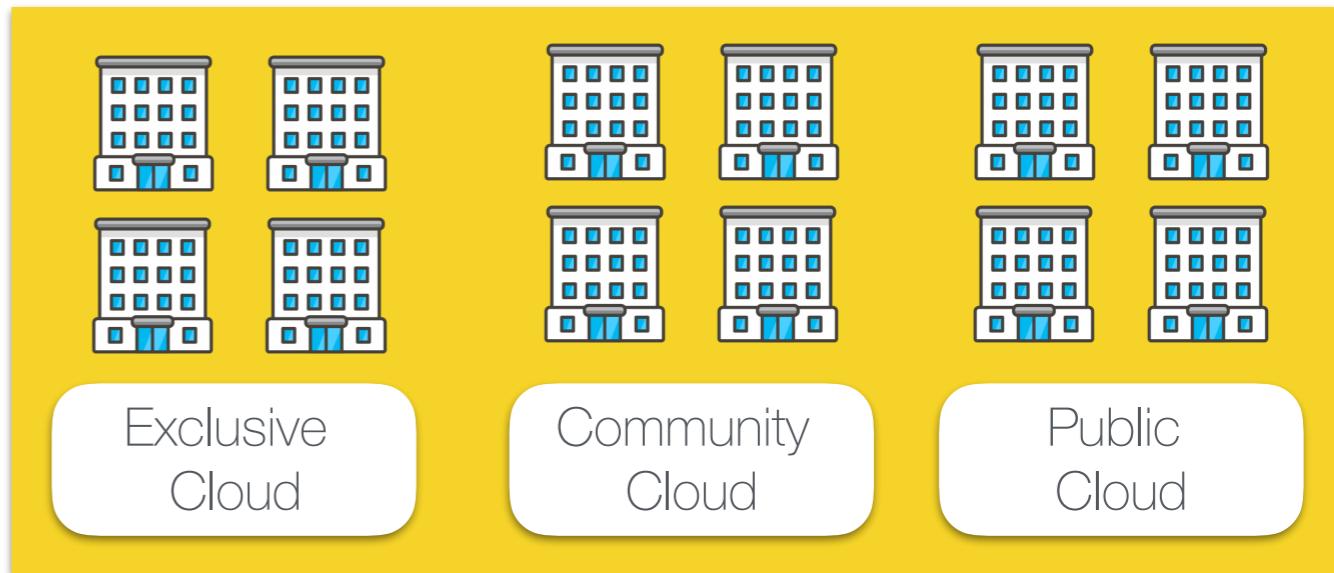
# Una vista organica



# Modelli di deployment



Privato Pubblico



# Private



A stylized icon of a four-story building. Each floor has four blue square windows. The bottom floor features a central entrance with a grey horizontal bar above it.

# Azienda A

# Custom Private Cloud

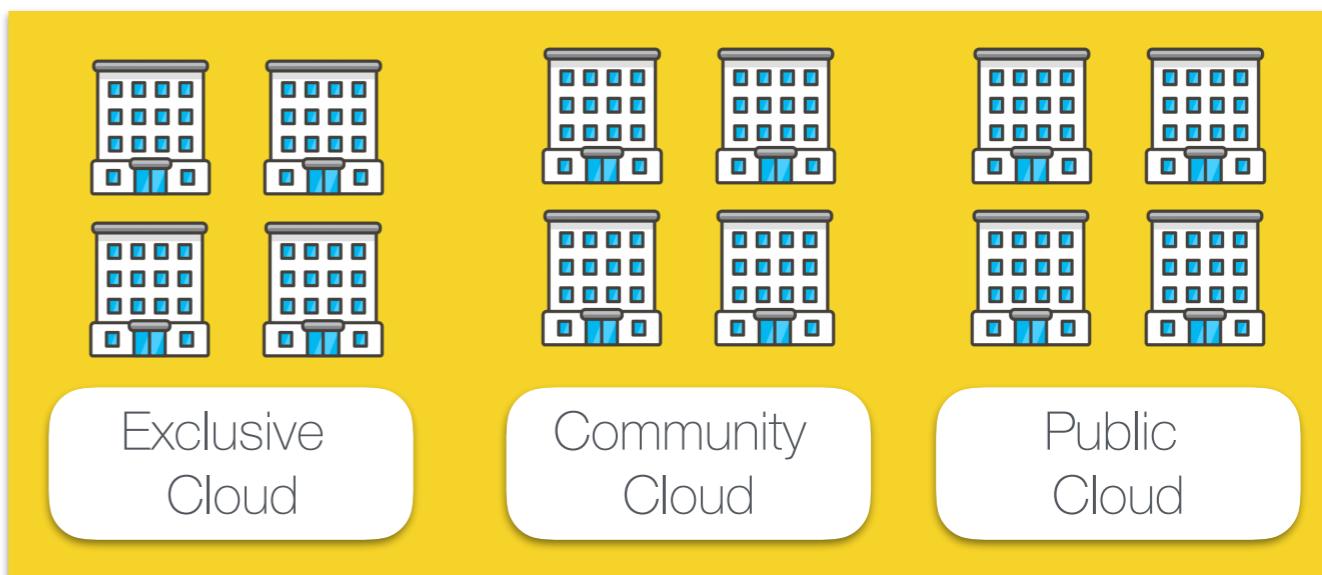


Azienda B

Packaged Private  
Cloud

# Packaged Private Cloud

Azienda B



# Exclusive Cloud

# Community Cloud

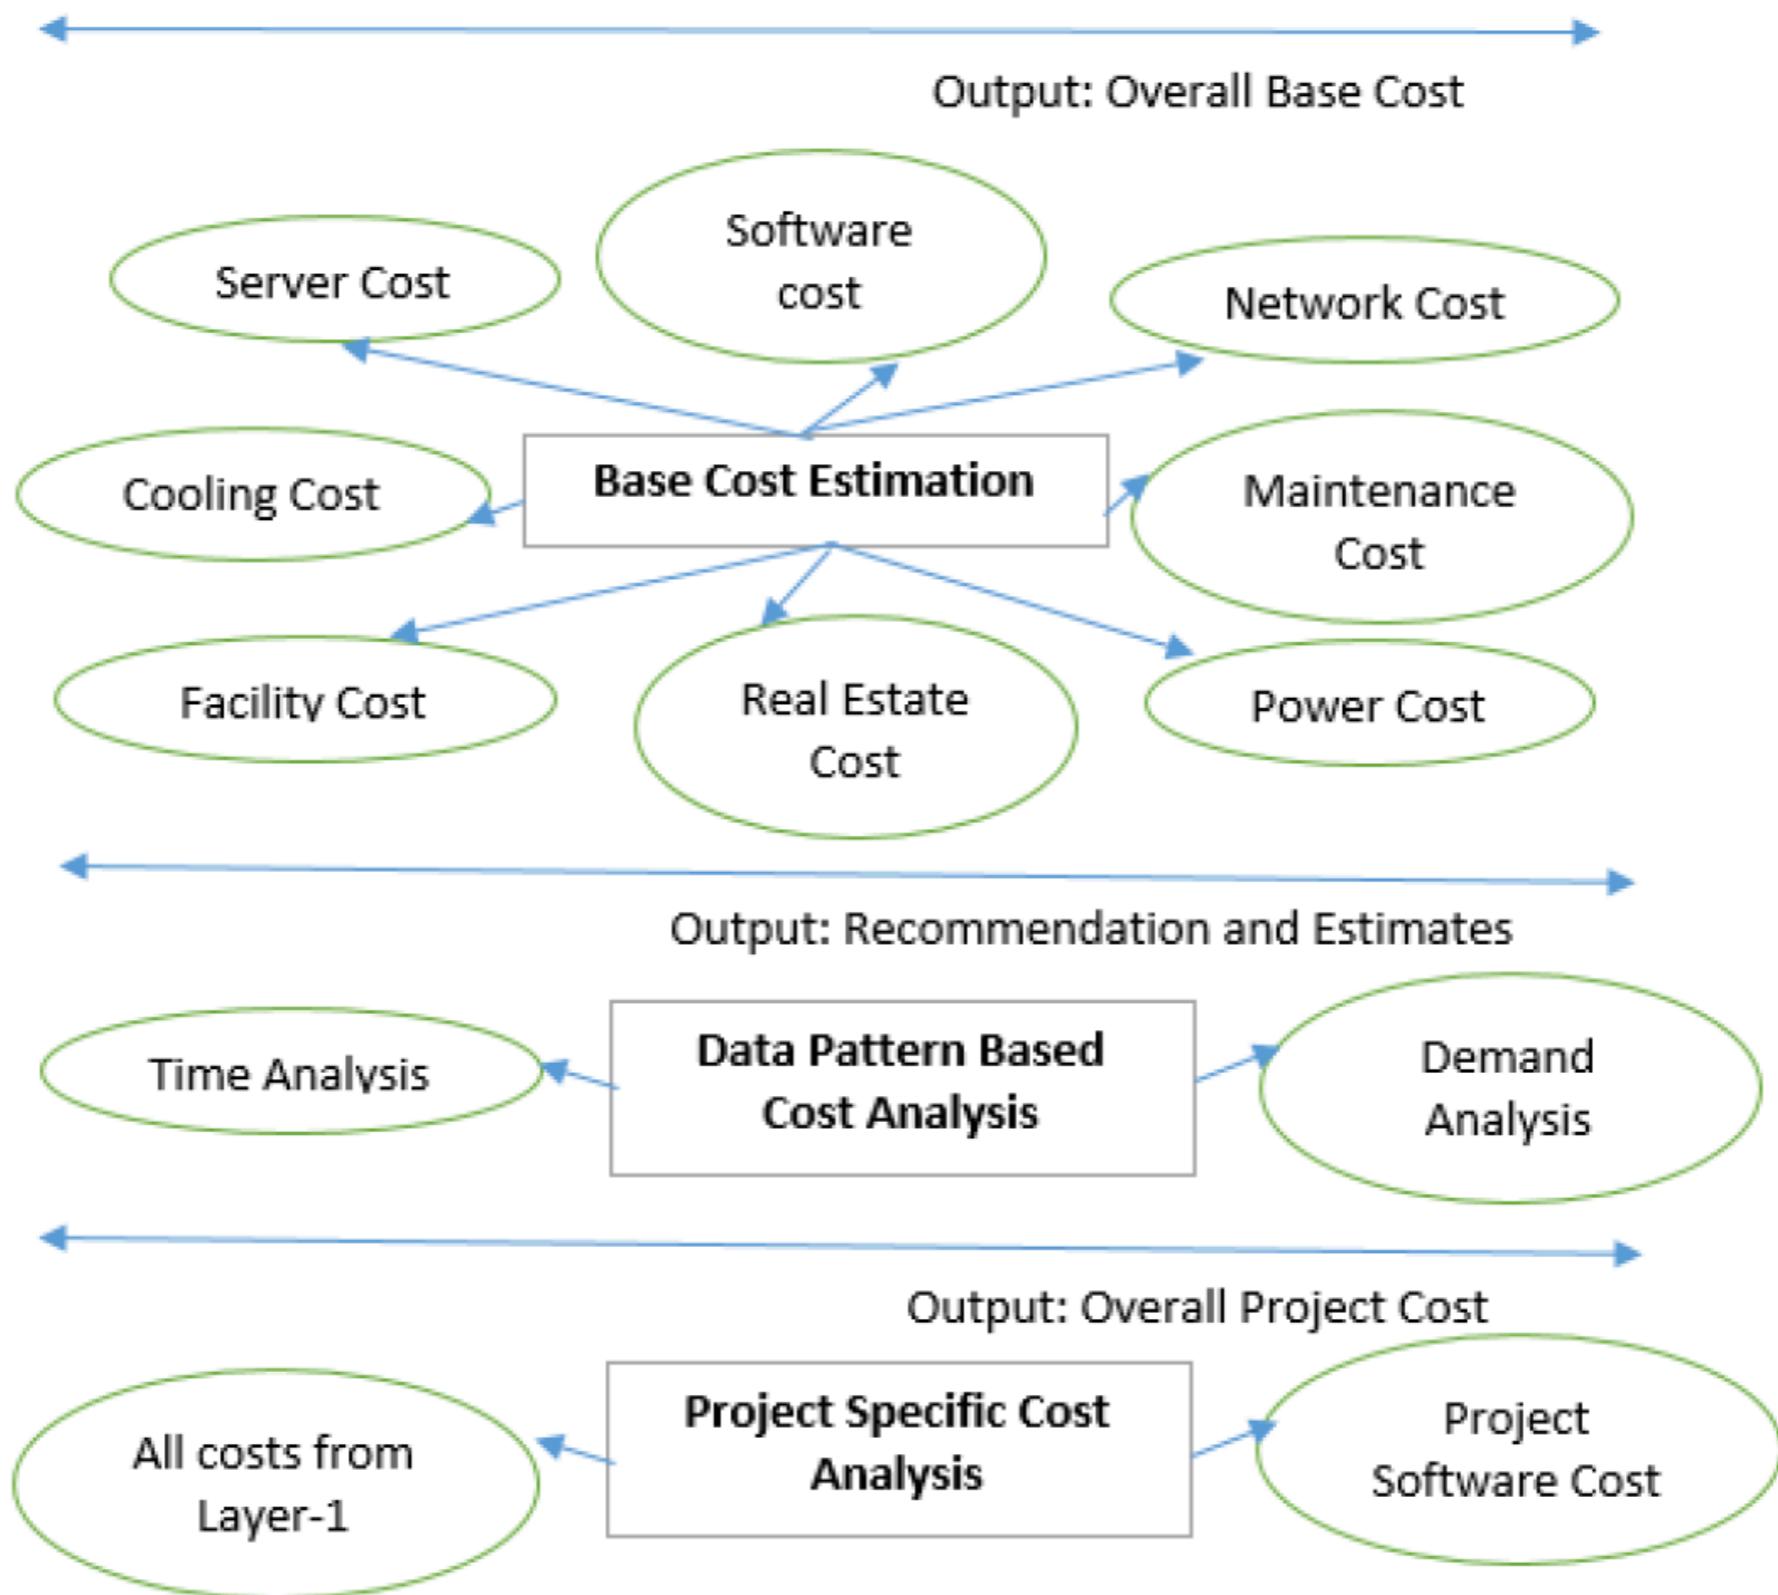
# Public Cloud

# Pubblico

# Flessibilità Sicurezza personalizzata

Costi ridotti  
Nessuna manutenzione  
Scalabilità quasi illimitata  
Affidabilità elevata  
Semplicità

Quali sono i costi da tenere in considerazione  
per un servizio in cloud?

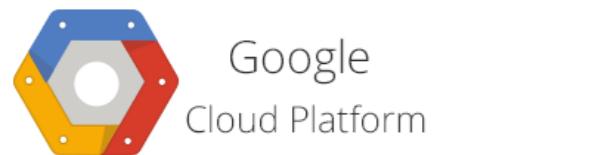


# Cloud pricing model

- Market-driven: variety of pricing models for each of IaaS, PaaS, SaaS service models
- Quanto costa 1 ora di computazione su un server “ready” per i Big Data?
- Quanto costa trasferire 1GB di dati?
- Quanto costa al mese 1GB di spazio occupato?



	vCPU	ECU	Memoria (GiB)	Storage dell'istanza (GB)	Utilizzo di Linux/UNIX
m5.large	2	8	8 GiB	Solo EBS	0,107 USD all'ora
m5.xlarge	4	16	16 GiB	Solo EBS	0,214 USD all'ora
m5.2xlarge	8	31	32 GiB	Solo EBS	0,428 USD all'ora
m5.4xlarge	16	60	64 GiB	Solo EBS	0,856 USD all'ora
m5.12xlarge	48	173	192 GiB	Solo EBS	2,568 USD all'ora
m5.24xlarge	96	345	384 GiB	Solo EBS	5,136 USD all'ora



Tipo di macchina	CPU virtuali	Memoria	Prezzo (USD)	Prezzo prilasciabile (USD)
n1-standard-1	1	3.75GB	\$0.0612	\$0.01230
n1-standard-2	2	7.5GB	\$0.1224	\$0.02460
n1-standard-4	4	15GB	\$0.2448	\$0.04920
n1-standard-8	8	30GB	\$0.4896	\$0.09840
n1-standard-16	16	60GB	\$0.9792	\$0.19680
n1-standard-32	32	120GB	\$1.9584	\$0.39360
n1-standard-64	64	240GB	\$3.9168	\$0.78720

# Modello di costo (semplice)

	Ore di utilizzo al giorno	Tipo di macchina	Totale (all'anno)
Database Server (NoSql)	24	m5.2xlarge	3500
ETL (Spark)	10	m5.4xlarge	2900
Presentation	10	m5.4xlarge	2900
			9300

Solo costi EC2!!!

Per risparmiare???

Per risparmiare???

Sfrutto i vantaggi del cloud....

# Istanze SPOT

Le istanze **Spot** ti permettono di sfruttare le capacità inutilizzate all'interno del cloud. Le istanze Spot sono disponibili con prezzi **scontati fino al 90%** inferiori rispetto ai prezzi delle istanze on demand.

Quando si utilizzano le Istanze Spot, è necessario essere **preparati alle interruzioni**. Il fornitore può interrompere le Istanze Spot quando il prezzo Spot supera il prezzo massimo, se la domanda di Istanze Spot aumenta o se l'offerta di Istanze Spot diminuisce.



	vCPU	ECU	Memoria (GiB)	Storage dell'istanza (GB)	Utilizzo di Linux/UNIX	Spot
m5.large	2	8	8 GiB	Solo EBS	0,107 USD all'ora	\$0.0365 all'ora
m5.xlarge	4	16	16 GiB	Solo EBS	0,214 USD all'ora	\$0.0721 all'ora
m5.2xlarge	8	31	32 GiB	Solo EBS	0,428 USD all'ora	\$0.1444 all'ora
m5.4xlarge	16	60	64 GiB	Solo EBS	0,856 USD all'ora	\$0.2916 all'ora
m5.12xlarge	48	173	192 GiB	Solo EBS	2,568 USD all'ora	\$0.8657 all'ora
m5.24xlarge	96	345	384 GiB	Solo EBS	5,136 USD all'ora	\$1.7314 all'ora

# Istanze riservate

Le **istanze riservate** ti danno la possibilità di effettuare un pagamento unico per ogni istanza che desideri prenotare e ricevere uno sconto significativo sulla tariffa di utilizzo oraria per quell'istanza.

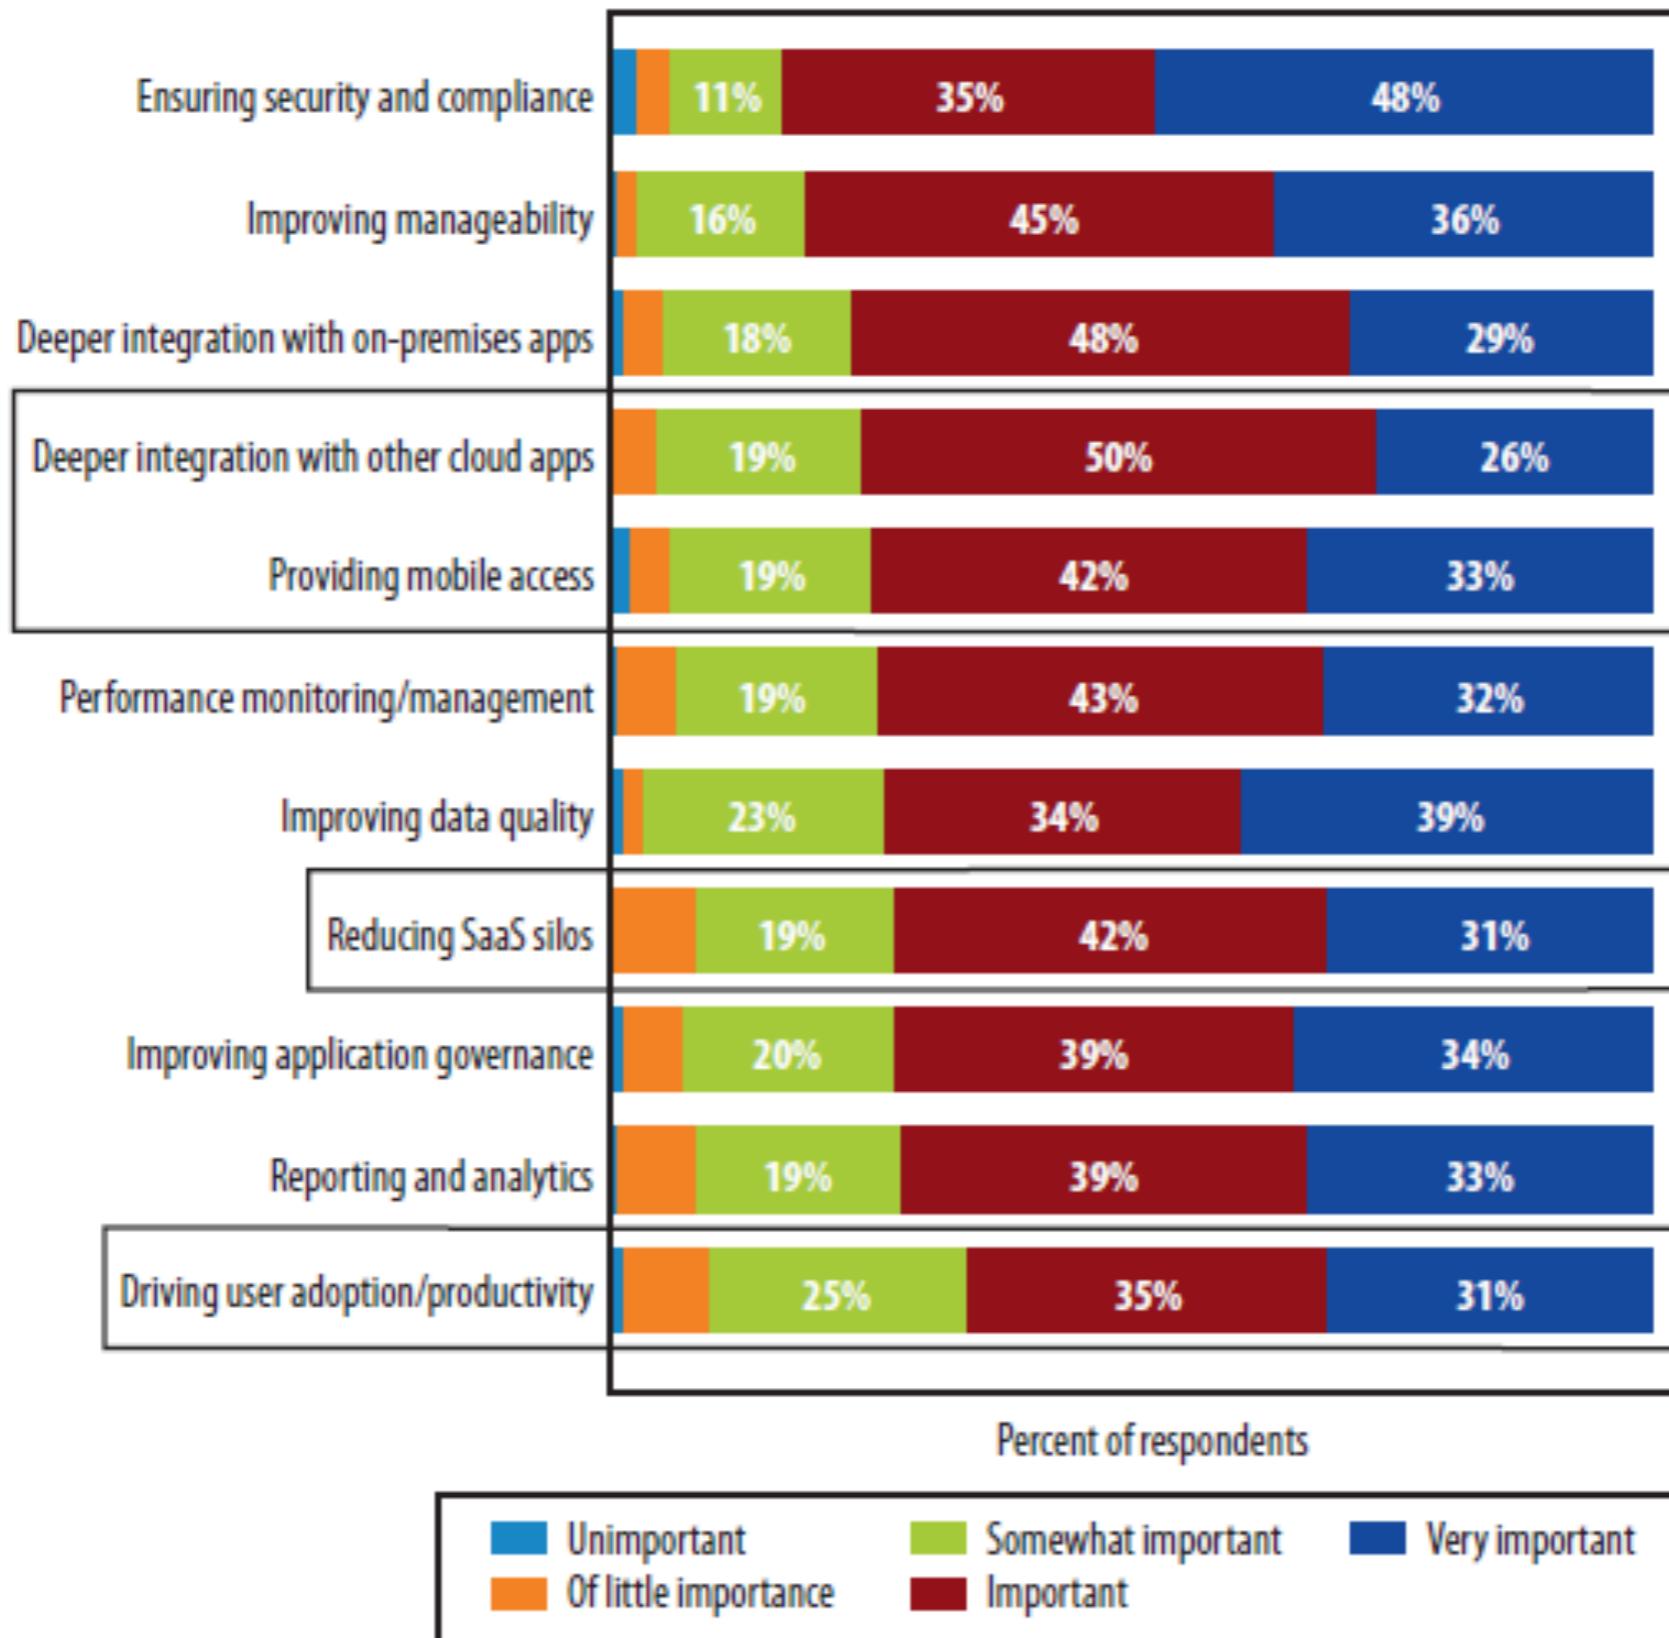
PERIODO DI 1 ANNO STANDARD					
Opzione di pagamento	Pagamento anticipato	Mensile*	Tariffa oraria**	Risparmio rispetto a tariffe on demand	Tariffa oraria on demand
Nessun pagamento anticipato	0,00 USD	219,73 USD	0,301 USD	30%	
Pagamento anticipato parziale	1.254,00 USD	104,39 USD	0,286 USD	33%	0,428 USD
Pagamento anticipato intero costo	2.457,00 USD	0,00 USD	0,28 USD	34%	

# Recap

## Concetti e parole chiave

- Risorse IT on demand
- Migliori benefici in un contesto affidabile
- Insieme di risorse informatiche virtualizzate
- Assegnamento rapido di risorse a run-time
- Sistemi ad alta scalabilità architetturale
- On demand
- Affidabilità
- Virtualizzazione
- Provisioning
- Scalabilità

# Cloud emerging challenges



# AWS

Amazon Web Services (AWS) è una piattaforma di servizi cloud sicura, in grado di offrire potenza di elaborazione, storage di database, distribuzione di contenuti e altre funzionalità per aiutare le imprese a ricalibrare le proprie risorse e crescere in modo organico.

# Gartner Magic Quadrant for Cloud Infrastructure as a Service 2018



**Gartner**<sup>®</sup>



# Region

AWS provides customers with the flexibility to place instances and store data within multiple geographic regions called Region. Each region is an independent collection of AWS resources in a defined geography.

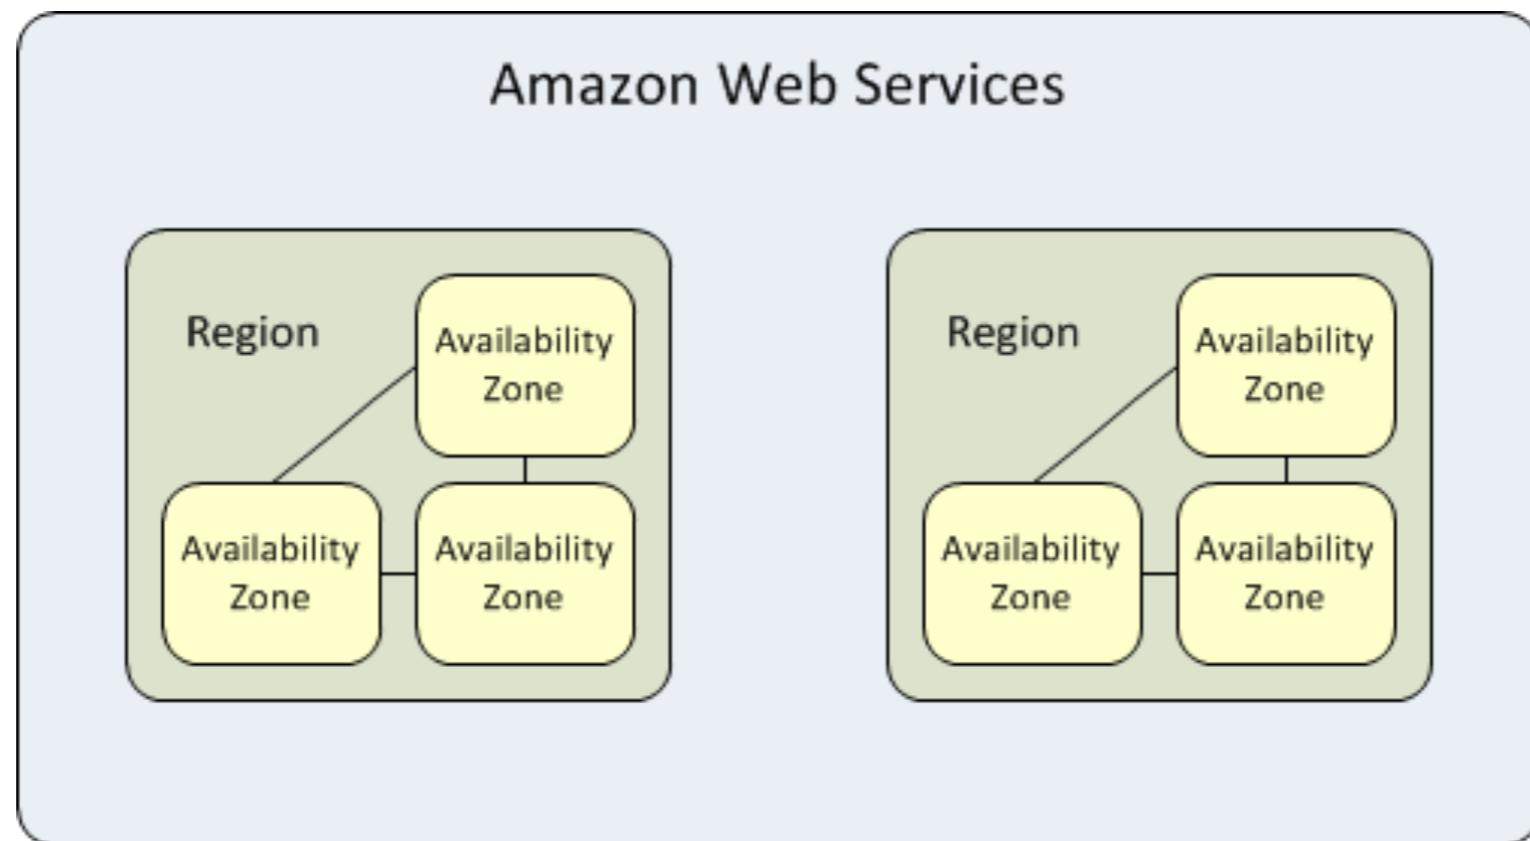


● Regioni

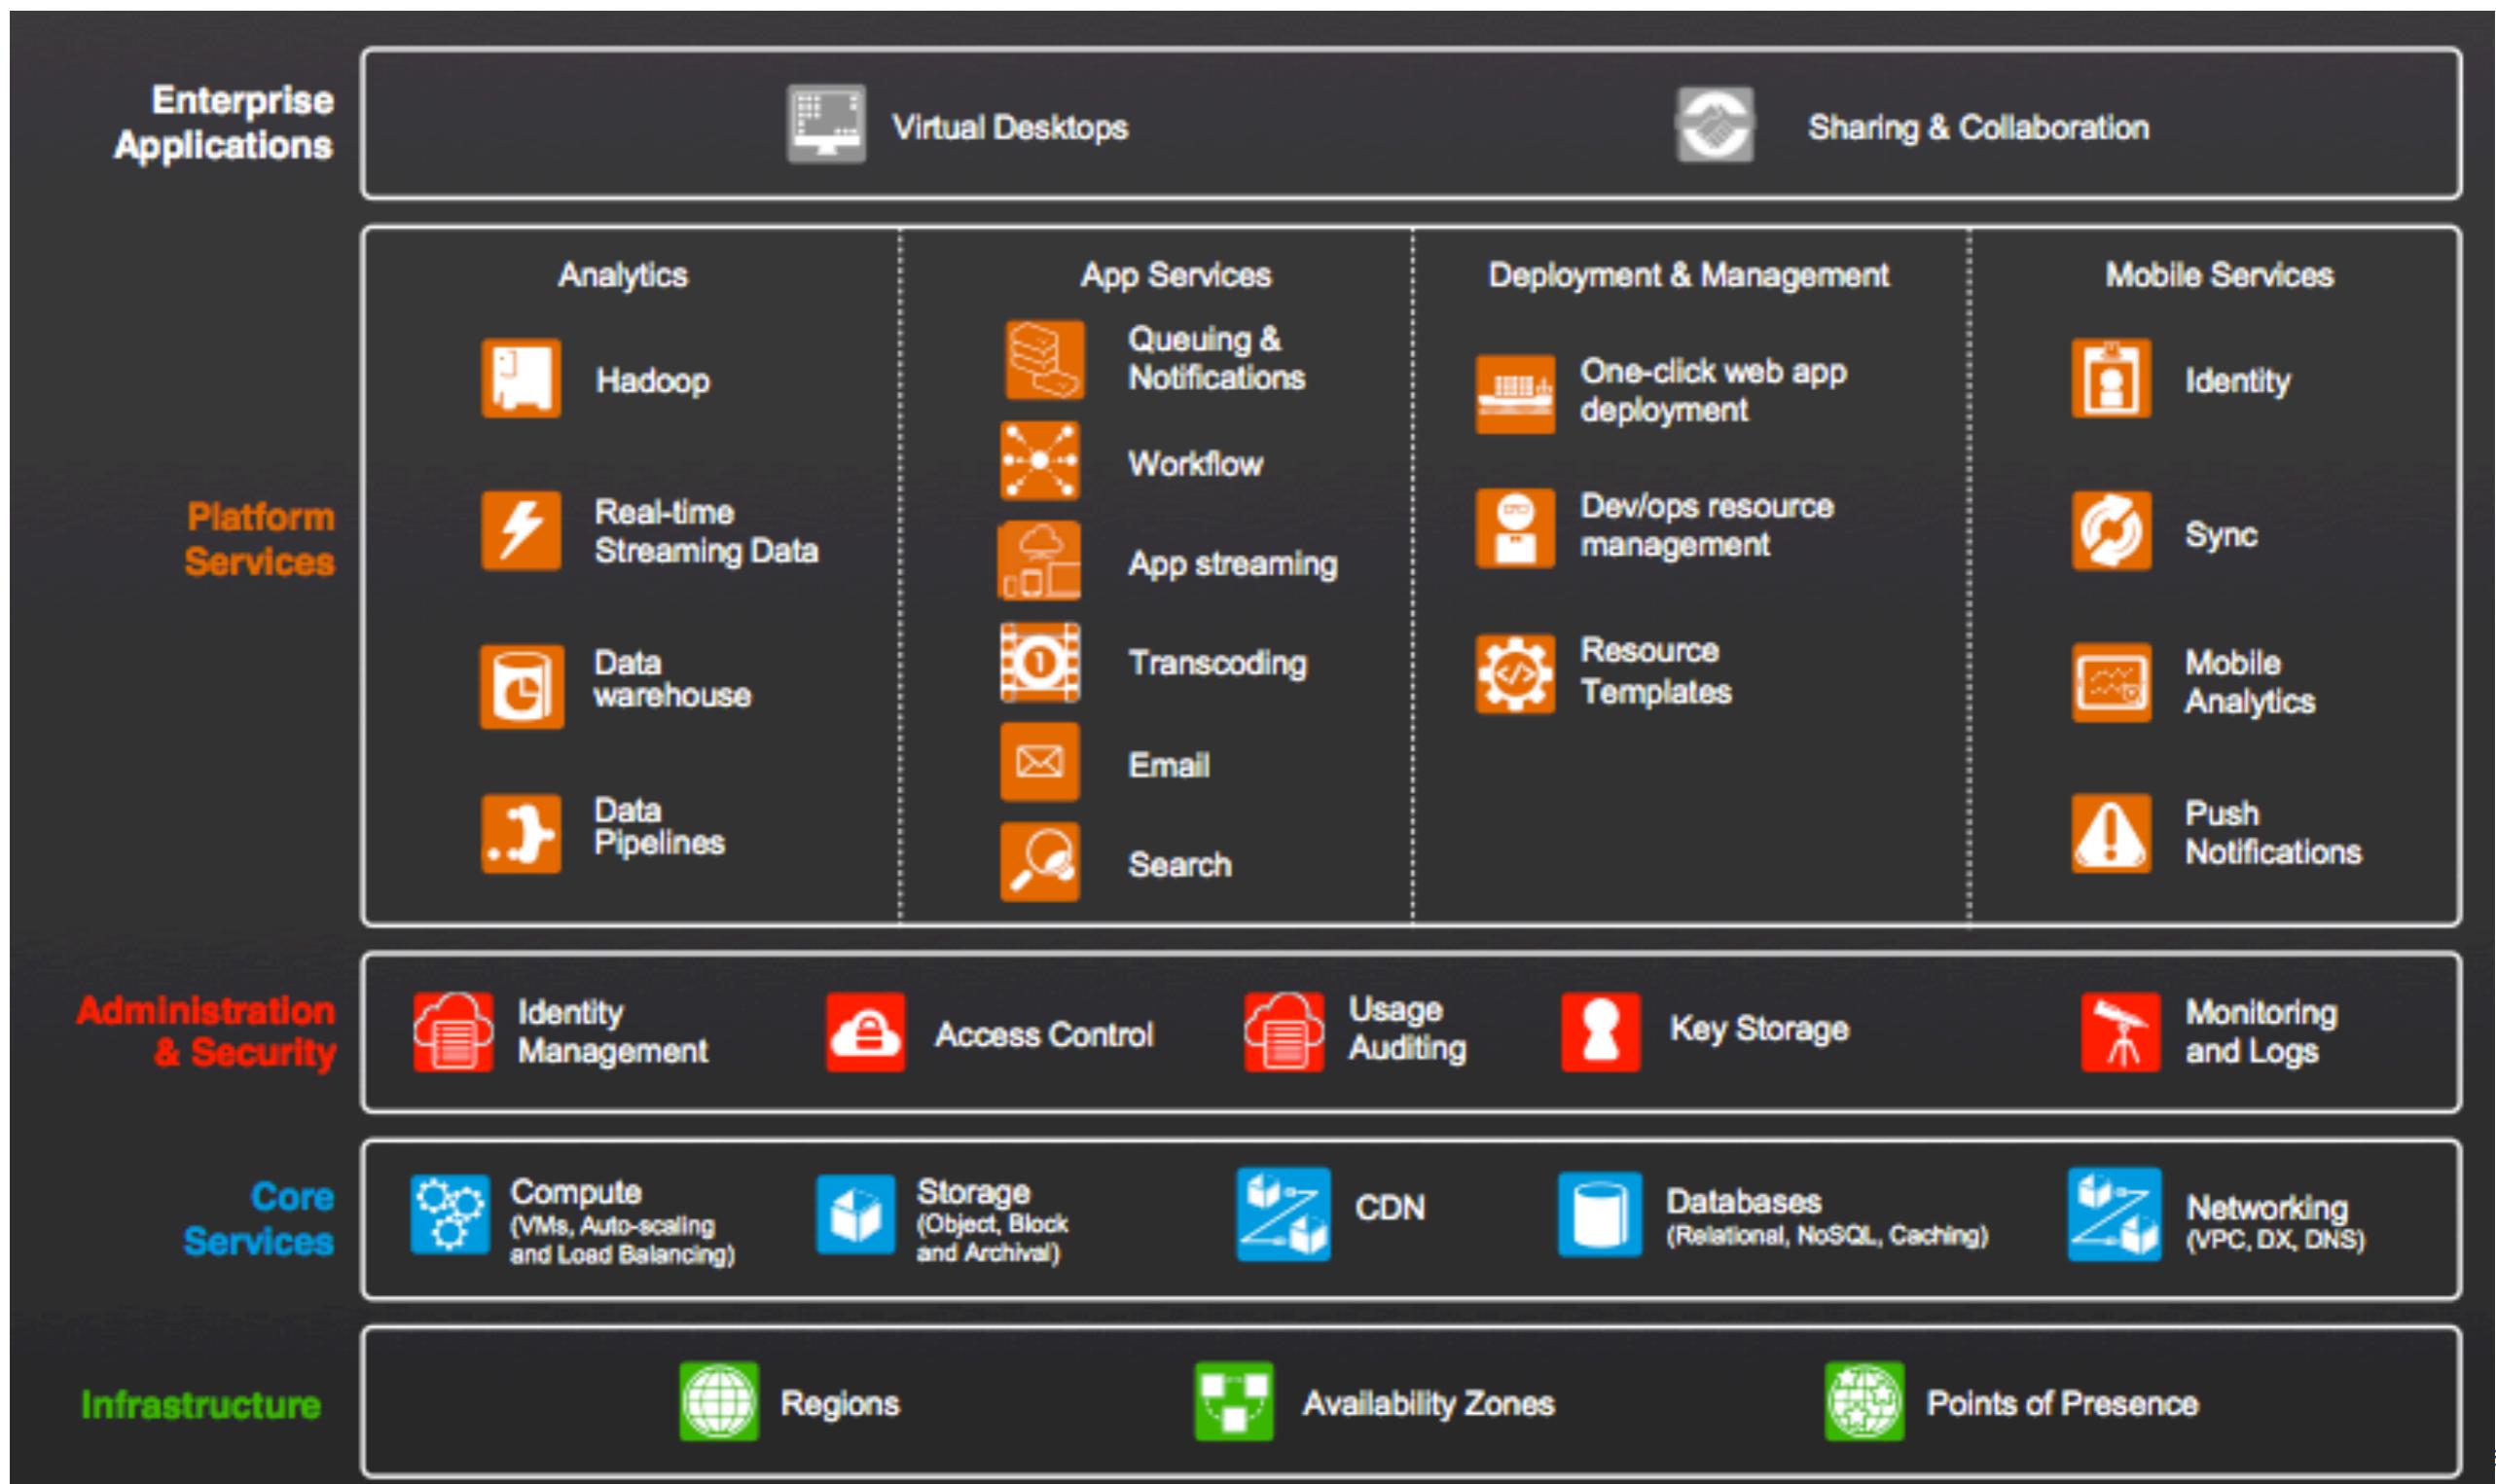
● Disponibile a breve

# Availability Zones

Each Region has multiple, isolated locations known as **Availability Zones**. Amazon EC2 provides you the ability to place resources, such as instances, and data in multiple locations.



# AWS Servizi





Quicksight



Rekognition



ML



SageMaker



Redshift



EC2



EMR



Athena



Lambda



S3



Storage

Alta  
scalabilità

Servizio

Web Store,  
No file system

Accesso via  
API

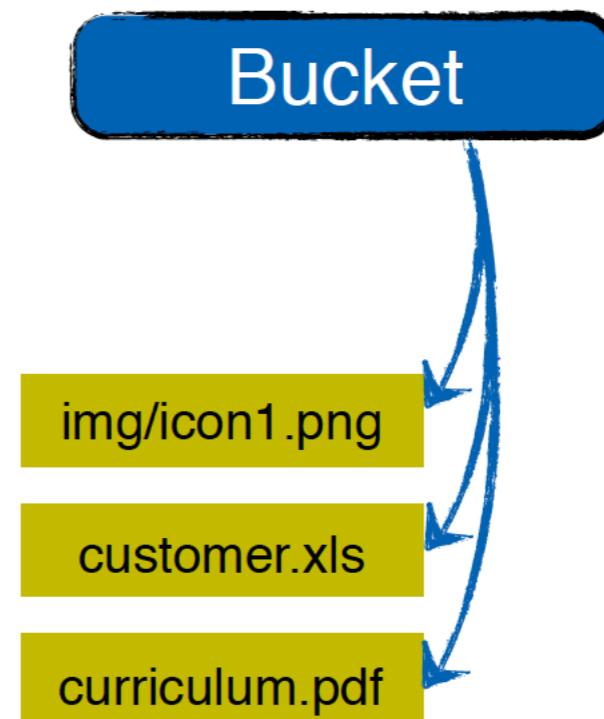
Veloce

Alta  
disponibilità

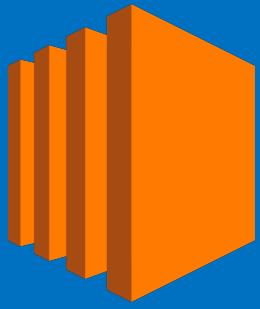
Economico



- Object Store
- Nativamente supportato da Spark, Hive, ...
- E' un servizio (non mi servono cluster)
- Oggetti (file illimitati)
- Basso costo
- Sicurezza

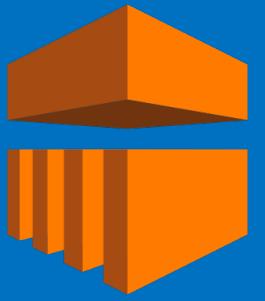


# AWS Elastic Compute Cloud (EC2)



- Virtual Machines
- AWS has everything from the tiny to gigantic
  - T2.Nano: 1 VCPU 512 MB Ram
  - X1.32xlarge: 128 VCPU 2000 GB Ram
  - GPUS! Useful for deep learning
- Priced per-second
- Options for On-Demand and “Spot Instances”
  - Spot instance: Auction for unused EC2 capacity; generally much cheaper than On-Demand

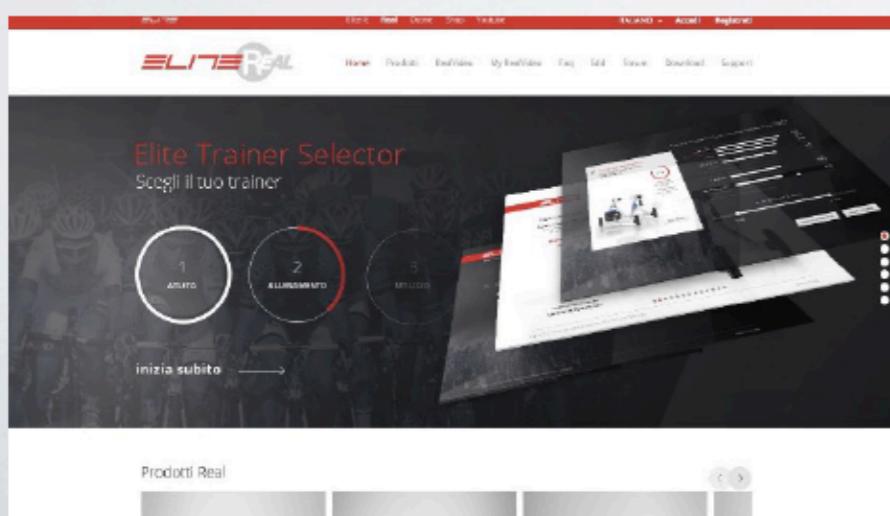
# Amazon Elastic Map Reduce



- Hosted Hadoop, Spark, HBase, Presto, Hive clusters
- Performs all necessary cluster scaling / provisioning automatically



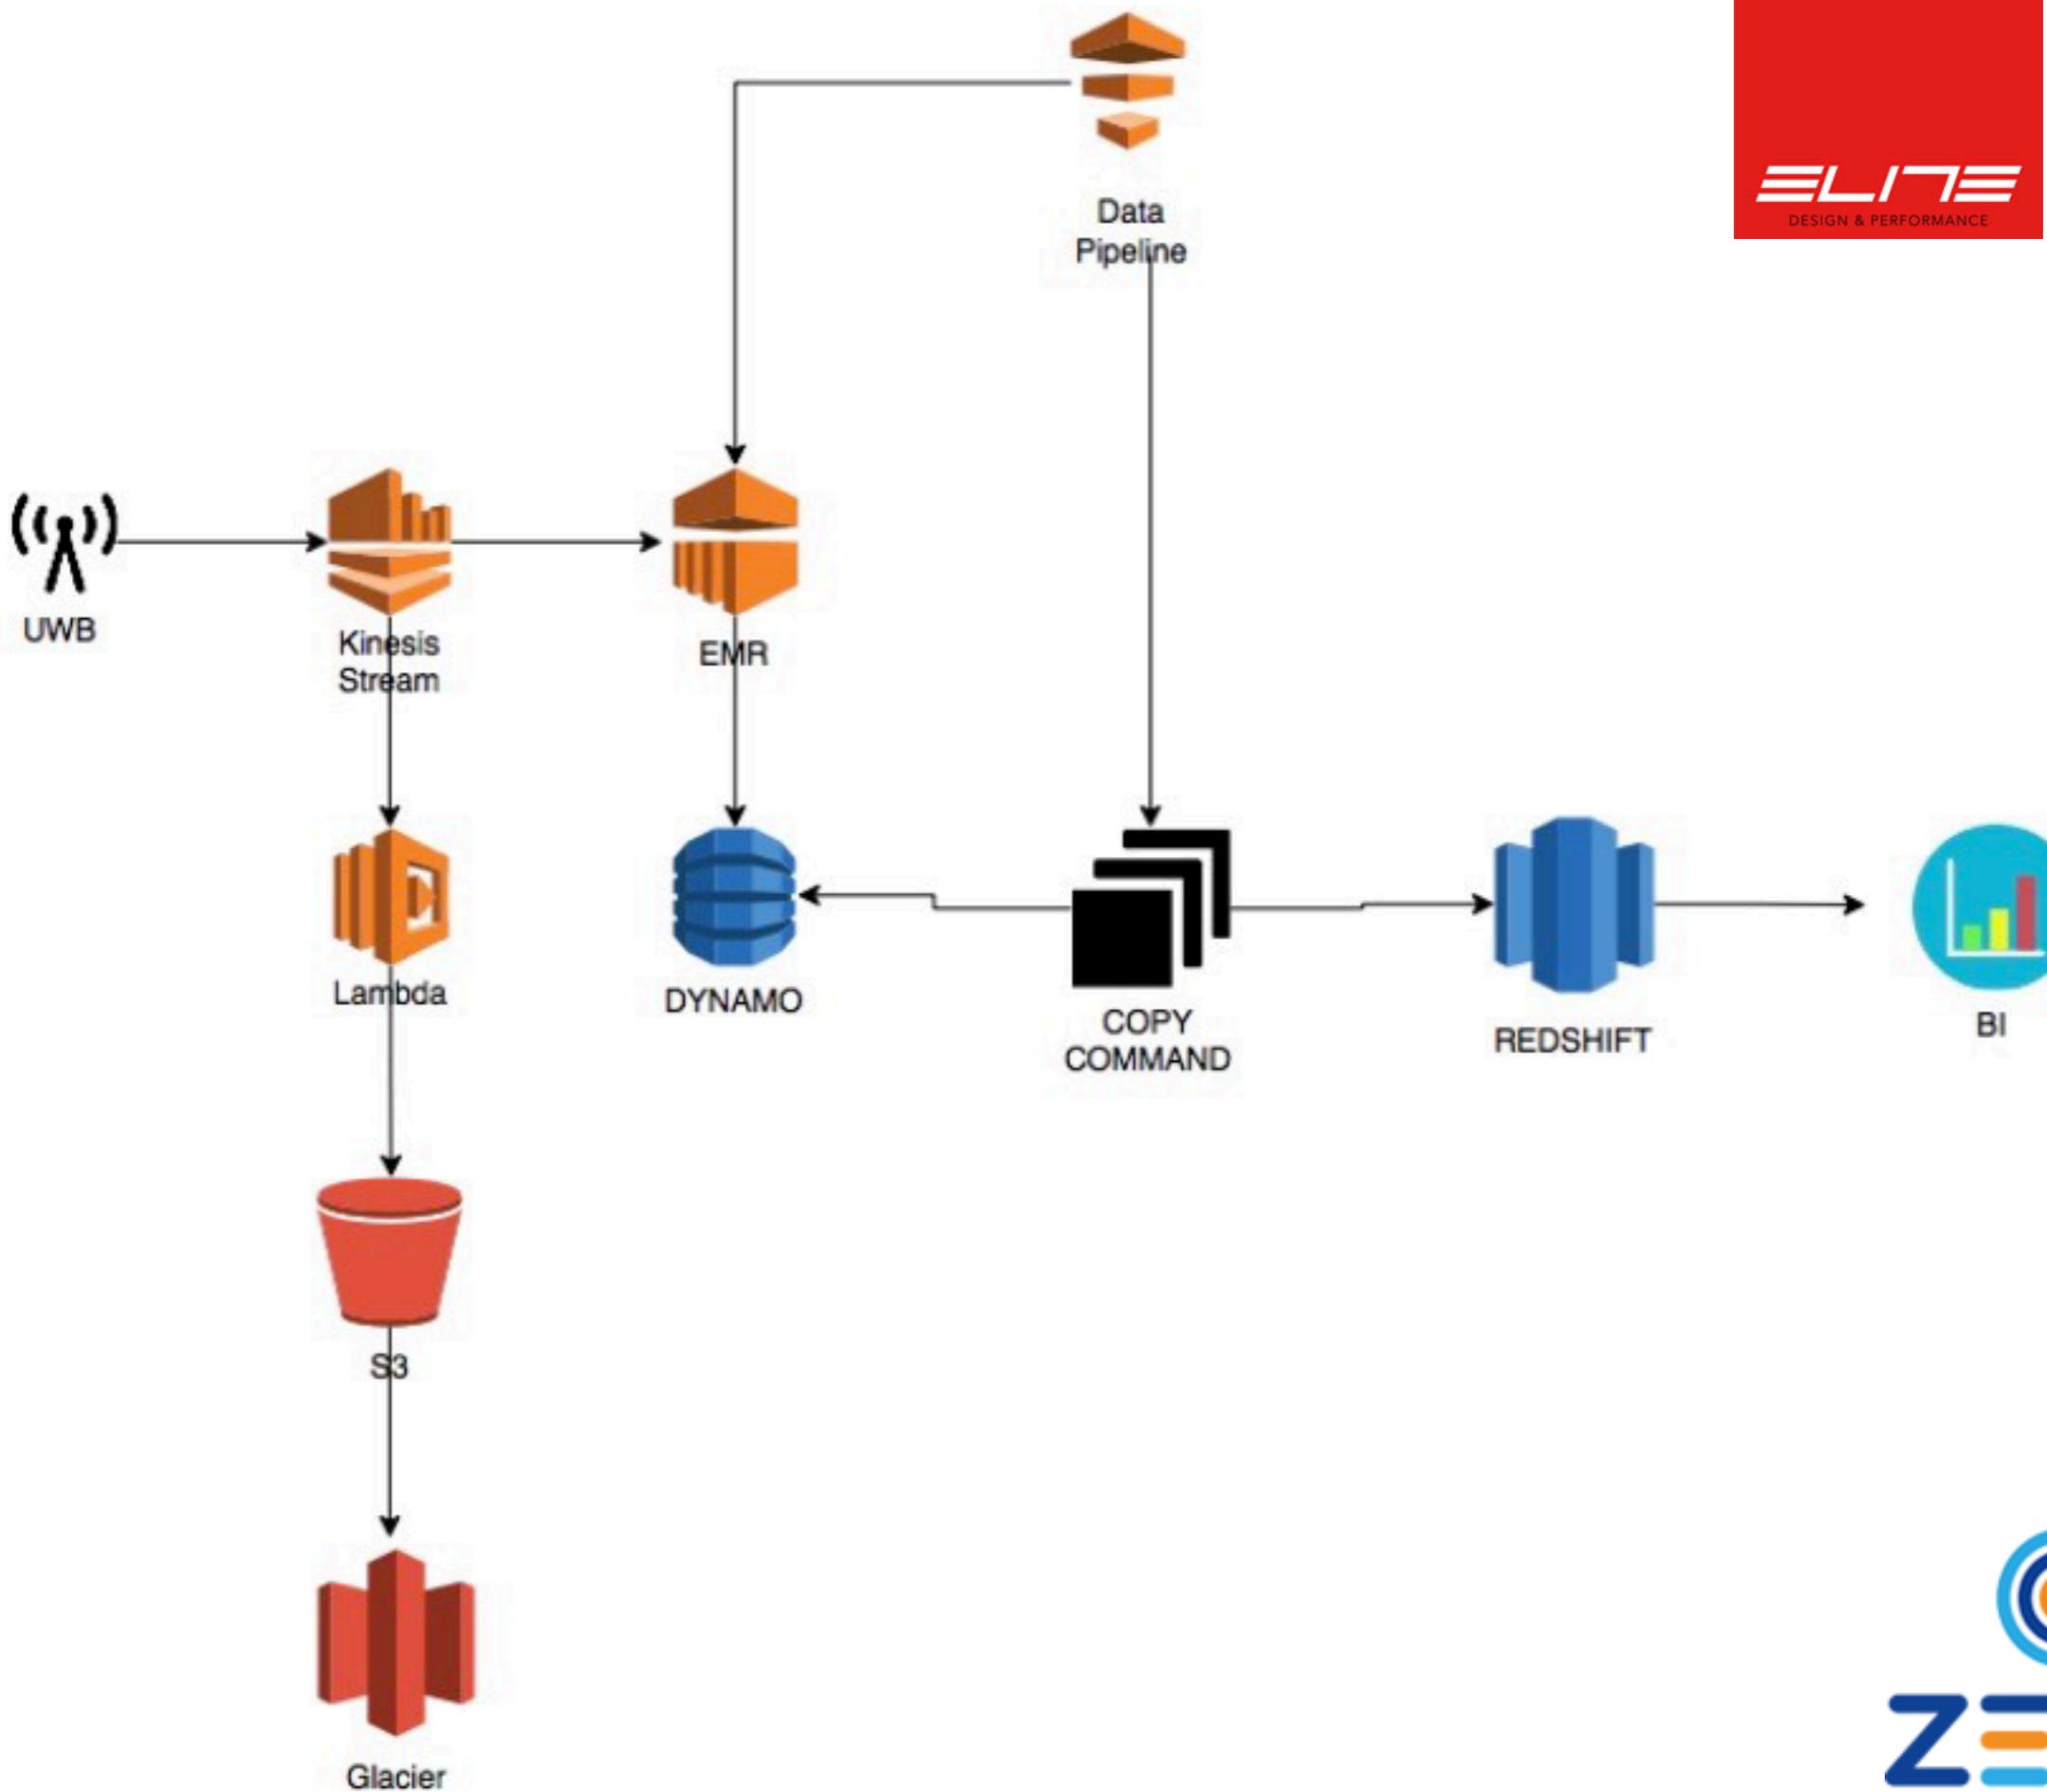
Recording  
video &  
Subscription

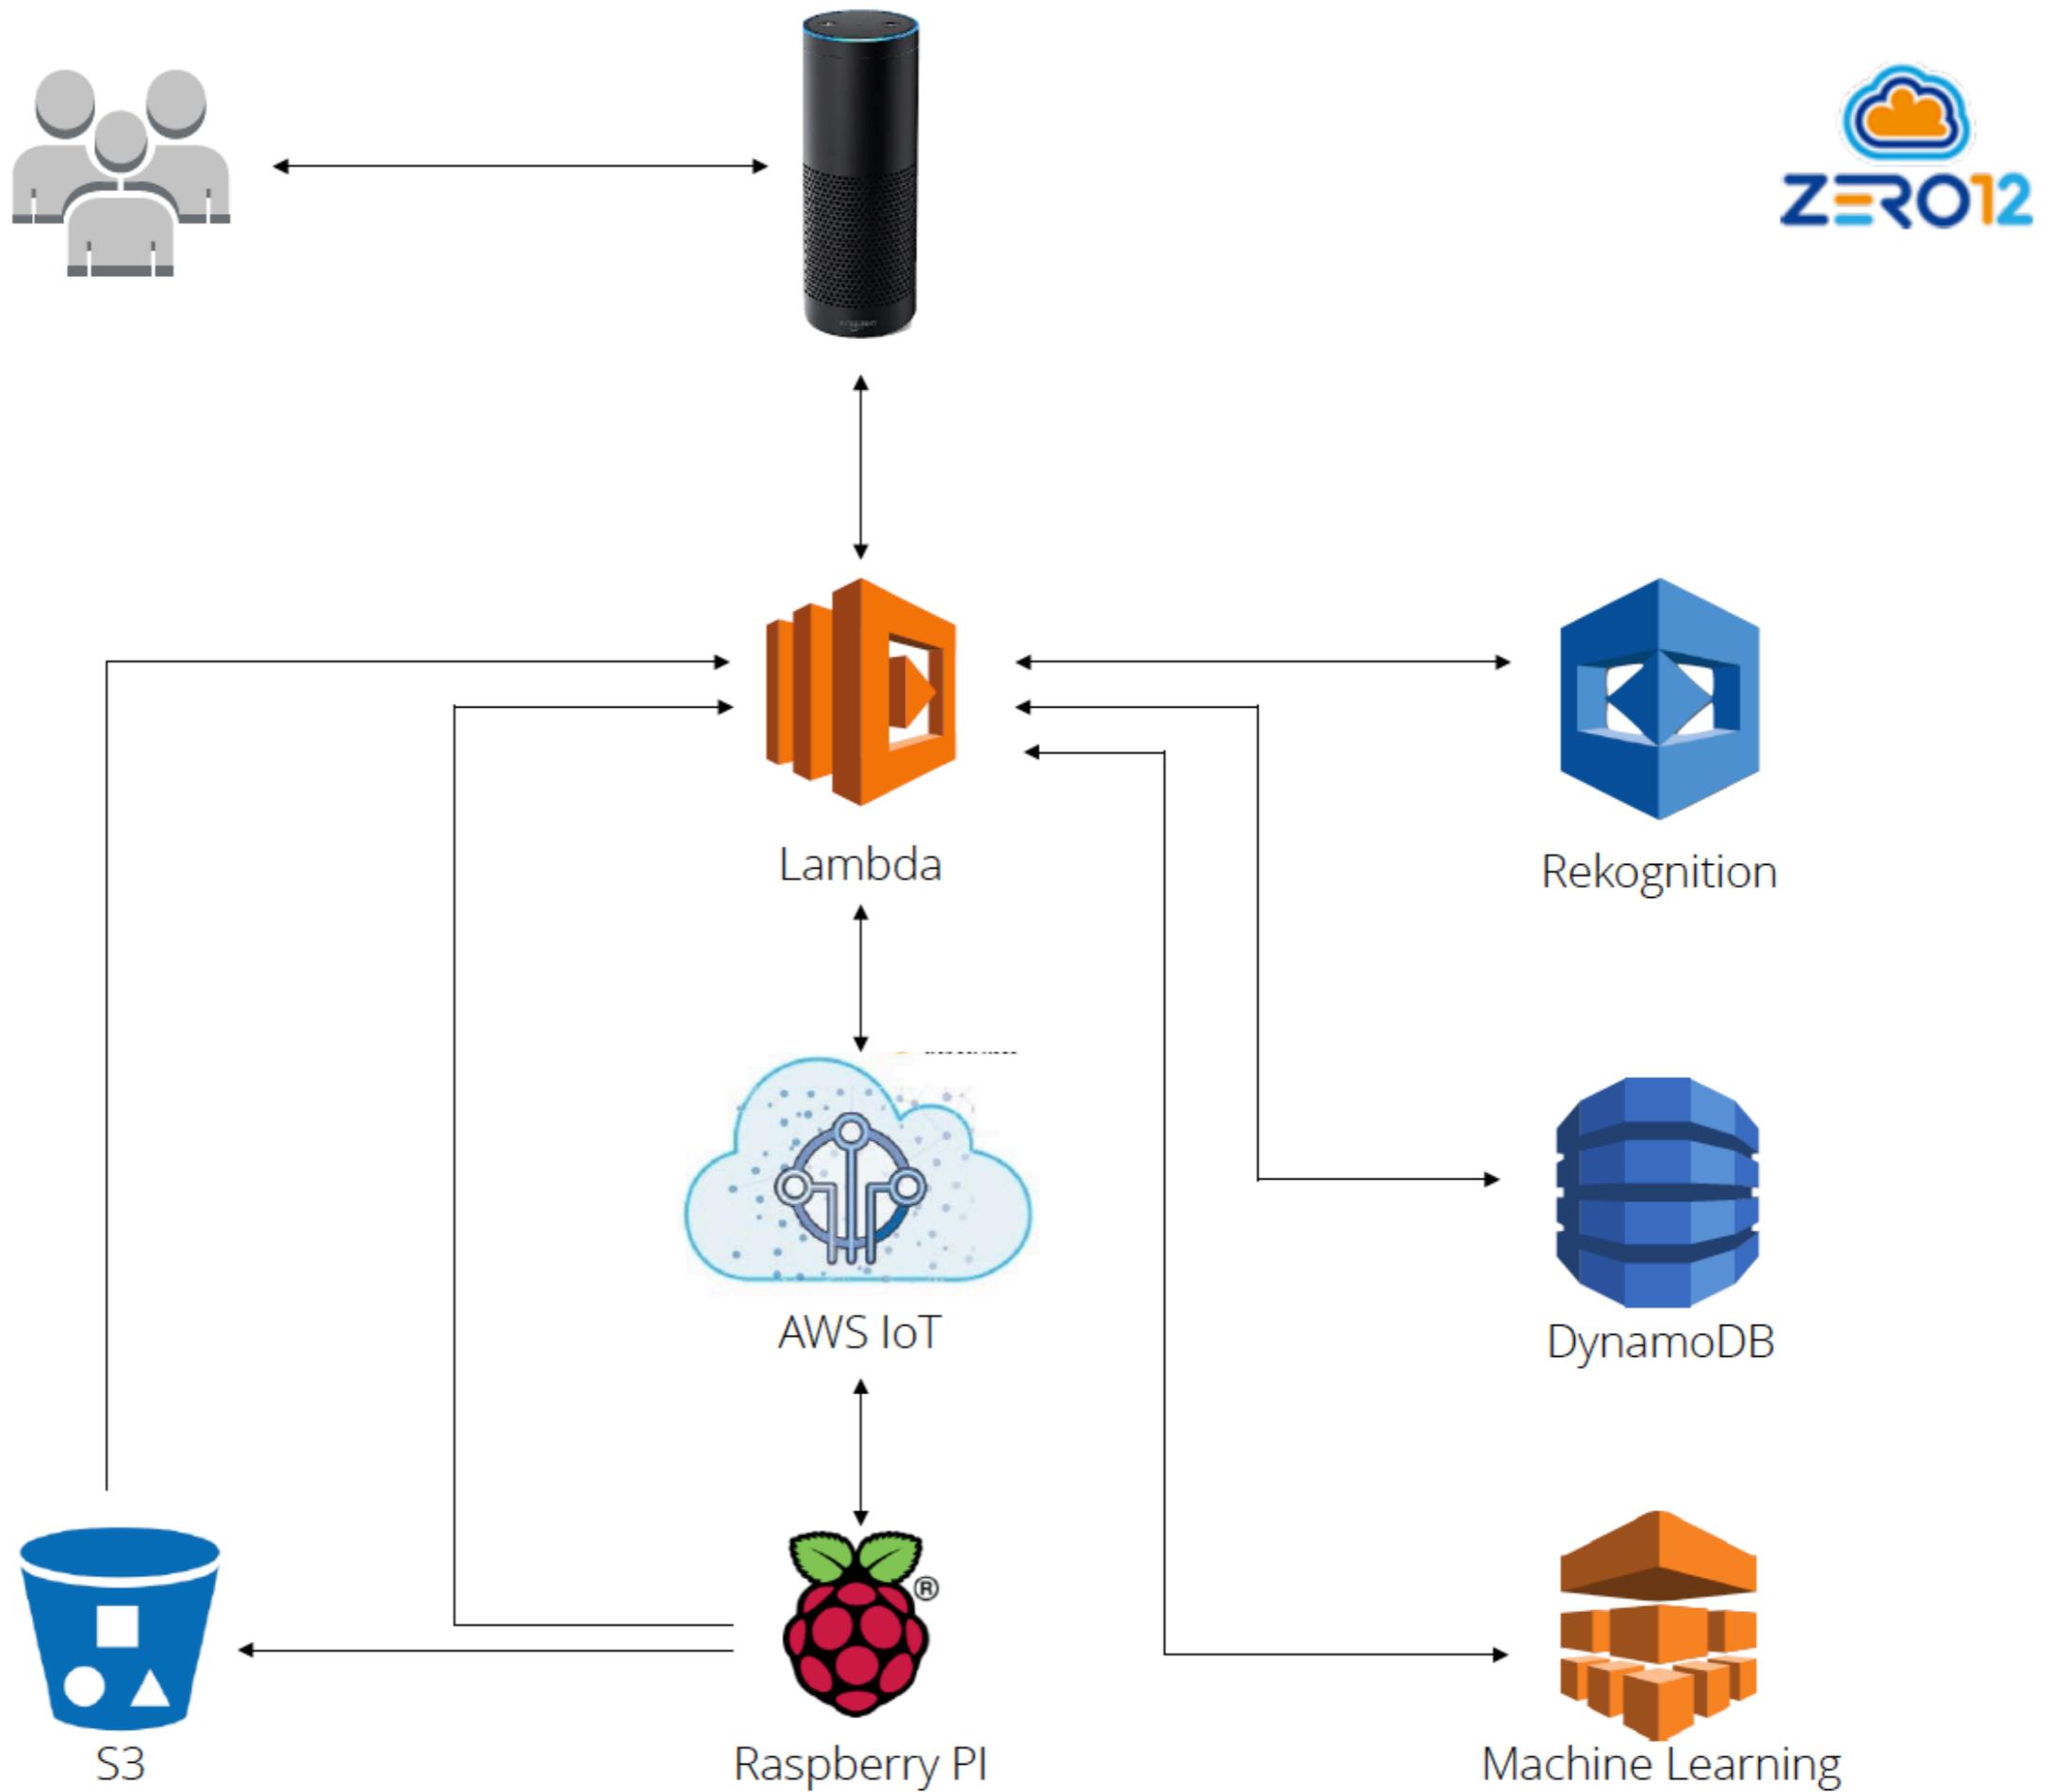


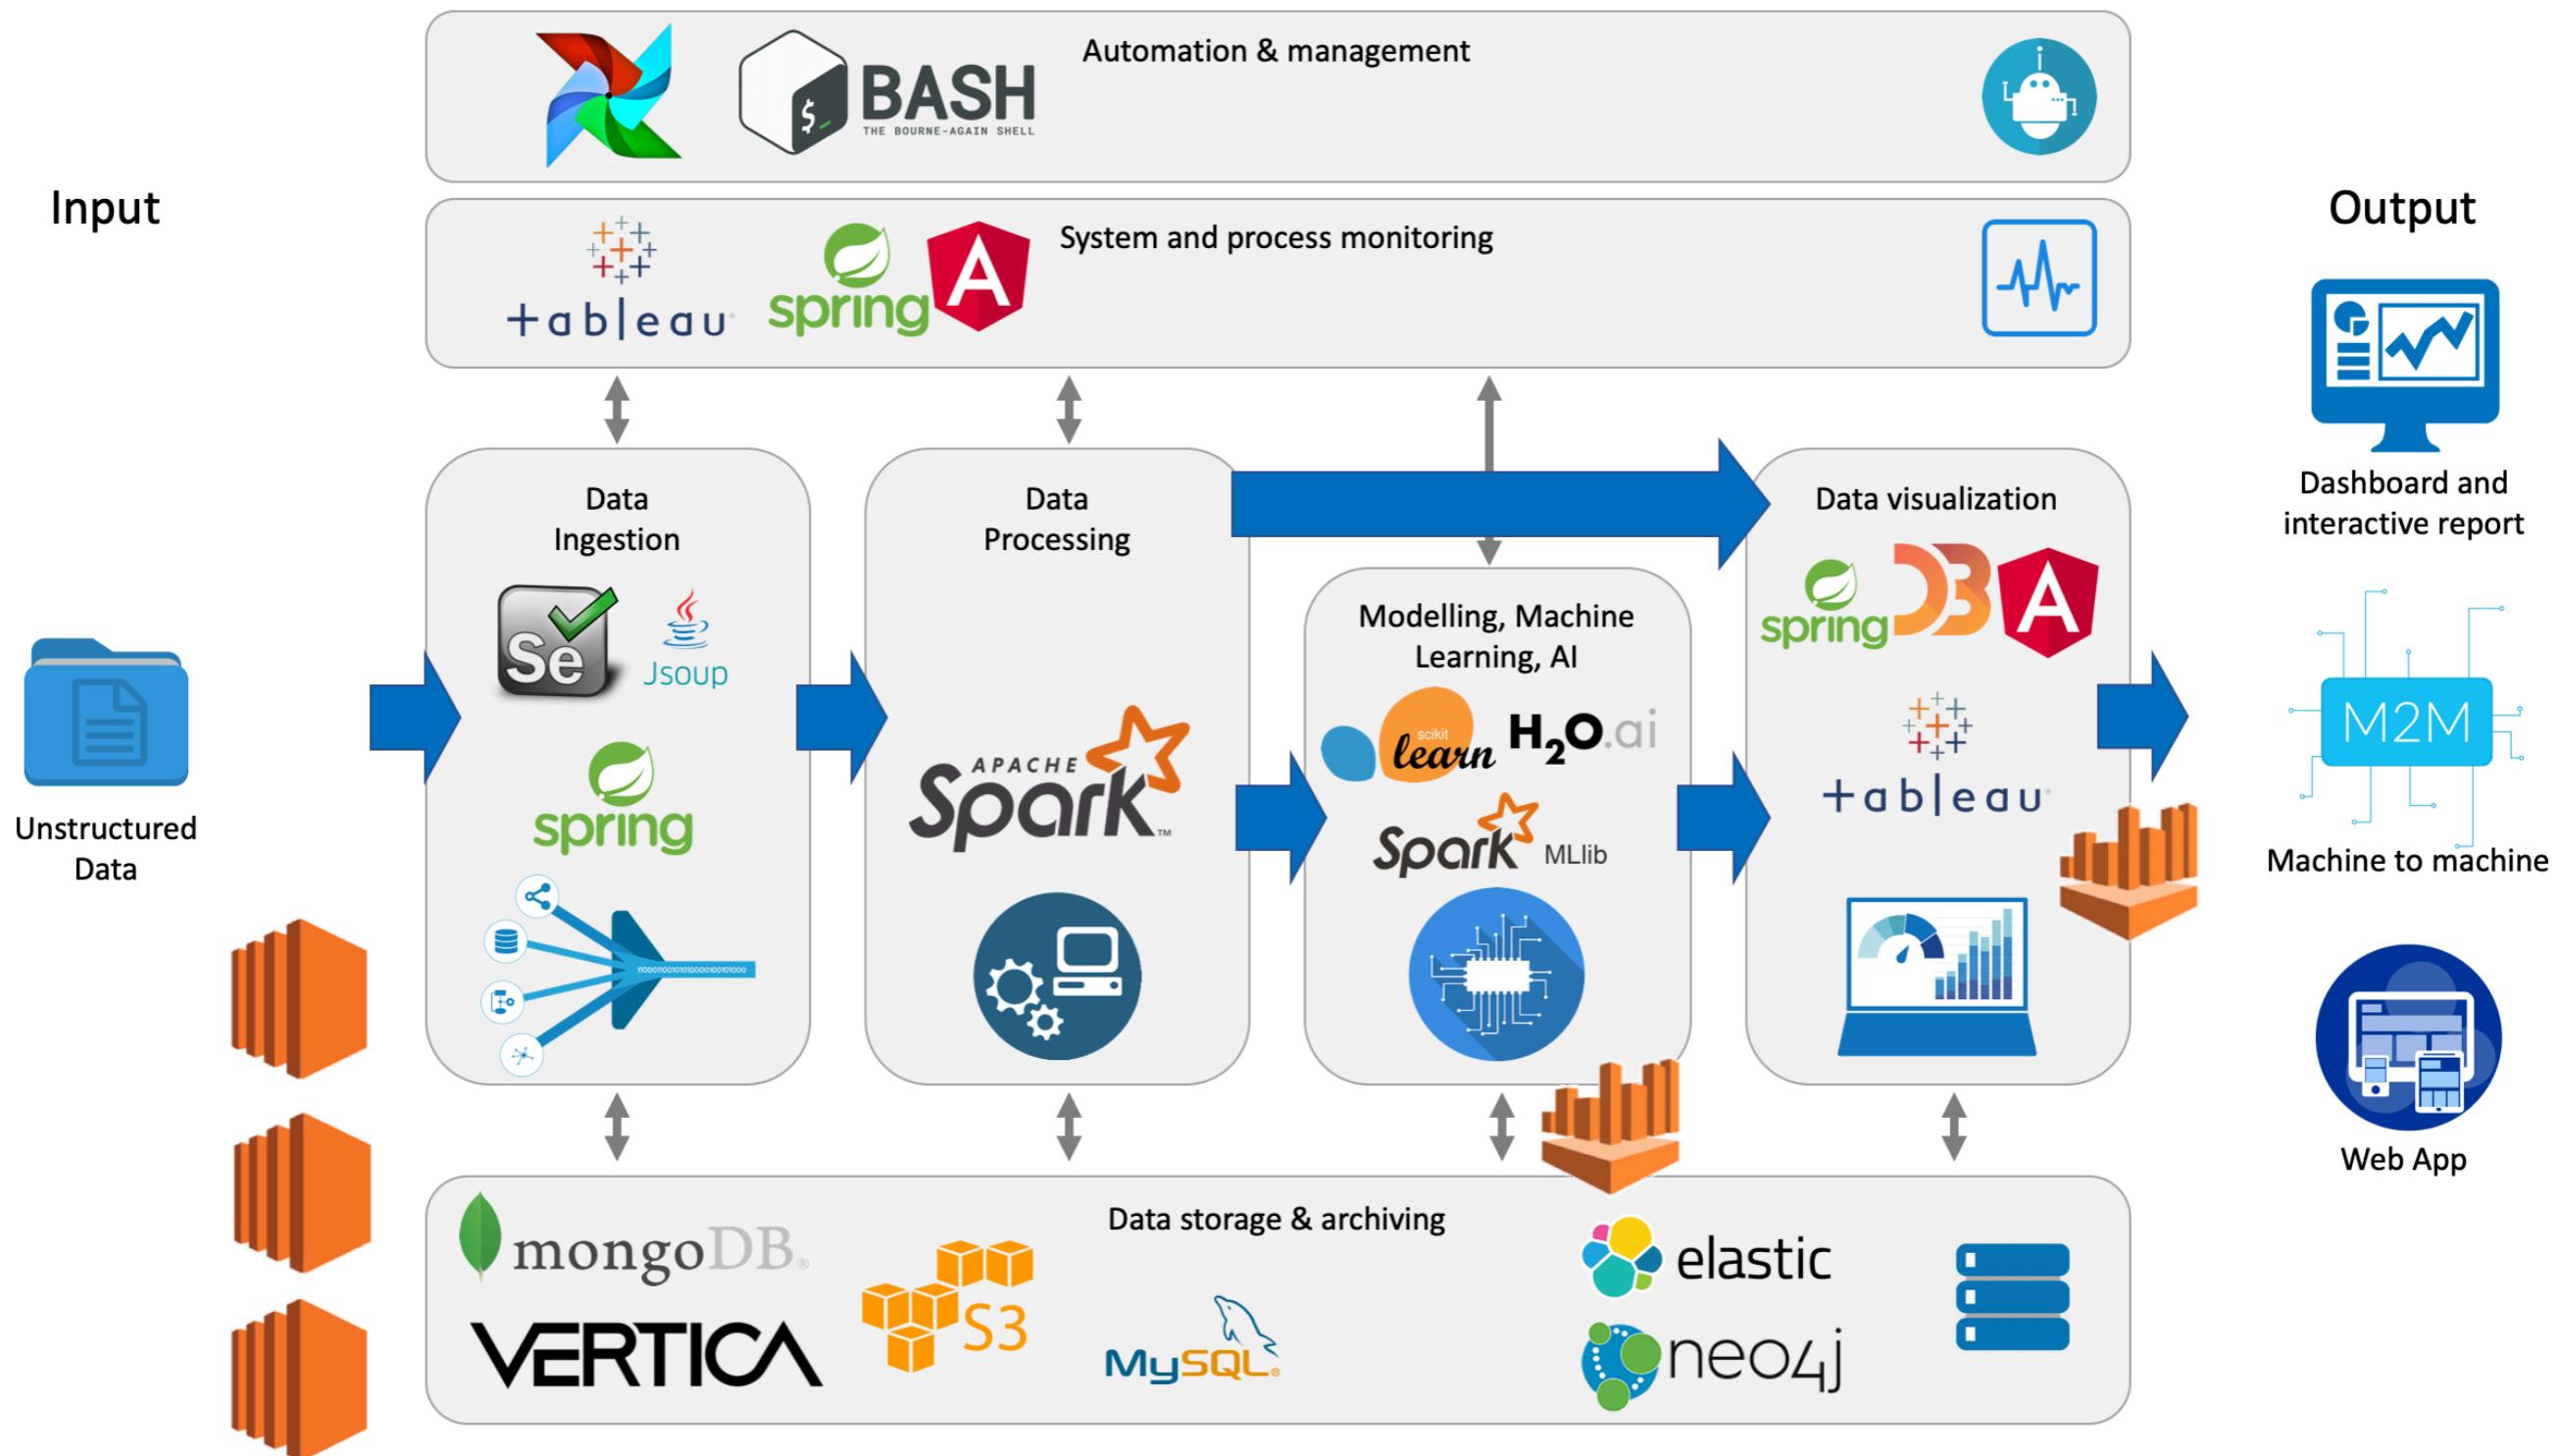
Configurator



VPC, EC2, S3, CDN, CodeDeploy, Autoscaling







# Micro-services

● docker

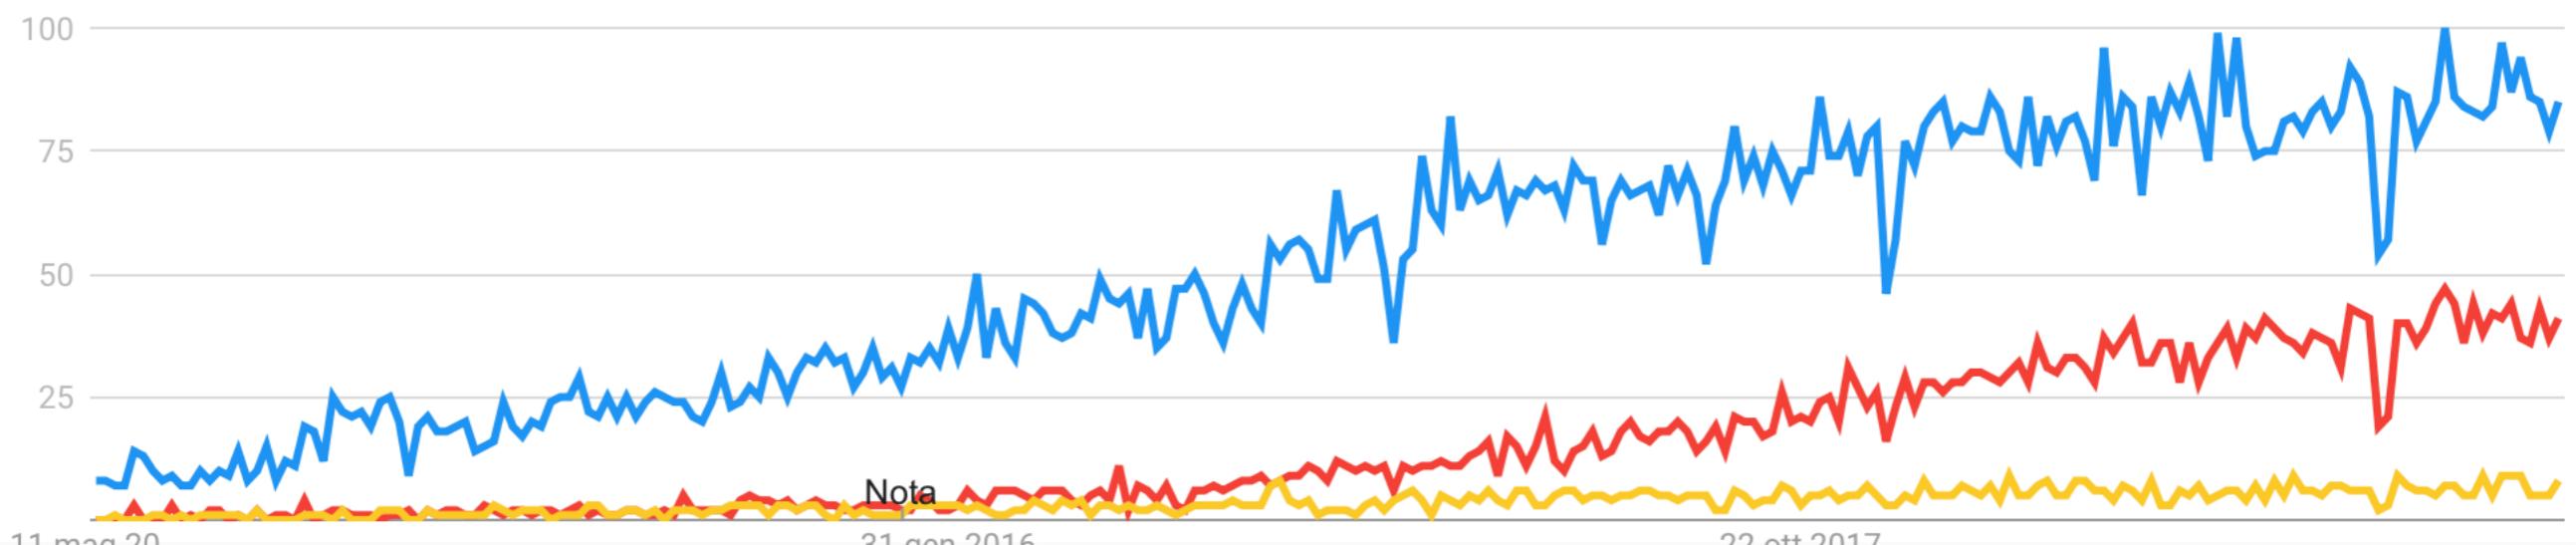
Termine di ricerca

● kubernetes

Termine di ricerca

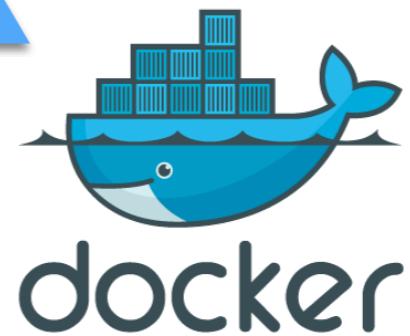
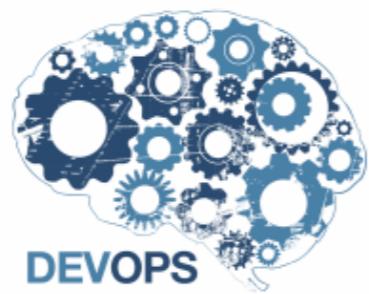
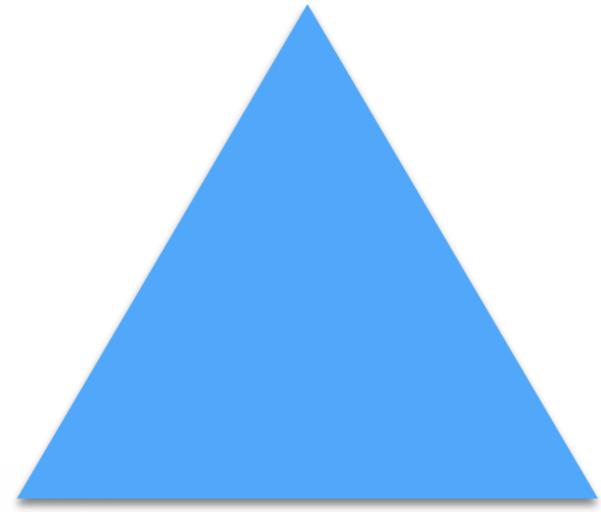
● microservices

Termine di ricerca



# Contesto

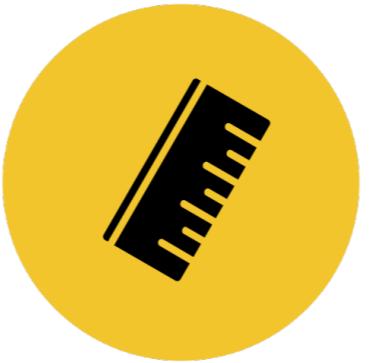
Microservices



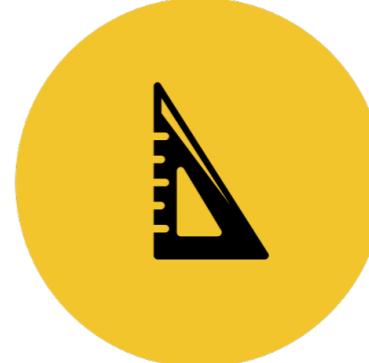
# Contesto



Manutenibilità



Monitoring



Scalability



Updates



Onboarding

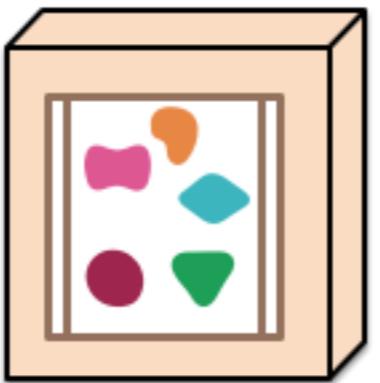
# Micro-sevices

James Lewis and Martin Fowler

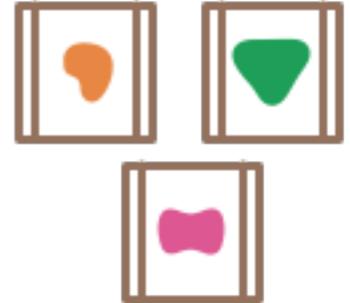
*A monolithic application puts all its functionality into a single process...*



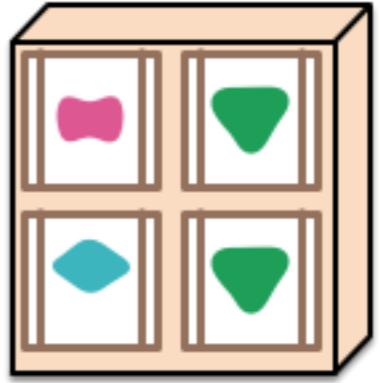
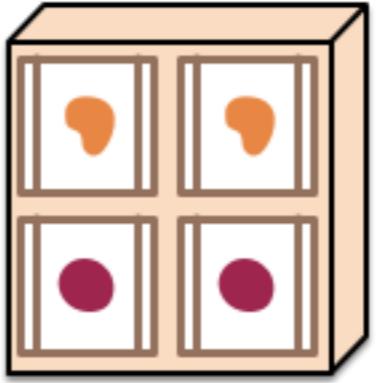
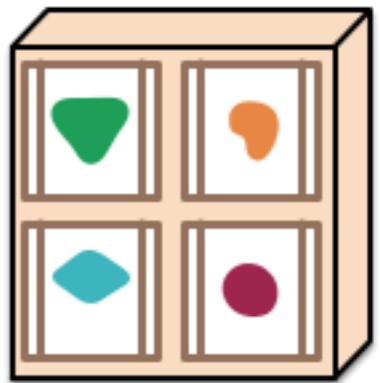
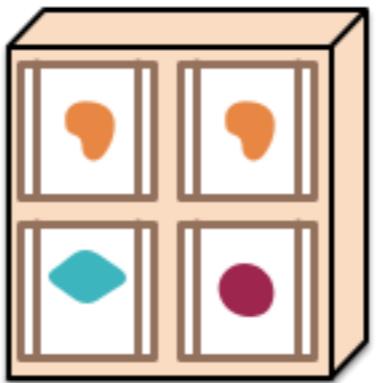
*... and scales by replicating the monolith on multiple servers*



*A microservices architecture puts each element of functionality into a separate service...*

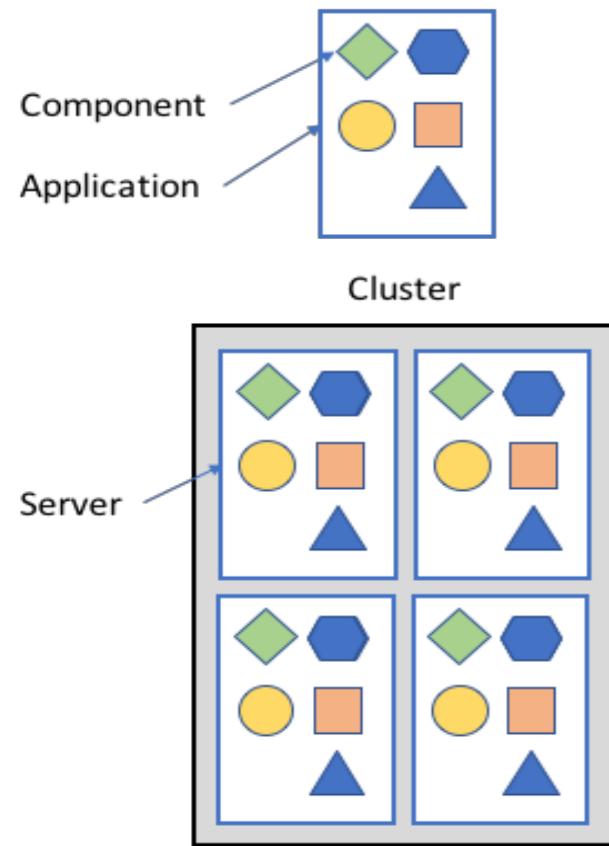


*... and scales by distributing these services across servers, replicating as needed.*

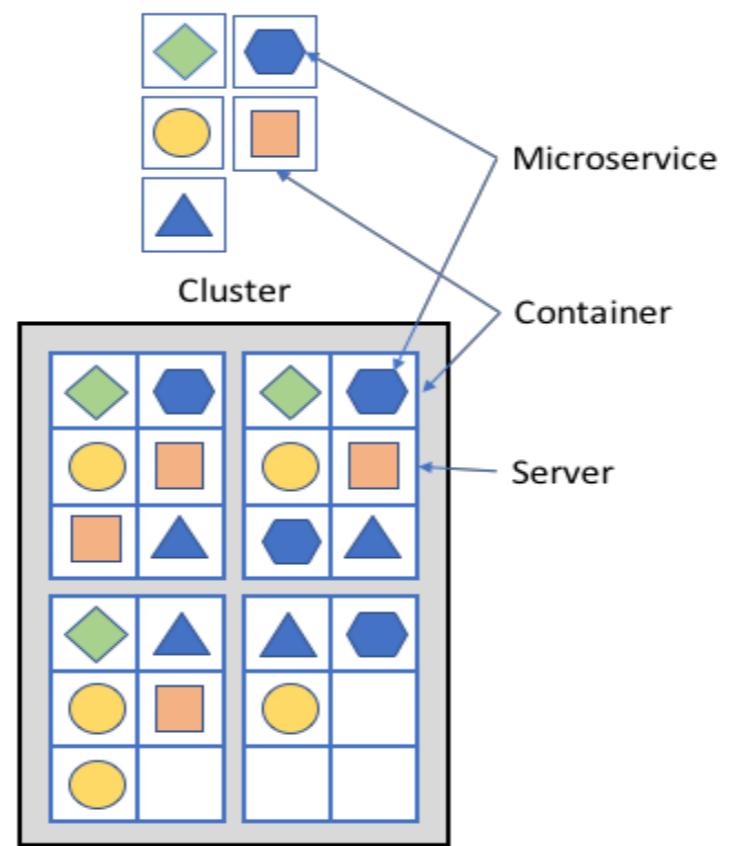


Functional system decomposition into manageable and independently deployable components

## Traditional Monolithic Architecture



## Microservice Architecture



"Monoliths and microservices are not a simple binary choice."

Both are fuzzy definitions that mean many systems would lie in a blurred boundary area among the two"

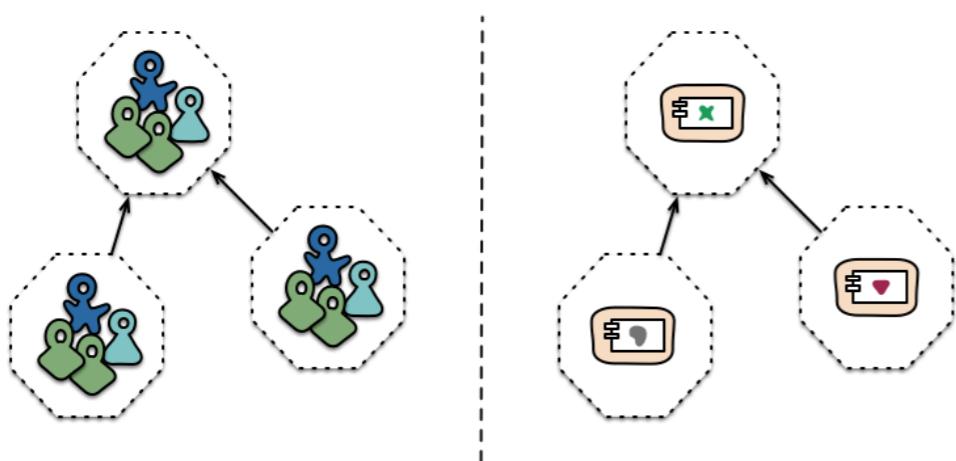
# Caratteristiche



Componentization  
via services



Products,  
not projects

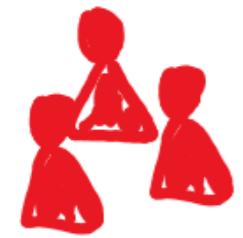


Organized around  
business capabilities

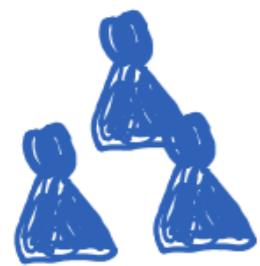


Decentralized

Organize around business services



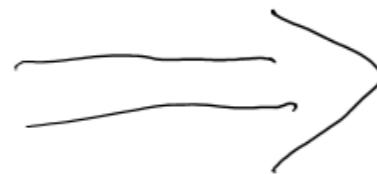
UI Team



Middleware  
guys



DBAs



Supply



Orders

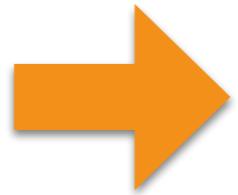


Recommender





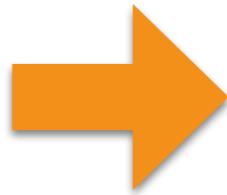
Monolithic  
app



Scritta in un  
linguaggio  
(C++)



Tiered  
architecture  
(MVC pattern)



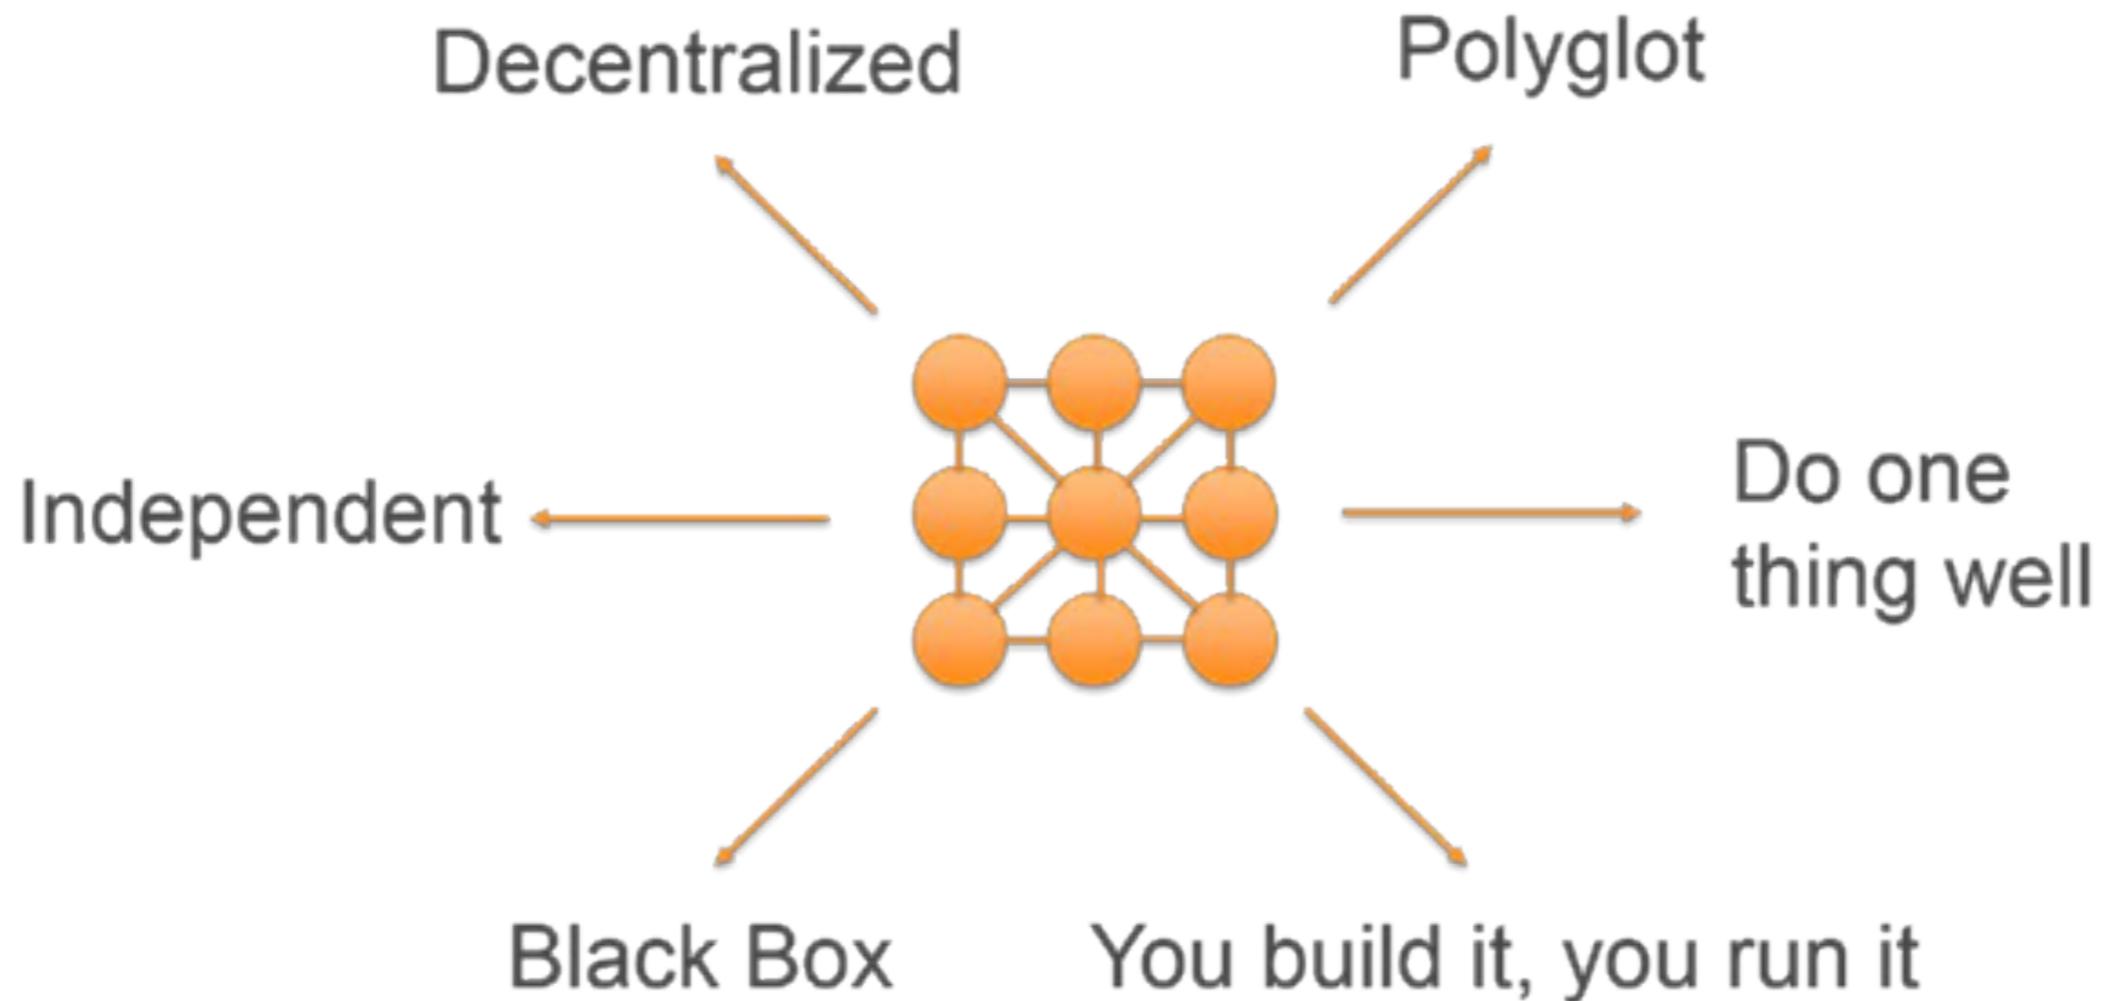
Java

Polyglot  
microservices

Mix di  
linguaggi



# Recap



# Docker

Docker tries to solve the "run anywhere" problem using  
(easily) Linux containers.

Perchè?

# Development environment



Your script / program

python3

Lib1  
v2

Lib2  
v2

# Production environment



Your script / program

python3

Lib1  
v2

Lib2  
v2

python2

Lib1  
v1

Lib2  
v1

## Development environment



Your script / program



## Production environment



Your script / program



Conflict!!!

# Development environment



Your script / program

python3

Lib1  
v2

Lib2  
v2

# Production environment



Your script / program

python3

Lib1  
v2

Lib2  
v2

Docker container

My script / program

python2

Lib1  
v1

Lib2  
v1

Docker container

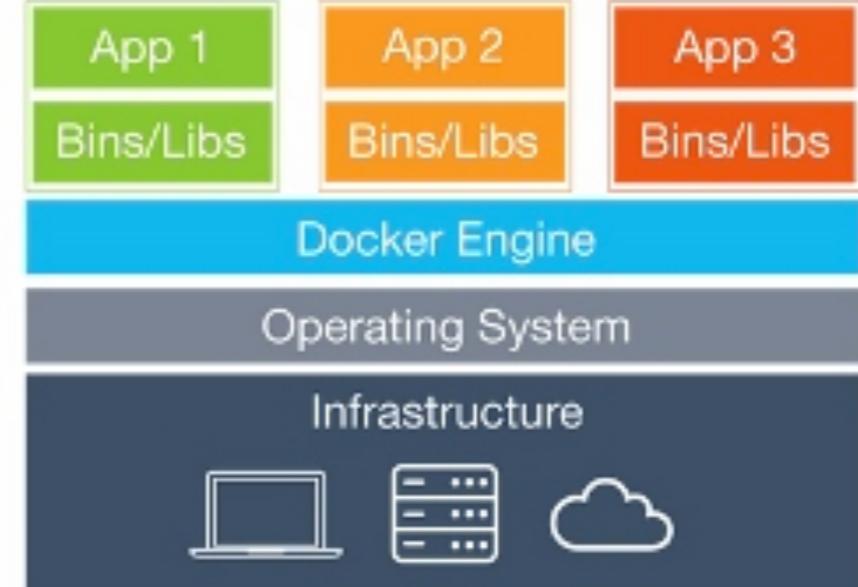
Docker Engine

# Efficienza

Perché è diverso da una virtual machine?  
(VMware, Xen, ...)



Virtual Machines



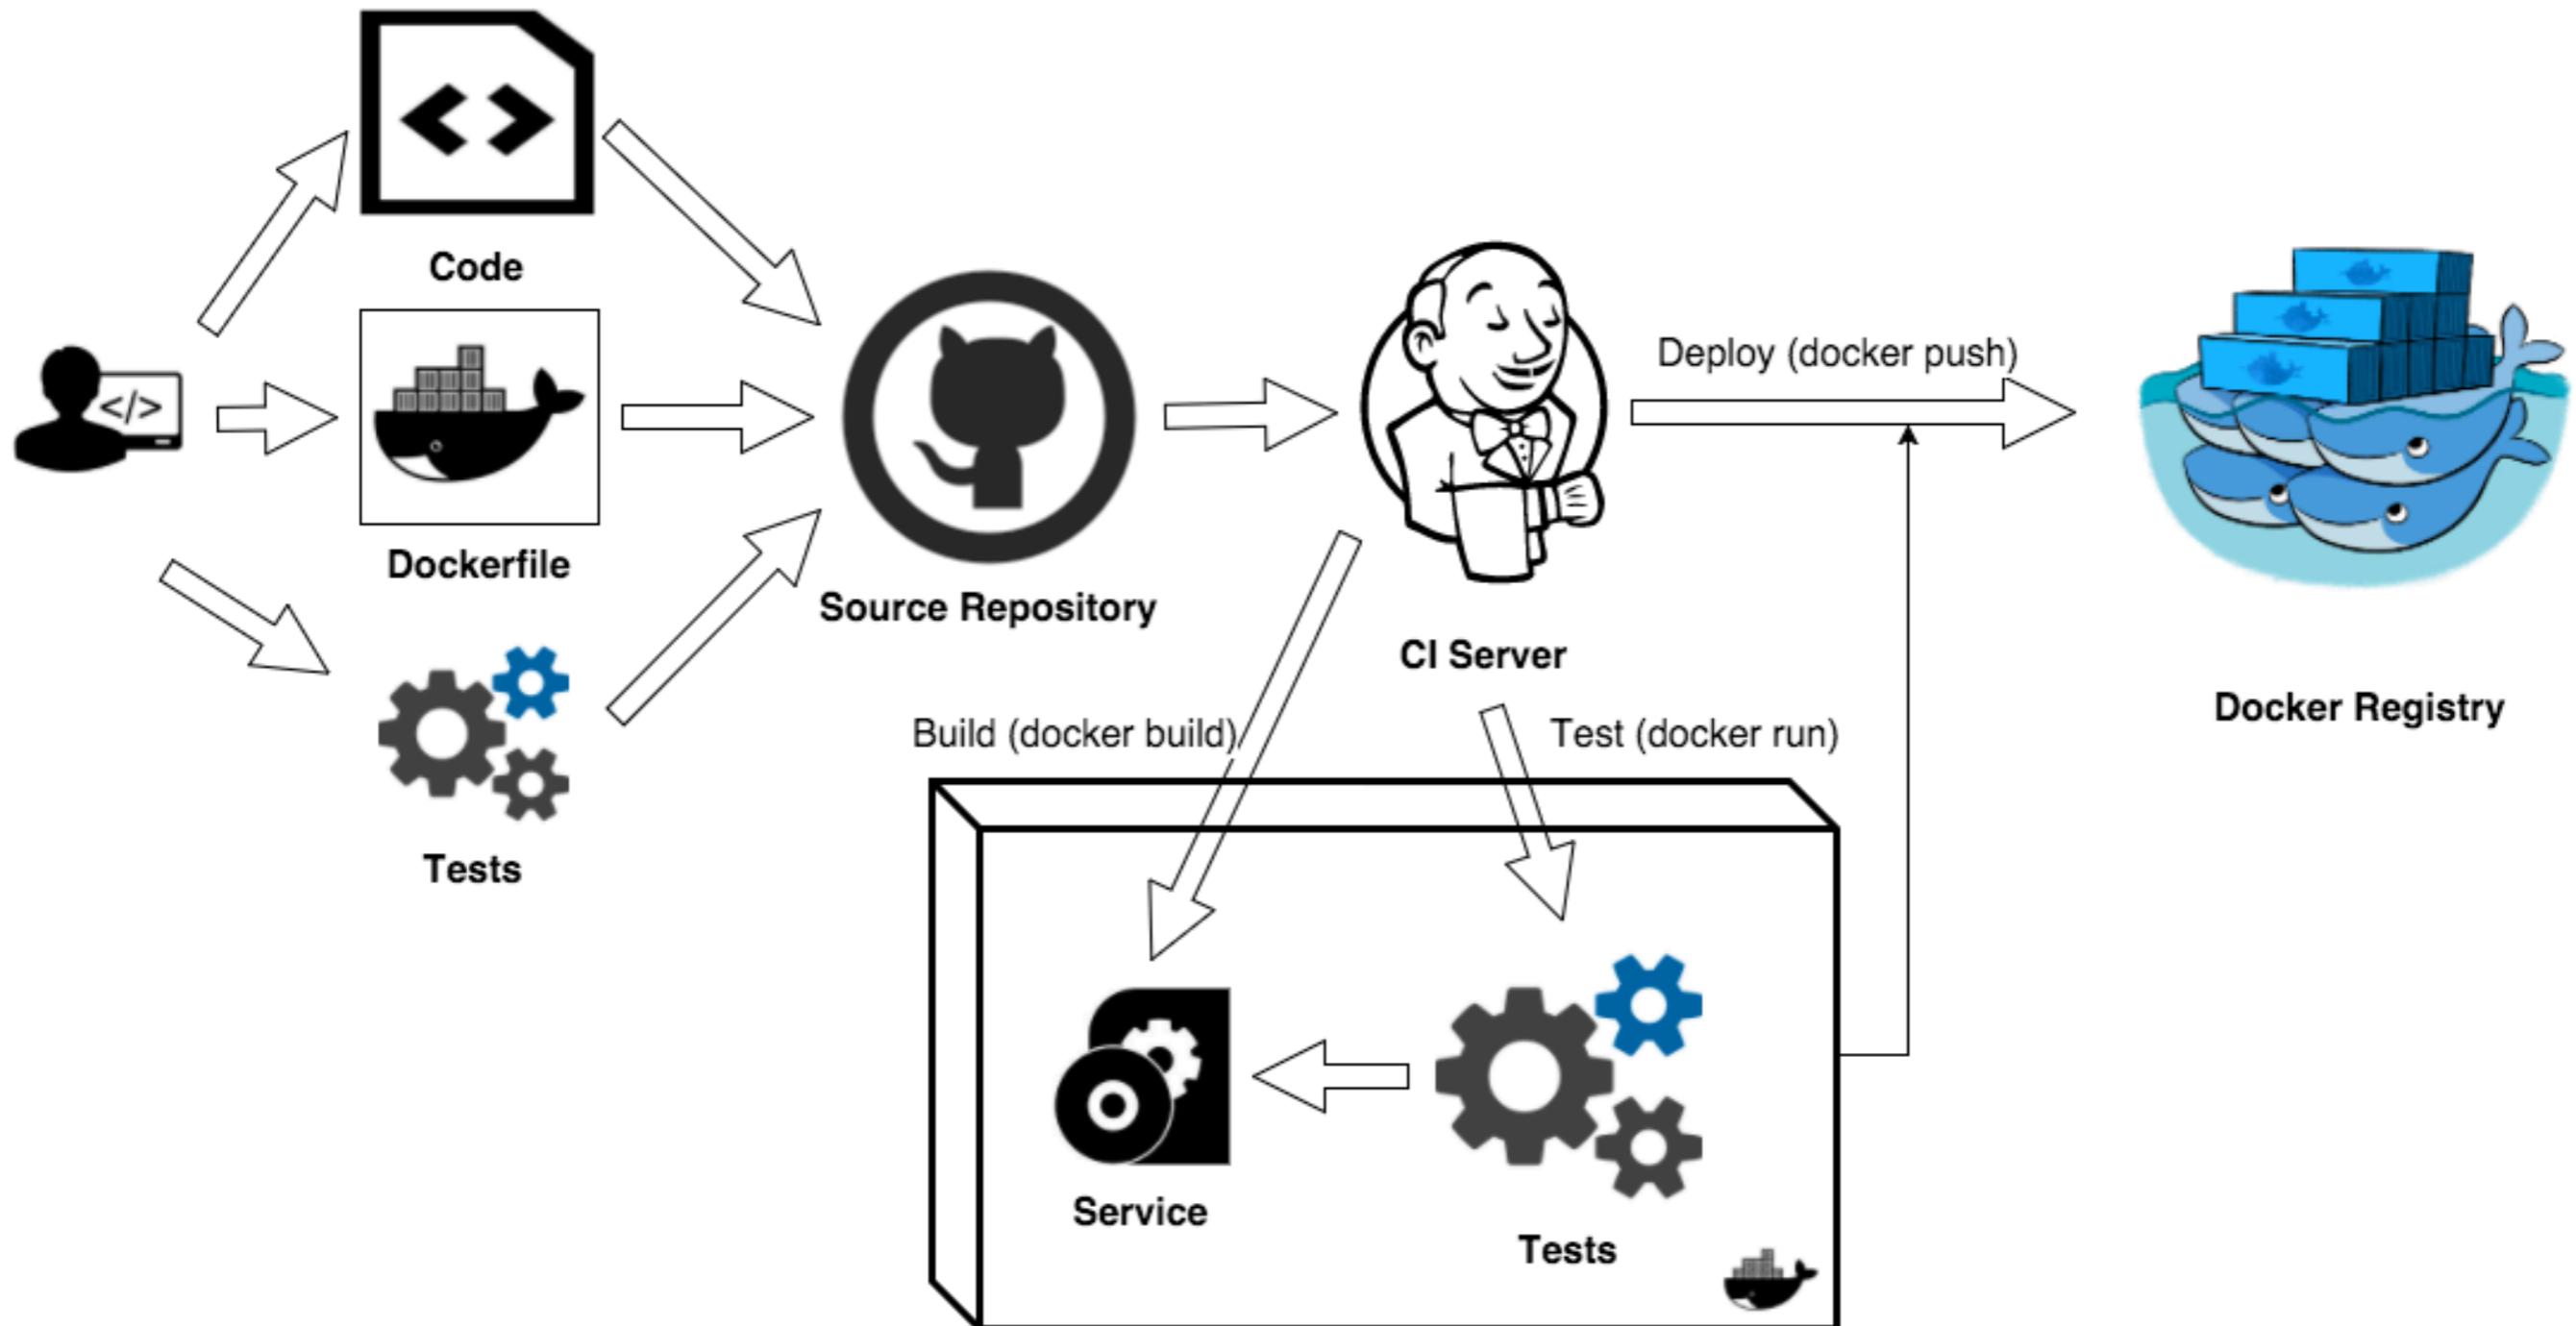
Containers

# Portabilità

		?	?	?	?	?	?
			?	?	?	?	?
		?	?	?	?	?	?
		?	?	?	?	?	?
		?	?	?	?	?	?
		?	?	?	?	?	?

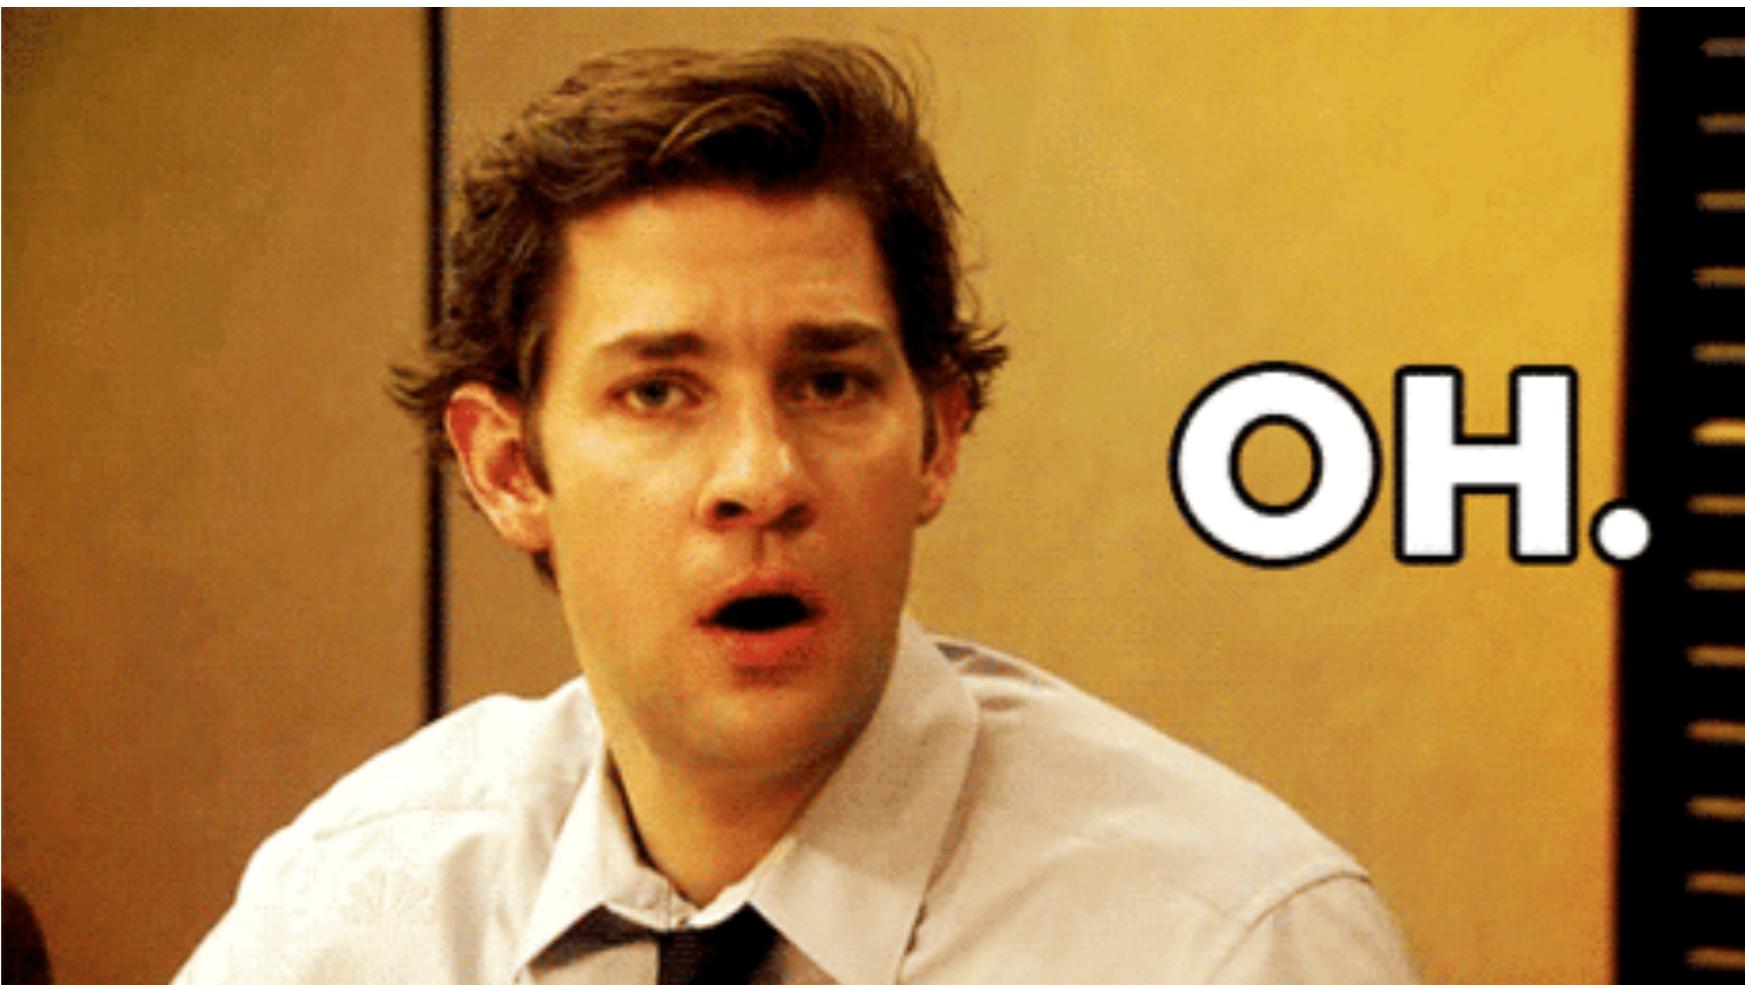
The matrix from hell

# Flessibilità



Build, ship and run any app, anywhere

# Sicurezza



# Sicurezza

- Don't give root
- If the application needs root, give looks-like-root
- If that's not sufficient, give root, but build another wall

# Docker

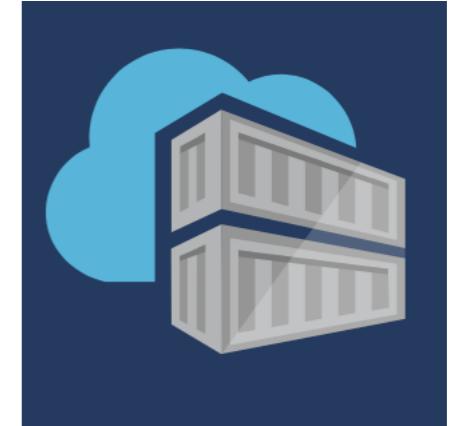
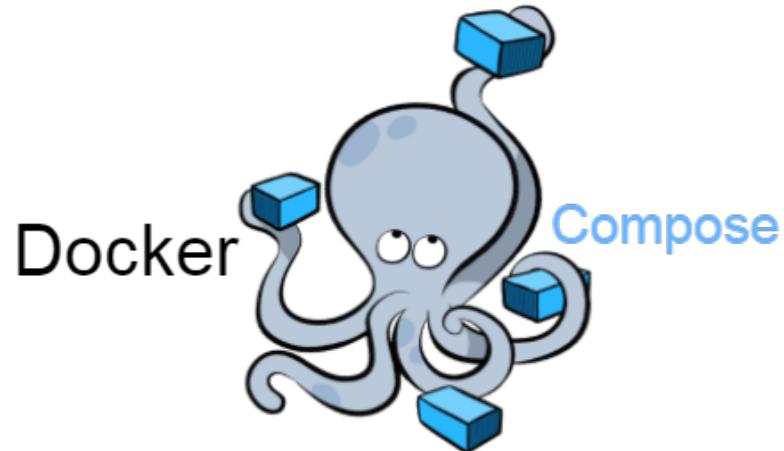
## Immagini e container

An image is an inert, immutable, file that's essentially a snapshot of a container.

Images are created with the build command, and they'll produce a container when started with run. Images are stored in a Docker registry



# Immagini



Immagini private  
(Dockerfile)

Docker Hub

Registri privati  
<https://github.com/docker/distribution>

# Come definire un'immagine e avviarla?



Scrivere un  
dockerfile

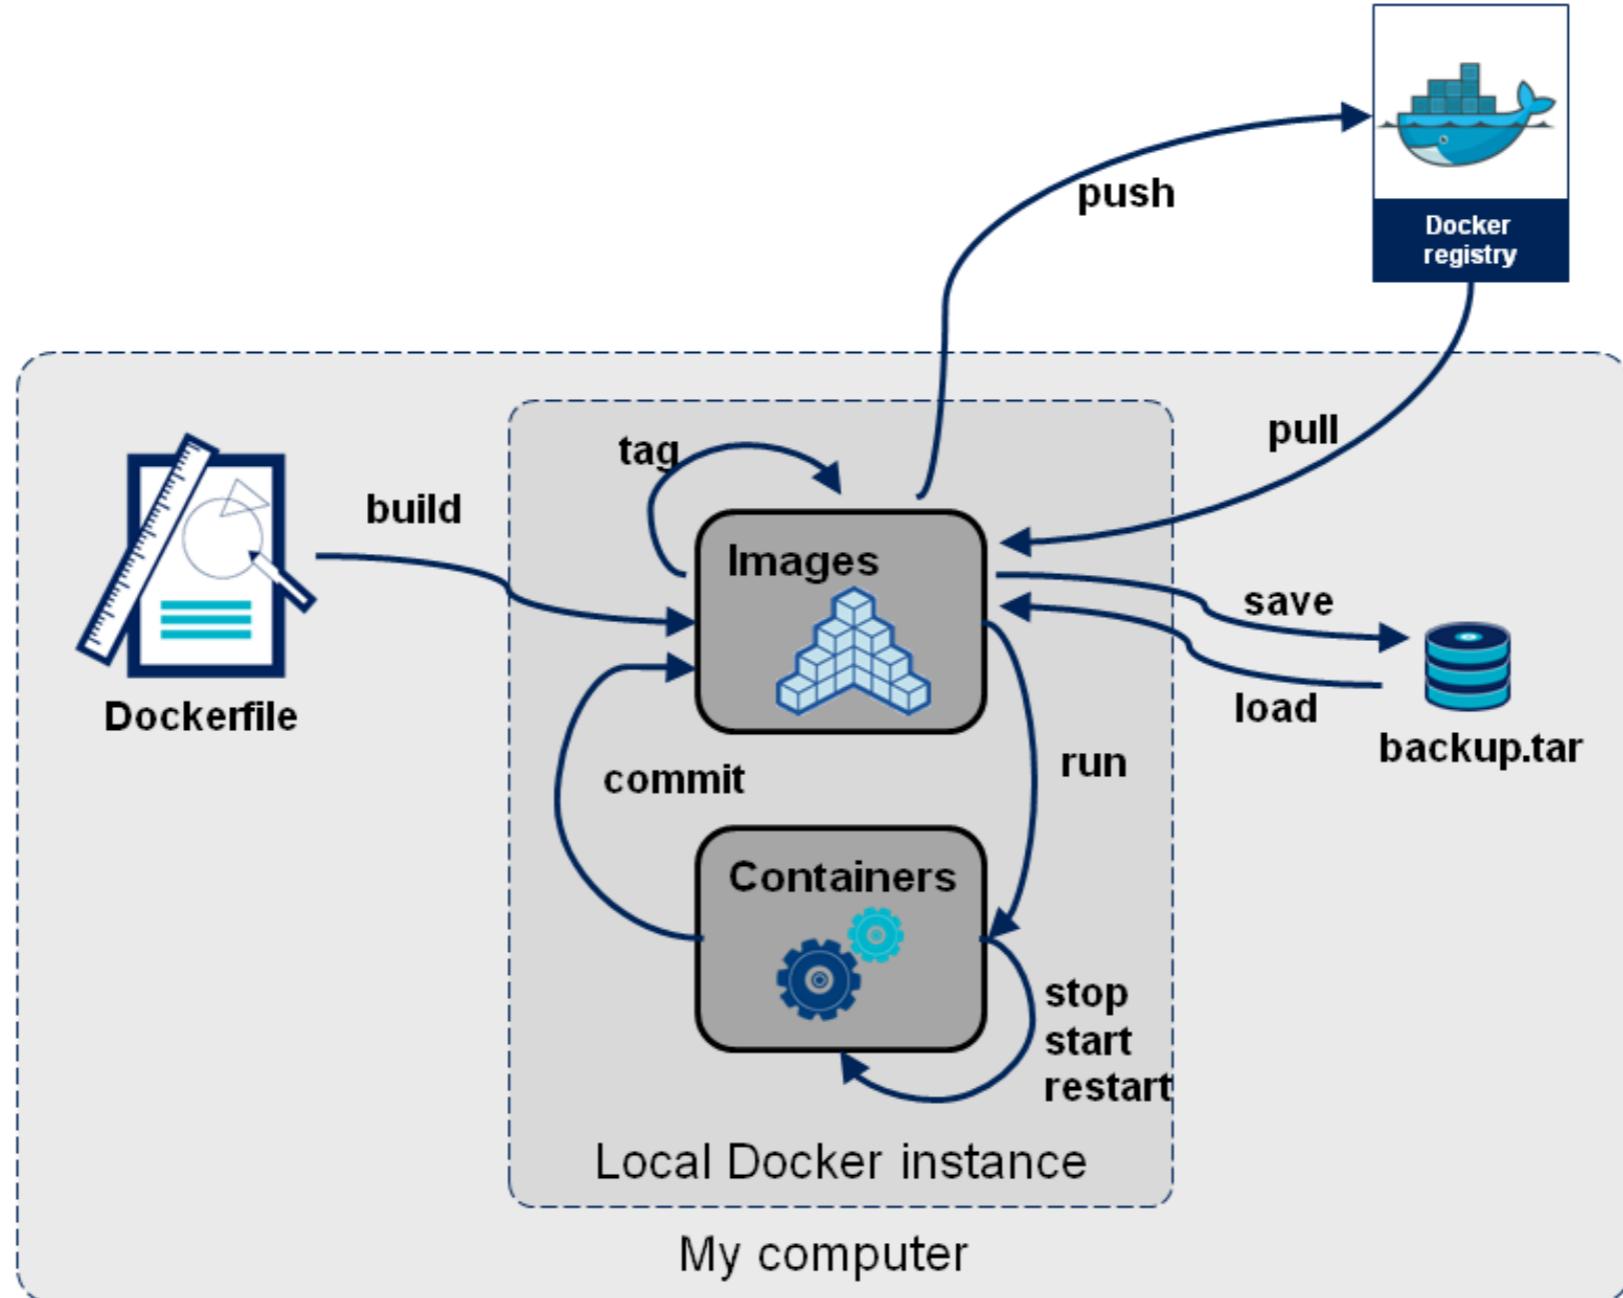


Fare il built  
dell'immagine



Avviare  
il container

# Come definire un'immagine e avviarla?



# Container programming

## Dockerfile

```
# Version: 0.0.1
FROM sequenceiq/hadoop-docker:2.7.1
LABEL maintainer="alessandrov87 (https://github.com/AlessandroVaccarino)"
# Download Pig
RUN curl http://apache.mirror.anlx.net/pig/latest/pig-0.17.0.tar.gz | tar -zx -C /usr/local
ENV PATH /usr/local/pig-0.16.0/bin:$PATH
ENV PATH /usr/local/hadoop/bin:$PATH
# Download and Init Hive
RUN curl http://apache.mirror.anlx.net/hive/stable/apache-hive-1.2.2-bin.tar.gz | tar -zx -C /usr/local
ENV PATH /usr/local/apache-hive-1.2.1-bin/bin:$PATH
# Download Zookeeper (for HBase)
RUN curl http://apache.mirror.anlx.net/zookeeper/stable/zookeeper-3.4.12.tar.gz | tar -zx -C /usr/local
ENV PATH /usr/local/zookeeper-3.4.12/bin:$PATH
# Download HBase
RUN curl http://apache.mirror.anlx.net/hbase/stable/hbase-1.2.6-bin.tar.gz | tar -zx -C /usr/local
ENV PATH /usr/local/hbase-1.2.6/bin:$PATH
# Configure HBase
RUN rm /usr/local/hbase-1.2.6/conf/hbase-site.xml
COPY hbase-site.xml /usr/local/hbase-1.2.6/conf/
# Configure bootstrap file
RUN rm /etc/bootstrap.sh
COPY bootstrap.sh /etc/
RUN chown root:root /etc/bootstrap.sh
RUN chmod 700 /etc/bootstrap.sh
# Expose Ports
EXPOSE 10020
EXPOSE 50070
EXPOSE 50030
EXPOSE 8088
EXPOSE 16000
EXPOSE 16010
EXPOSE 16020
EXPOSE 16030
EXPOSE 16100
```

# Dockerfile

FROM	Sets the base image for subsequent instructions
ENV	Set environment variable in container
RUN	Execute command in the image and commit results
ADD	Add file from host/URL to container
VOLUME	Specify directory that lives outside union fs
EXPOSE	Specify ports to open between containers
CMD	Default command when executing container

# Docker Gestione

docker ps	Mostra l'elenco dei container attivi
docker stop ubuntu	Ferma un container
docker rm ubuntu	Rimuove un container
docker run ubuntu ...	Avvia un container

# Docker

## Hello World

```
sudo apt install docker.io
```

```
sudo docker run ubuntu:14.04 /bin/echo 'Hello world'
```

```
master@master-pc:~$ sudo docker run ubuntu:14.04 /bin/echo 'Hello world'  
Unable to find image 'ubuntu:14.04' locally  
14.04: Pulling from library/ubuntu  
d6fdcbe24ed5: Downloading [=====] 61.77MB/67.19MB  
1115217119d4: Download complete  
8180a23d2c44: Download complete  
ceab6b3c2a5e: Download complete
```

# Docker

## Hello World

```
docker run ubuntu:14.04 /bin/echo 'Hello world'  
Unable to find image 'ubuntu:14.04' locally  
Pulling repository ubuntu  
6b4e8a7373fe: Download complete  
511136ea3c5a: Download complete  
b18d0a2076a1: Download complete  
67b66f26d423: Download complete  
25c4824a5268: Download complete  
8b1c48305638: Download complete  
c900195dcbf3: Download complete  
Hello world
```

# Recap

- Docker permette di costruire sistemi modulari e portabili
- Possiamo distribuirlo ovunque: on premises o in cloud
- Dobbiamo orchestrare i container: Kubernetes
- Iniziate ad usare Docker e Kubernetes!!!
- Google ha introdotto i container più di 10 anni fa

# Server-less architecture

## **Serverless computing**

(also known as *functions as a service*)  
is a new cloud computing abstraction that makes it easier to  
write robust, large-scale web services.

In serverless computing, programmers write what are called  
serverless functions, which are programs that respond to  
external events.

When **demand for the serverless function spikes**,  
the platform automatically allocates additional hardware and  
manages load-balancing

When **demand falls**, the platform silently deallocates idle  
resources; and when the platform detects a failure, it  
transparently retries affected requests.

# Serverless

## Function as a service

- Allows for granularly billing by cloud service providers.
- An event-driven programming model and architecture.
- Handles most, if not all, operational concerns.
- Removes control over resources.
- Often short-lived.
- Auto-scaling and provisioning.
- Implicit high availability.
- Implicit fault tolerance.
- No servers

serverless

Search term

+ Compare

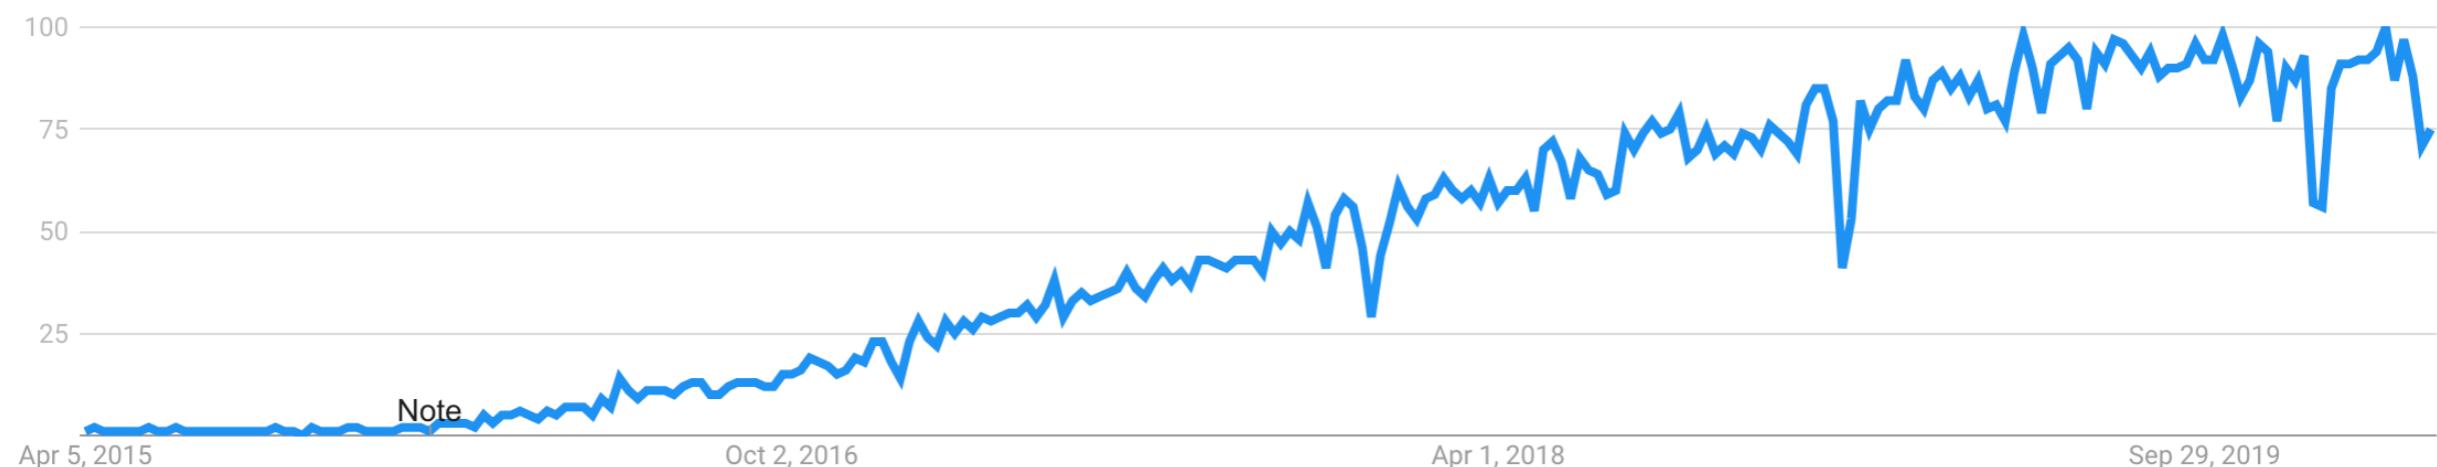
Worldwide ▾

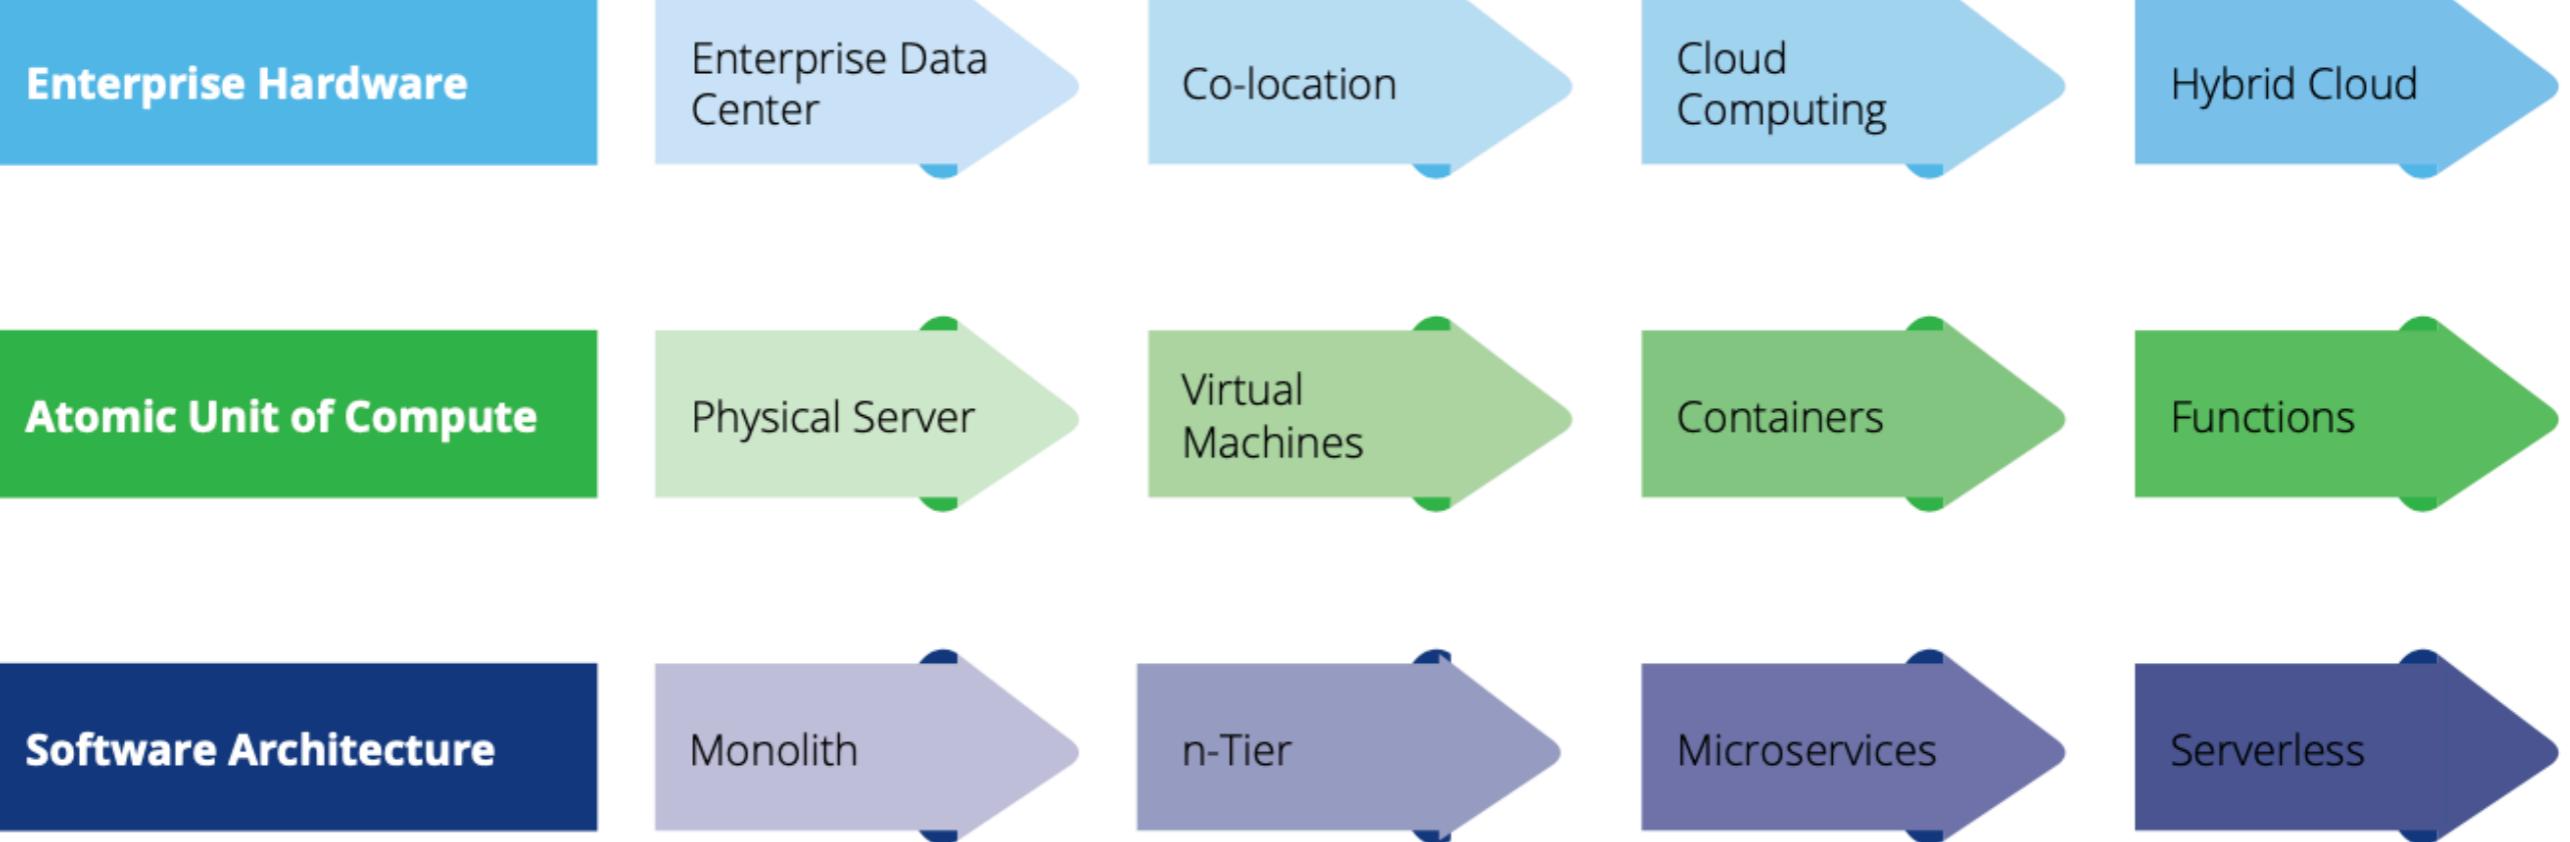
Past 5 years ▾

All categories ▾

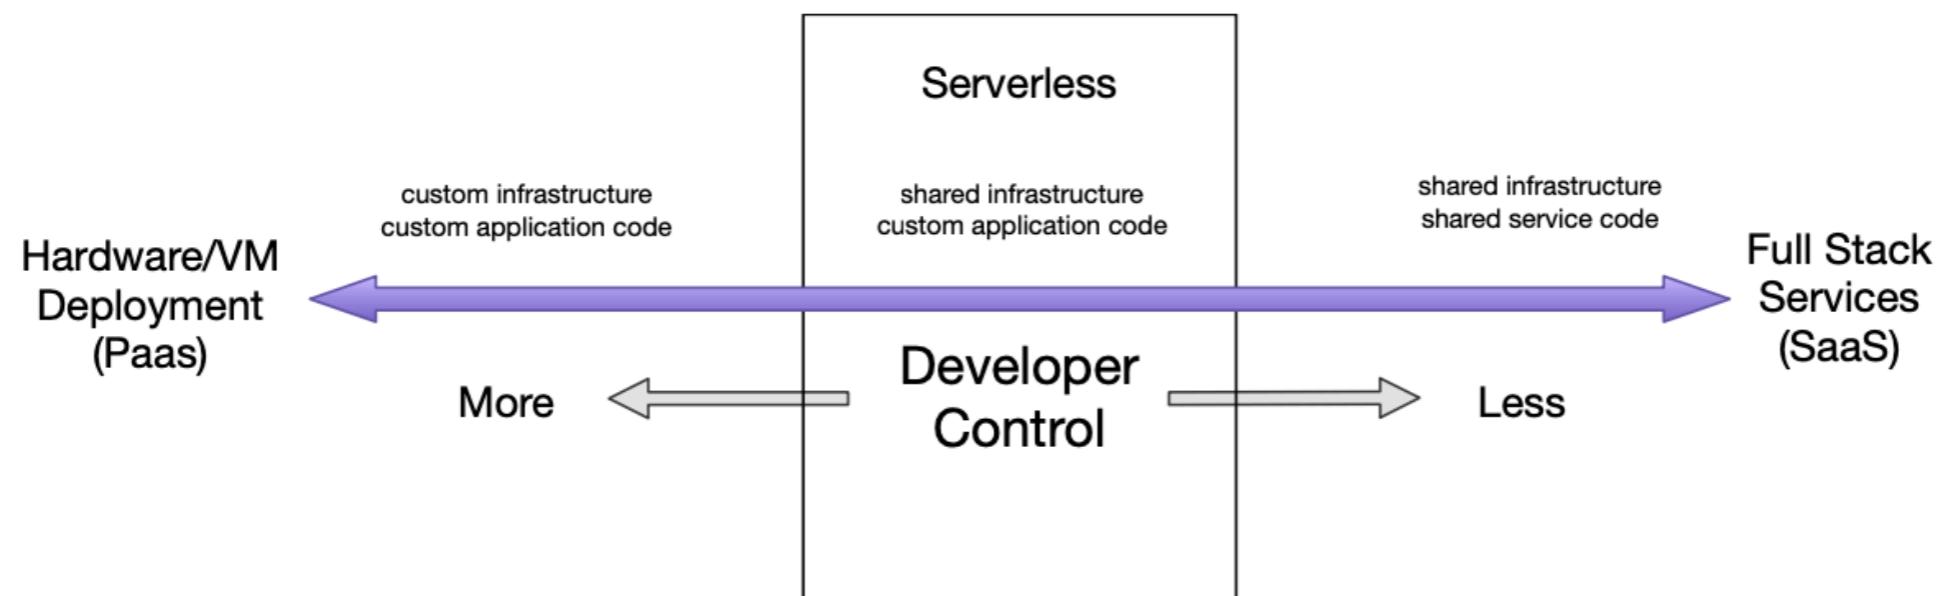
Web Search ▾

Interest over time ?

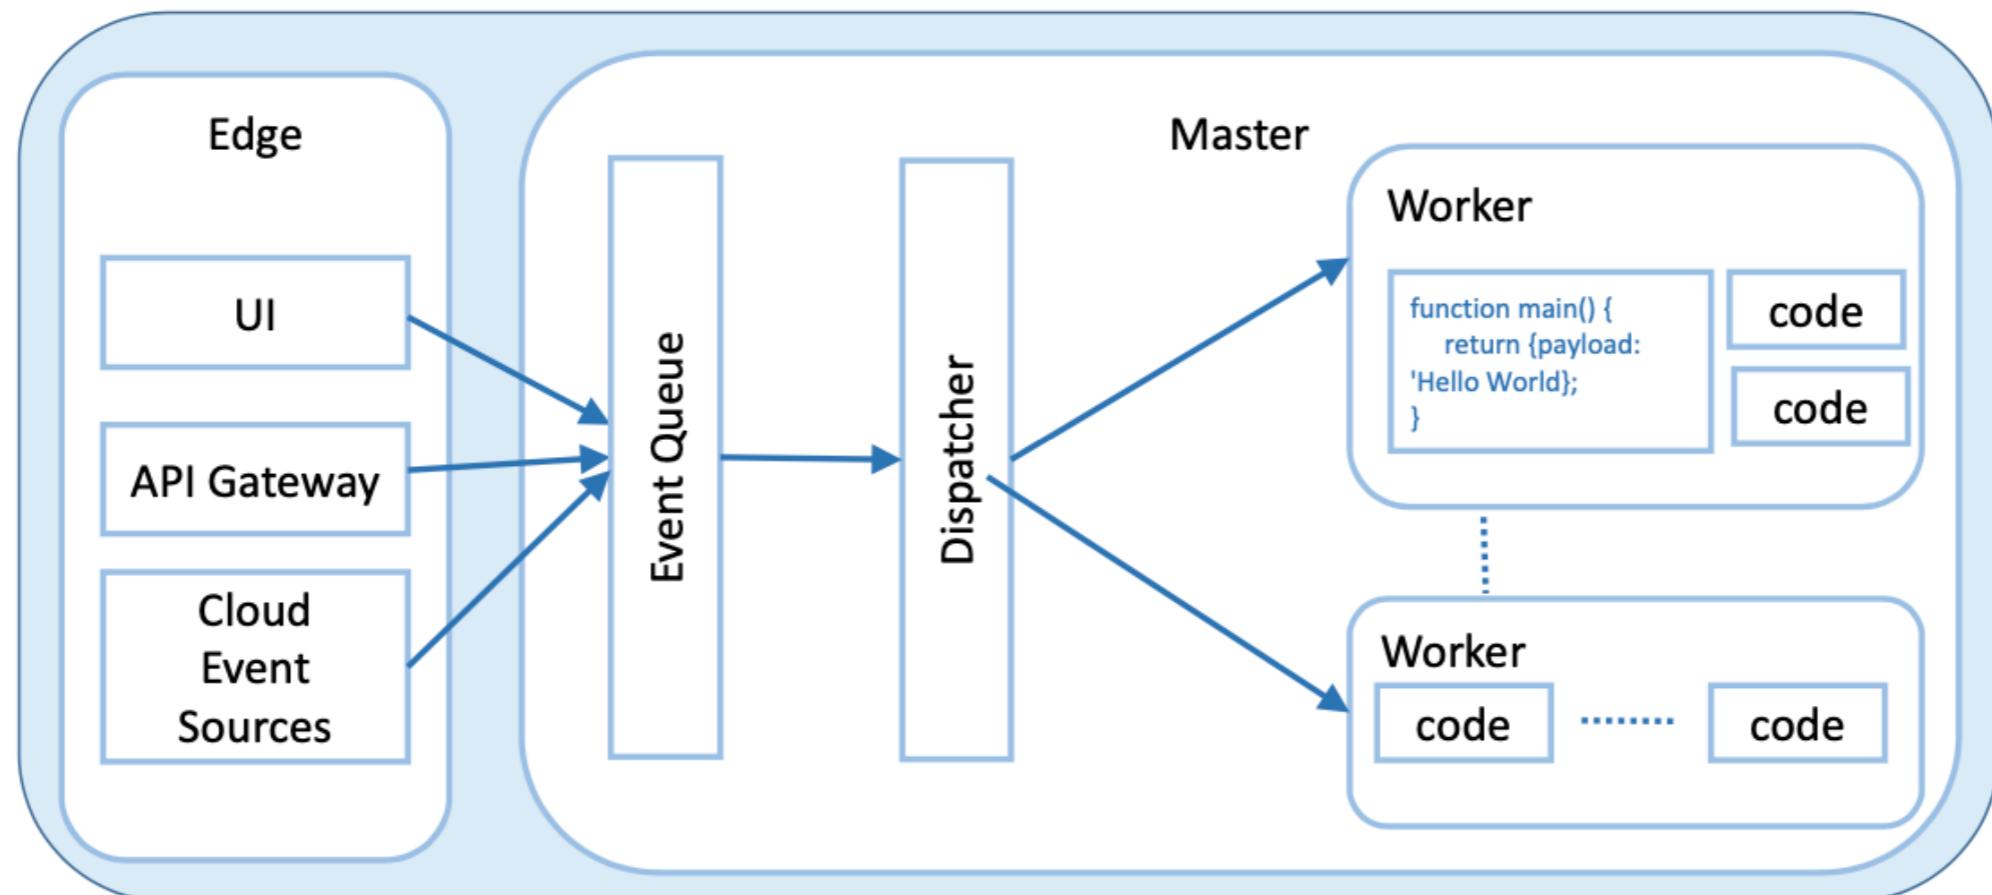




# Developer control and Serverless computing



# Serverless platform architecture



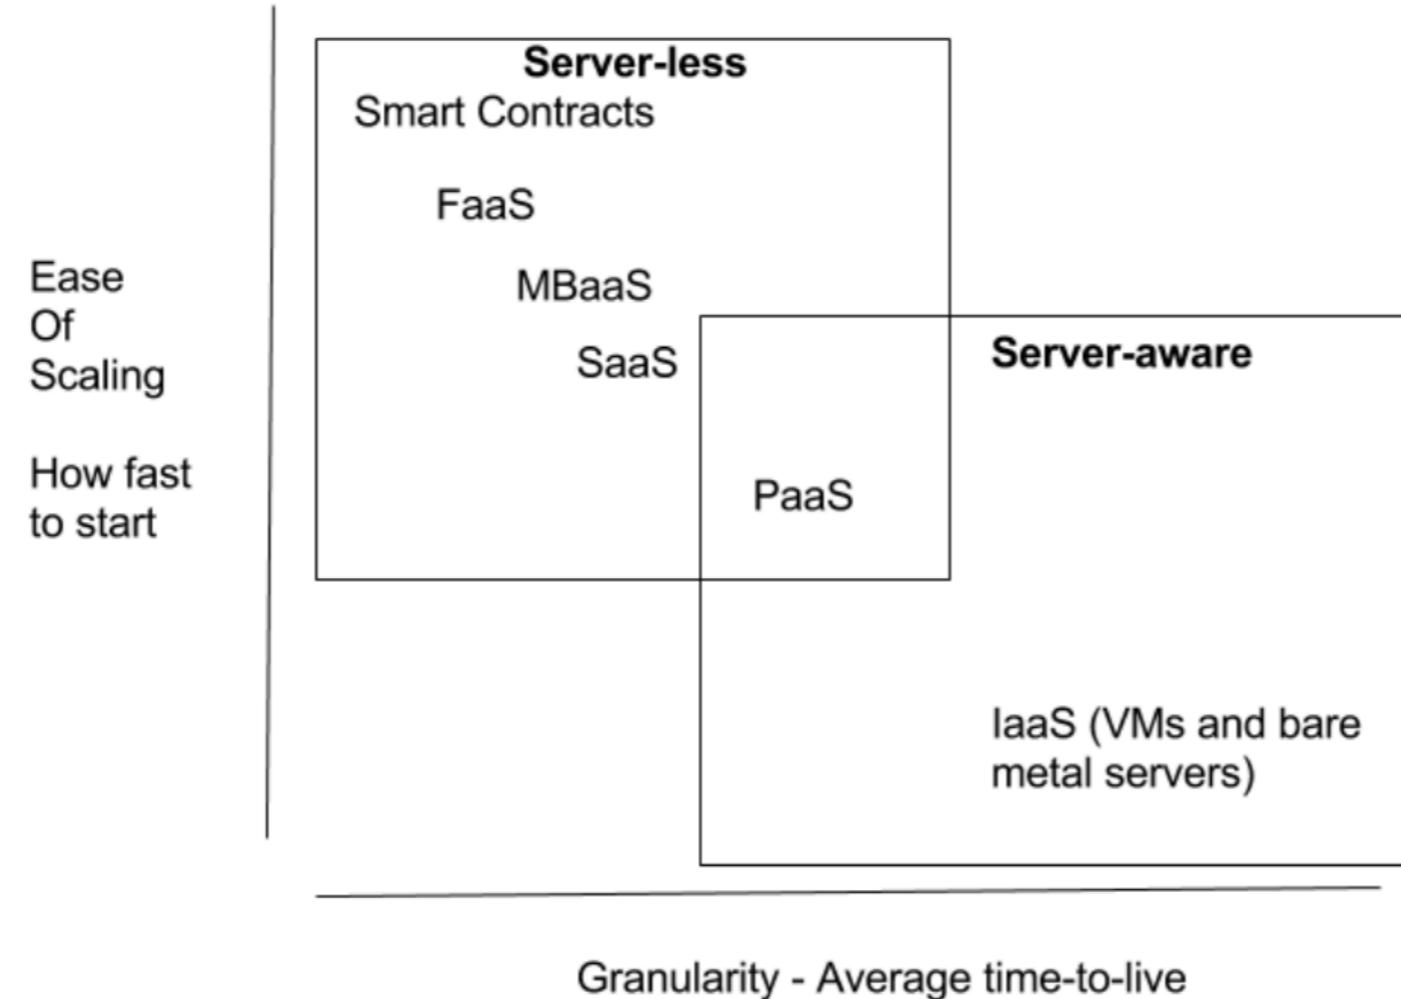
# Characteristics

- Cost
- Performance and limits
- Programming languages
- Programming model
- Composability
- Deployment
- Security and accounting
- Monitoring and debugging

# Programming model

```
function main(params, context) {  
    return {payload: 'Hello, ' + params.name  
           + ' from ' + params.place};  
}
```

# Time-to-live vs Scaling



The figure is showing relation between time-to-live (x-axis) and ease-of-scaling (y-axis). Server-aware compute (bare metal, VMs, IaaS) has long time to live and take longer to scale (time to provision new resources); server-less compute (FaaS, MBaaS, PaaS, SaaS) is optimized to work on multiple servers and hide server details.

@Serverless Computing: Current Trends and Open Problems

Ioana Baldini, Paul Castro, Kerry Chang, Perry Cheng, Stephen Fink, Vatche Ishakian, Nick Mitchell, Vinod Muthusamy, Rodric Rabbah, Aleksander Slominski, Philippe Suter



## 01

**API Gateway:** The API Gateway acts as the communication layer between the frontend and the FaaS layer. It maps REST API endpoints with the respective functions that runs the business logic. With servers out of equation there is no need for deploying and manage load balancers also in this model.

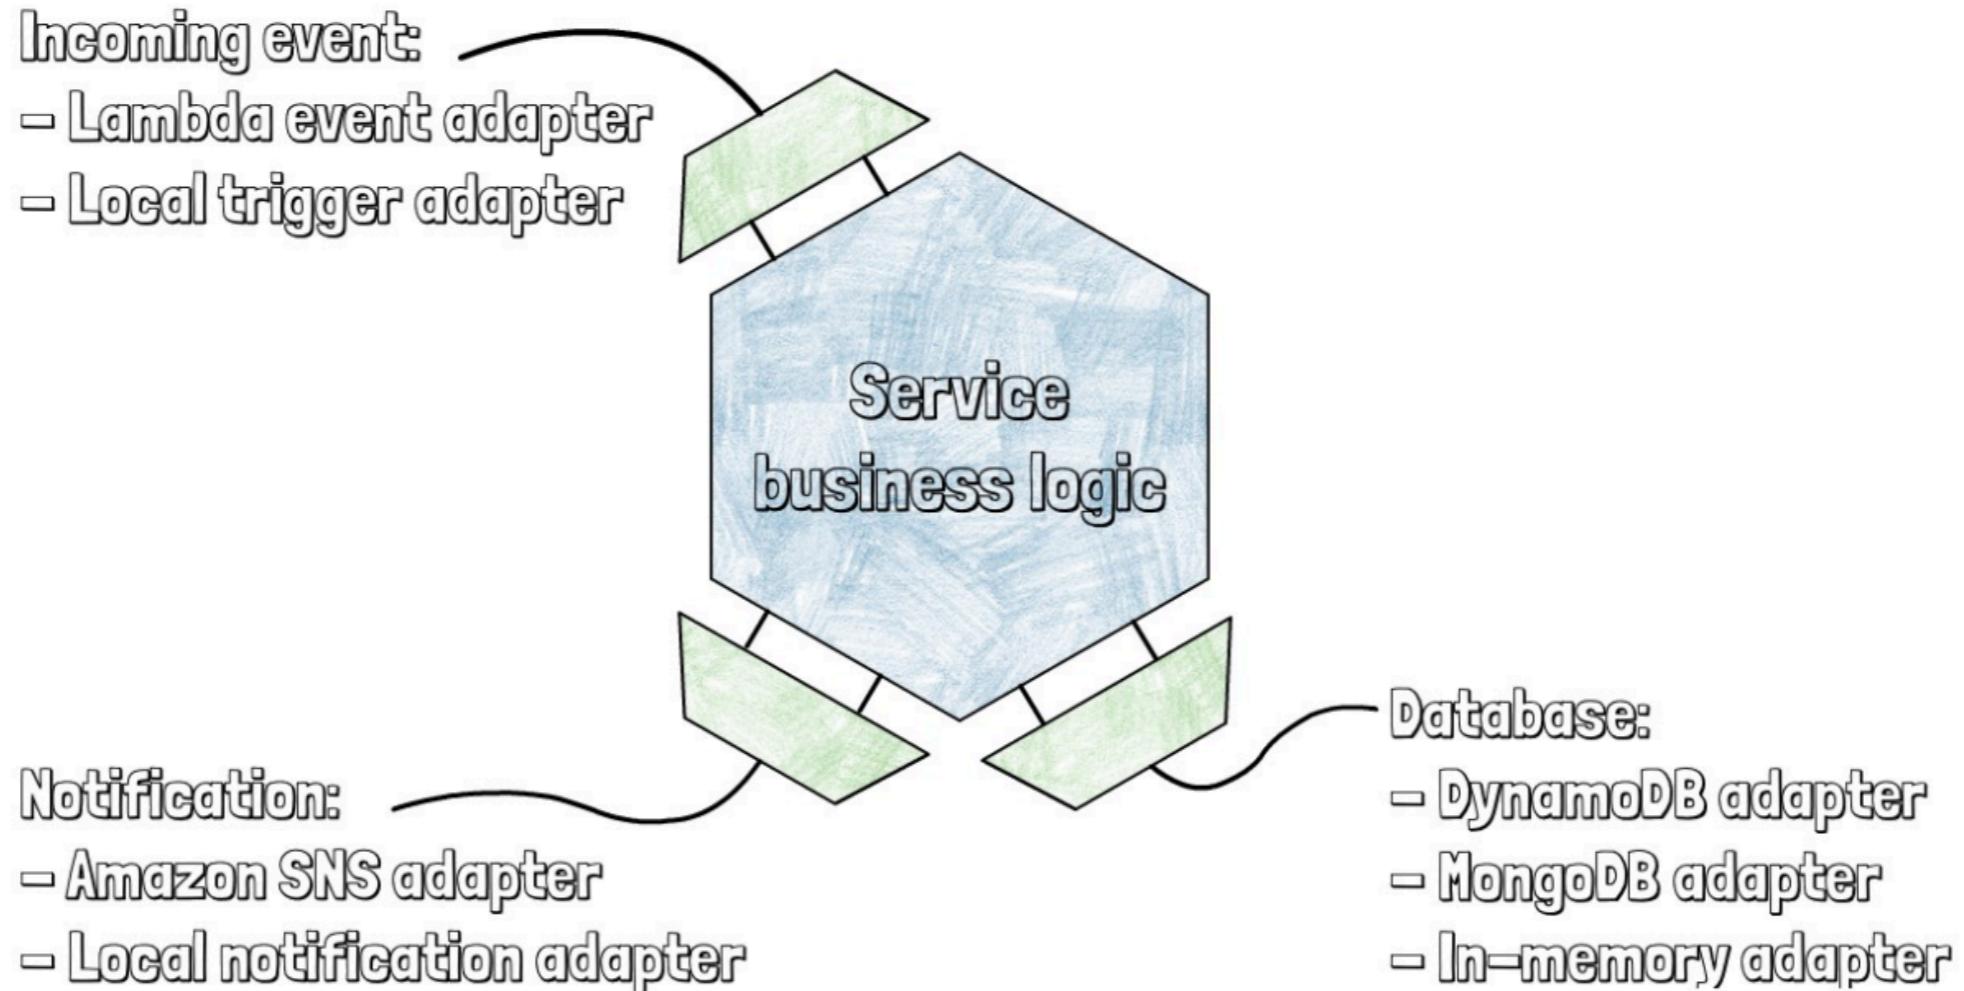
## 02

**Functions or Function as a Service (FaaS):** This is the layer that executes specific business logic (or code) with the cloud provider providing the level of abstraction in terms of executing the business logic.

## 03

**Backend as a Service (BaaS):** This is essentially a cloud based distributed NoSQL database which essentially removes database administration overheads.

# Service Design



**Serverless means:**

Greater agility

Less overhead

Better focus

Increased scale

More flexibility

Faster time to market

